



Chapter 2: Operating-System Structures





Chapter 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines
- System Boot





Objectives

- To describe the services an OS provides to users, processes, and other systems
- To discuss the various ways of structuring an OS
- To explain how OS boot





Operating System Services

- User interface
- Program execution
- I/O operations
- File-system manipulation
- Communications
- Error detection
- Resource allocation
- Accounting
- Protection and security





User Operating System Interface - CLI

- CLI allows direct command entry
 - Sometime implemented by kernel
 - Sometime implemented by program
 - Sometimes multiple flavors implemented – shells
 - Fetches a command from user and executes it

- Commands
 - Sometimes built-in
 - sometimes just names of programs
 - Adding new features doesn't require shell modification





User Operating System Interface - GUI

- **User-friendly desktop metaphor interface**
 - Usually mouse, keyboard, and monitor/touchscreen
 - WIMP-Window, Icon, Menu, Pointing device
 - Invented at Xerox PARC in 1973

- **Many systems include both CLI and GUI**
 - Windows is GUI with CLI “command” shell
 - Apple Mac OS X as “Aqua” GUI interface with shells
 - Linux and Solaris are CLI with optional GUI interfaces





System Calls (系统调用)

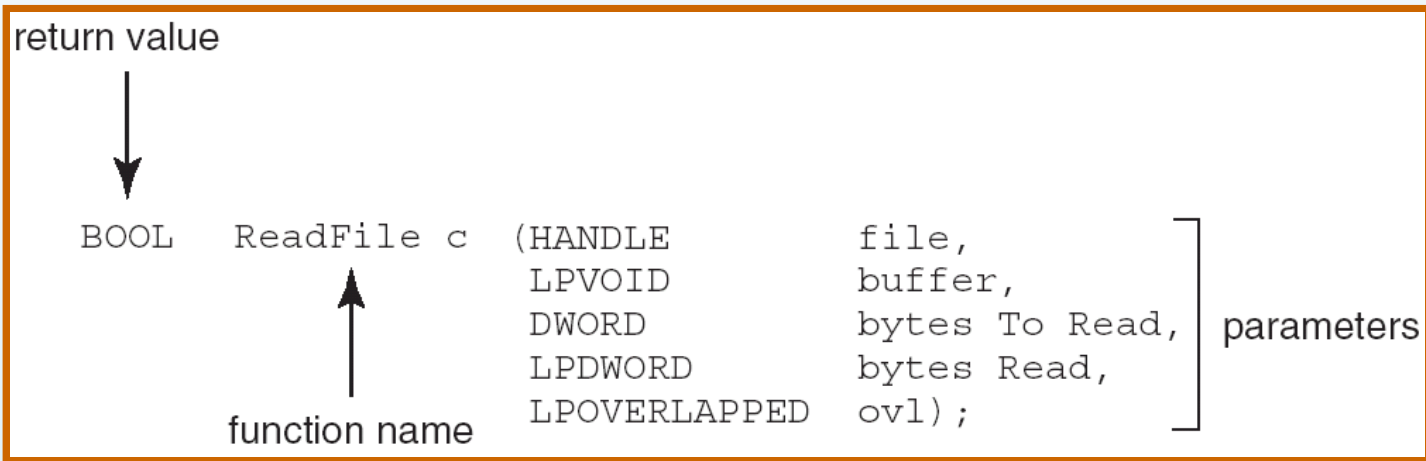
- Programming interface to the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level Application Program Interface (API)
- Three most common APIs
 - Win32 API for Windows
 - POSIX API for POSIX-based systems
 - ▶ including all versions of UNIX, Linux, and Mac OS X
 - Java API for the Java virtual machine (JVM)





Example of Standard API

■ ReadFile() in the Win32 API—a function for reading from a file



■ A description of the parameters passed to ReadFile()

- **HANDLE file**—the file to be read
- **LPVOID buffer**—a buffer where the data will be read into and written from
- **DWORD bytesToRead**—the number of bytes to be read into the buffer
- **LPDWORD bytesRead**—the number of bytes read during the last read
- **LPOVERLAPPED ovl**—indicates if overlapped I/O is being used





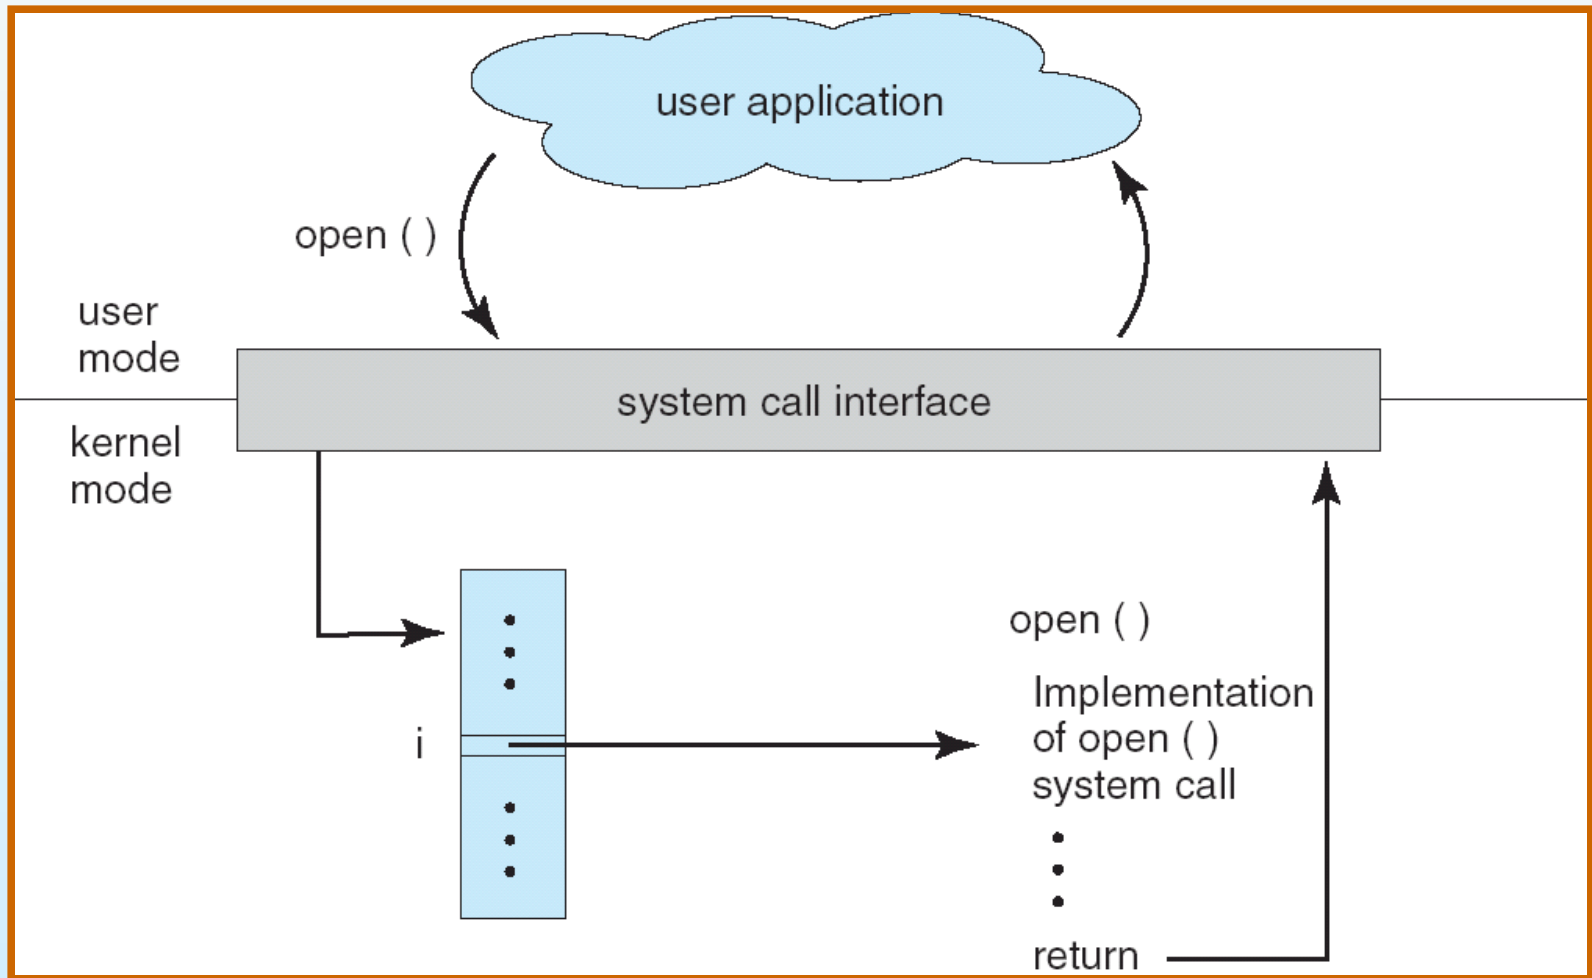
System Call Implementation

- A number associated with each system call
- System-call interface
 - Maintains a table indexed by system-call numbers
 - Invokes intended system call in OS kernel
 - Returns status of the system call and any return values
- The caller needn't know how the system call is implemented
 - Just obey API and understand what OS will do
 - APIs are managed by run-time library
 - ▶ Functions built into libraries included with compiler





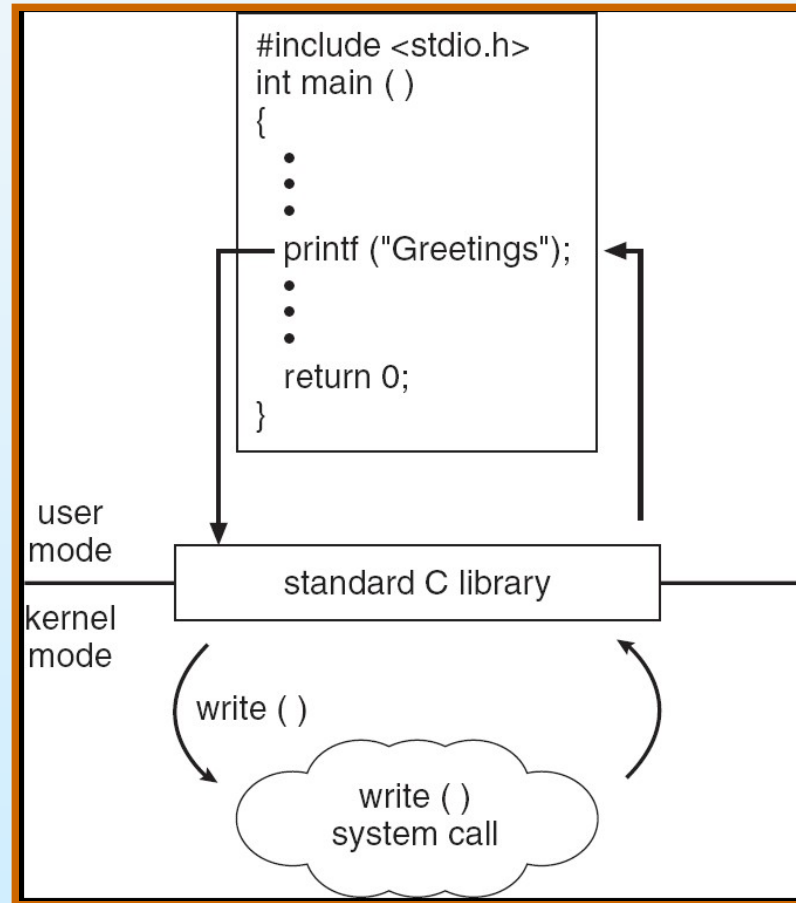
API – System Call – OS Relationship





Standard C Library Example

- C program invoking printf() library call, which calls write() system call





Linux strace Hello World

```
$ strace ./hello
```





System Call Parameter Passing

- **Pass parameters to the OS**
 - **Simplest: pass the parameters in *registers***
 - ▶ May be more parameters than registers
 - **Parameters stored in a *block*, or table, in memory, and address of block passed in a register**
 - ▶ This approach taken by Linux and Solaris
 - **Parameters *pushed* onto the *stack* by the program and *popped* by the OS**
 - **Block and stack do not limit the number or length of parameters**





Types of System Calls

- **Process control**
- **File management**
- **Device management**
- **Information maintenance**
- **Communications**





System Programs

- **System programs provide a convenient environment for program development and execution.**
 - **File manipulation**
 - **Status information**
 - **File modification**
 - **Programming language support**
 - **Program loading and execution**
 - **Communications**
 - **Application programs**
- **Most users' view of the OS is defined by system programs, not the system calls**





Operating System Design and Implementation

- Design and Implementation of OS not “solvable”, but some approaches have proven successful
- Internal structure of different OS can vary widely
- Start by defining goals and specifications
- Affected by choice of hardware, type of system
- *User* goals and *System* goals
 - User goals
 - ▶ Convenient to use, easy to learn, reliable, safe, and fast
 - System goals
 - ▶ Easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient





Operating System Design and Implementation (Cont.)

■ Important principle to separate

Policy: What will be done?

Mechanism: How to do it?

■ The separation of policy from mechanism is a very important principle

- it allows maximum flexibility if policy decisions are to be changed later

Which language is your OS implemented by?





Operating System Implementation

- Long long ago, written by ASM
 - MS-DOS
- Now, written by C and C++, minor ASM
 - Unix, Linux, Windows...
- Advantages of higher-level language
 - How about ASM vs. C?
 - Portability
- Disadvantages of higher-level language
 - Performance





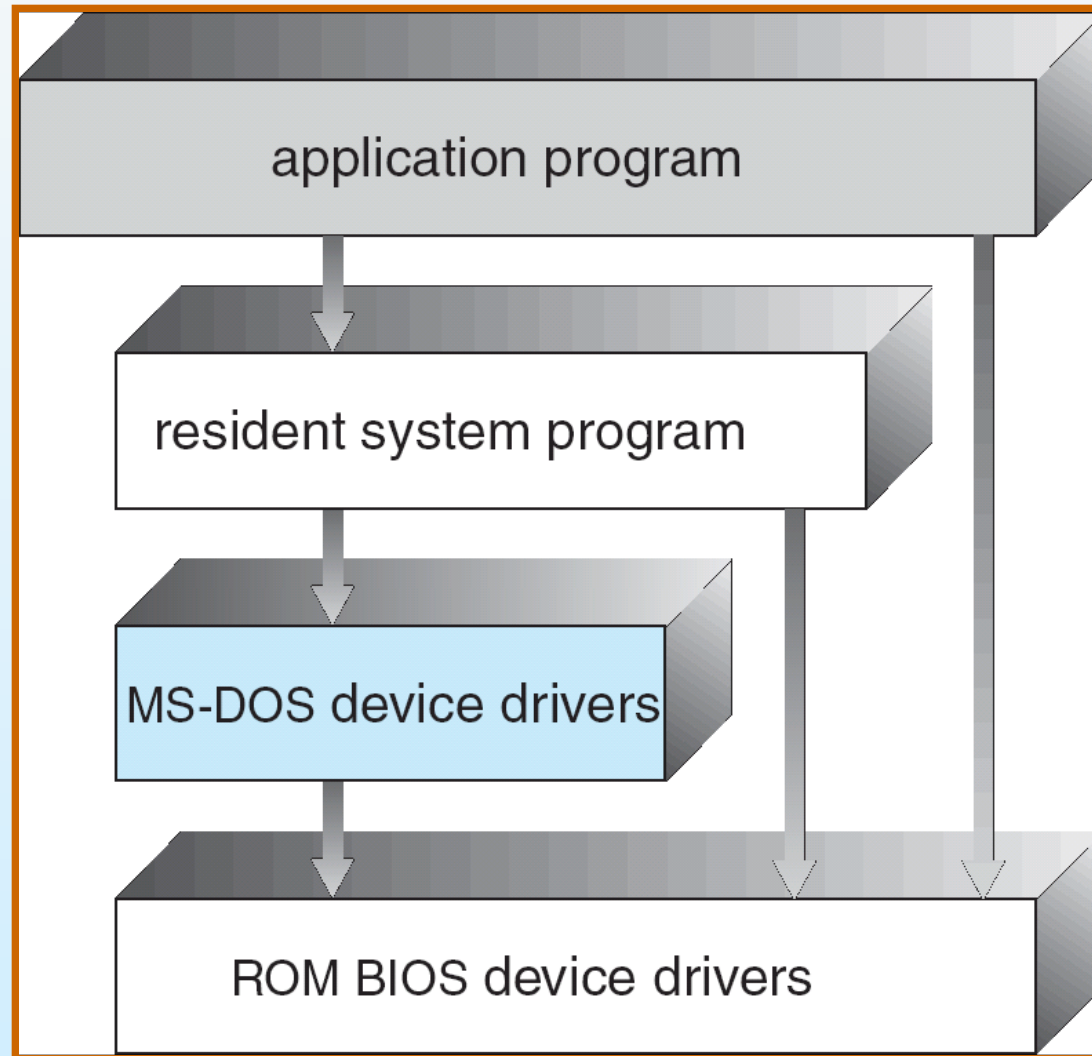
Operating System Structure

- Simple structure
- Layered approach
- Microkernel
- Modules
- Virtual machine





MS-DOS Layer Structure





Simple Structure

■ MS-DOS

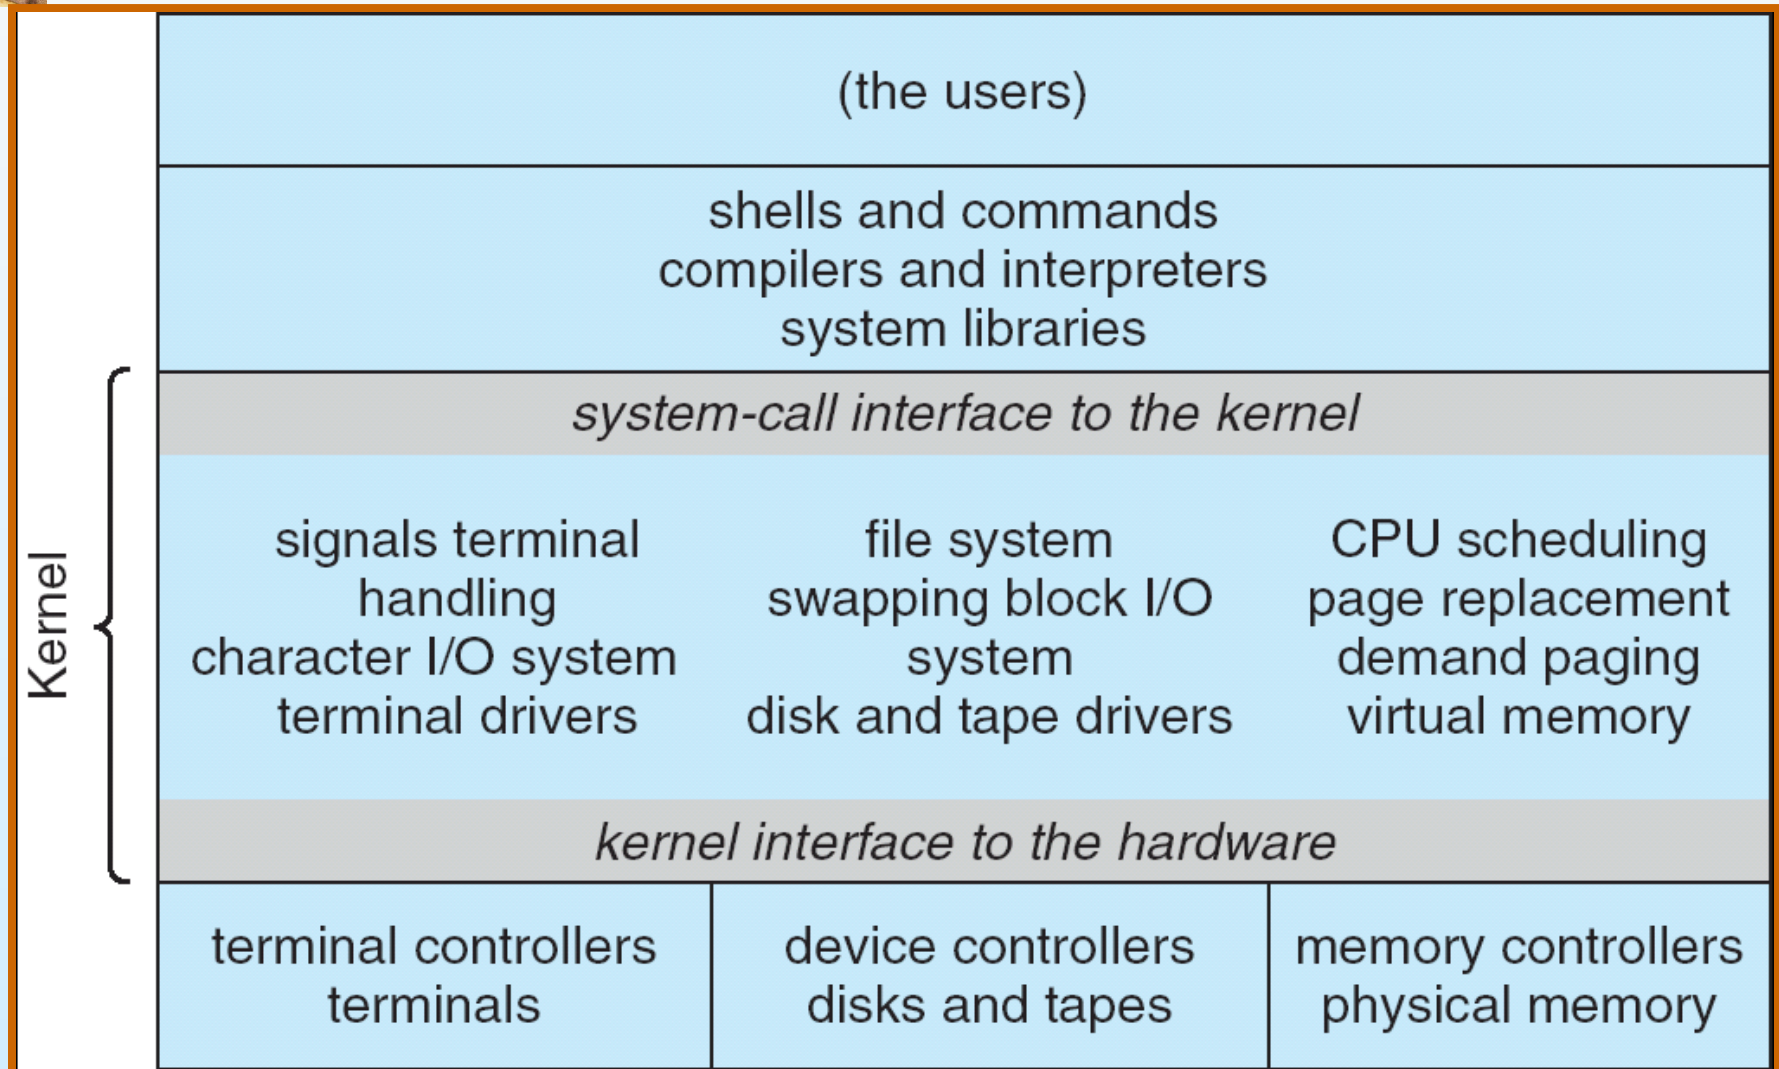
- Provide the most functionality in the least space
- Not divided into modules
- Although has some structure, its interfaces and levels of functionality are not well separated

■ MS-DOS was limited by hardware of its era





UNIX System Structure






UNIX

- **The original UNIX OS had limited structuring**
- **Consists of two separable parts**
 - **System programs**
 - **The kernel**
 - ▶ **Everything below the system-call interface and above the hardware**
 - **The file system, CPU scheduling, memory management, and other OS functions**
 - ▶ **A large number of functions for one level**





Layered Operating System





Layered Approach

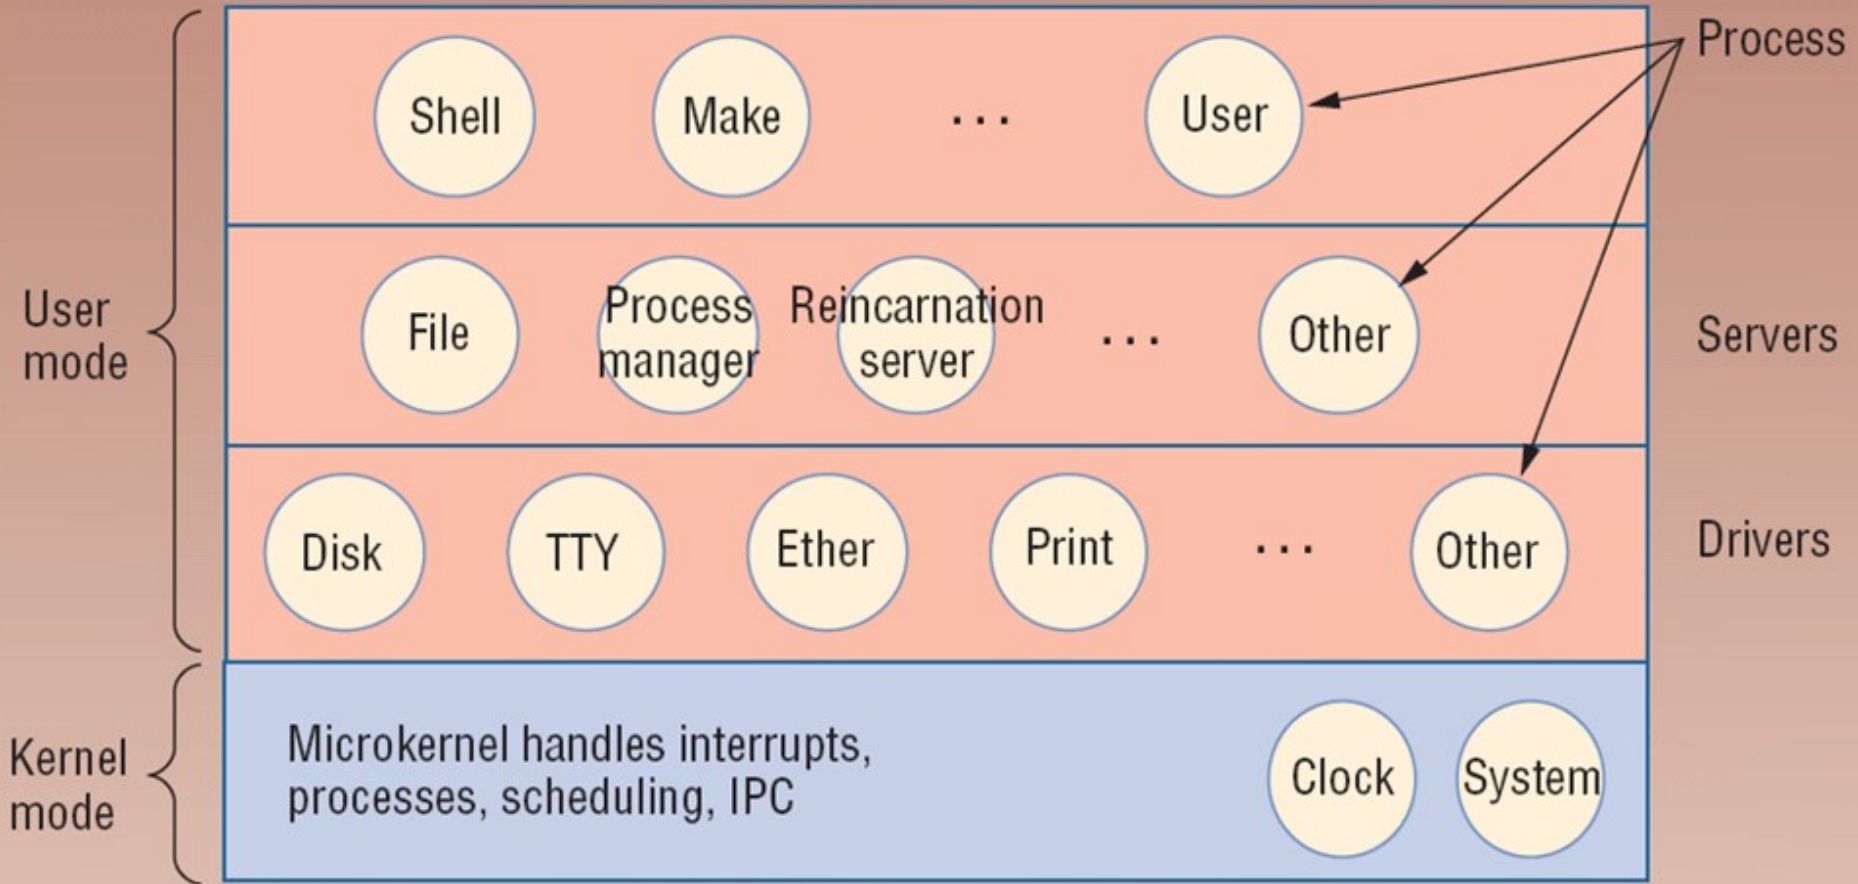
- The OS is divided into a number of layers (levels)
 - Each layer built on top of lower layers
 - The bottom layer (layer 0), is the hardware
 - The highest (layer N) is the user interface

- With modularity, layers are selected such that
 - Each uses functions (operations) and services of only lower-level layers





MINIX 3





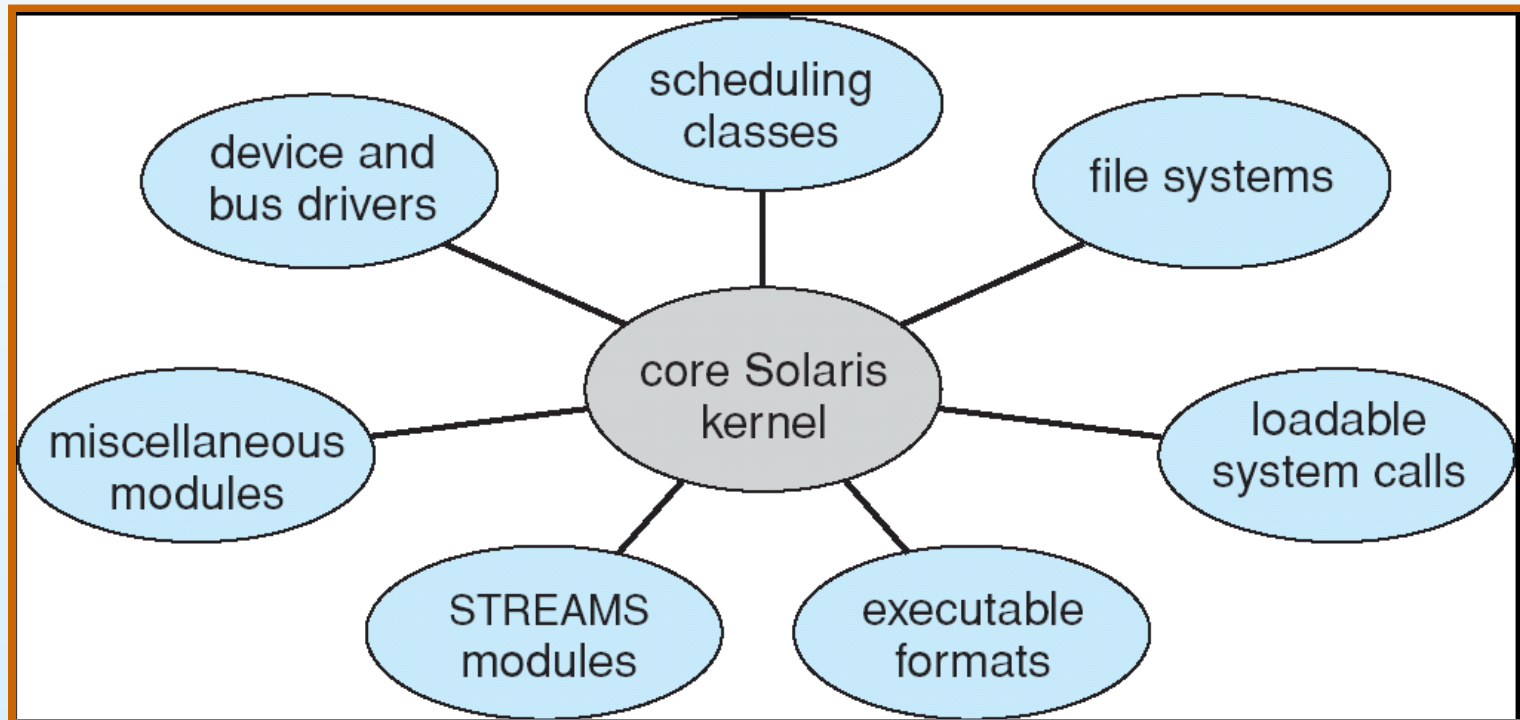
Microkernel System Structure

- Moves as much from the kernel into “*user*” space
- Communicate between user modules by message passing
- Benefits:
 - Easier to extend a microkernel
 - Easier to port the OS to new architectures
 - More reliable (less code is running in kernel mode)
 - More safe
- Detriments:
 - Performance overhead of user space to kernel space communication





Solaris Modular Approach





Modules

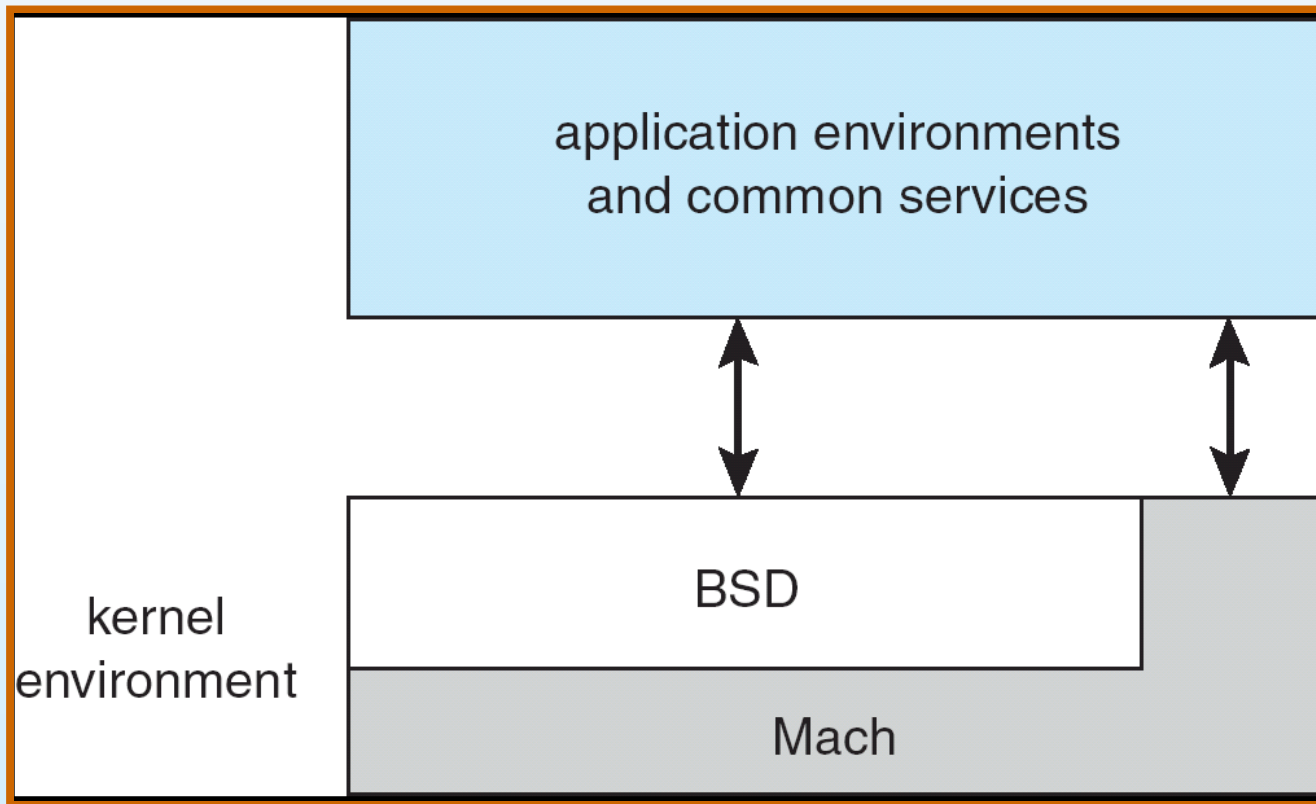
- **Most modern OS implement kernel modules**
 - **Uses object-oriented approach**
 - **Each core component is separate**
 - **Each talks to the others over known interfaces**
 - **Each is loadable as needed within the kernel**

- **Overall, similar to layers but with more flexible**



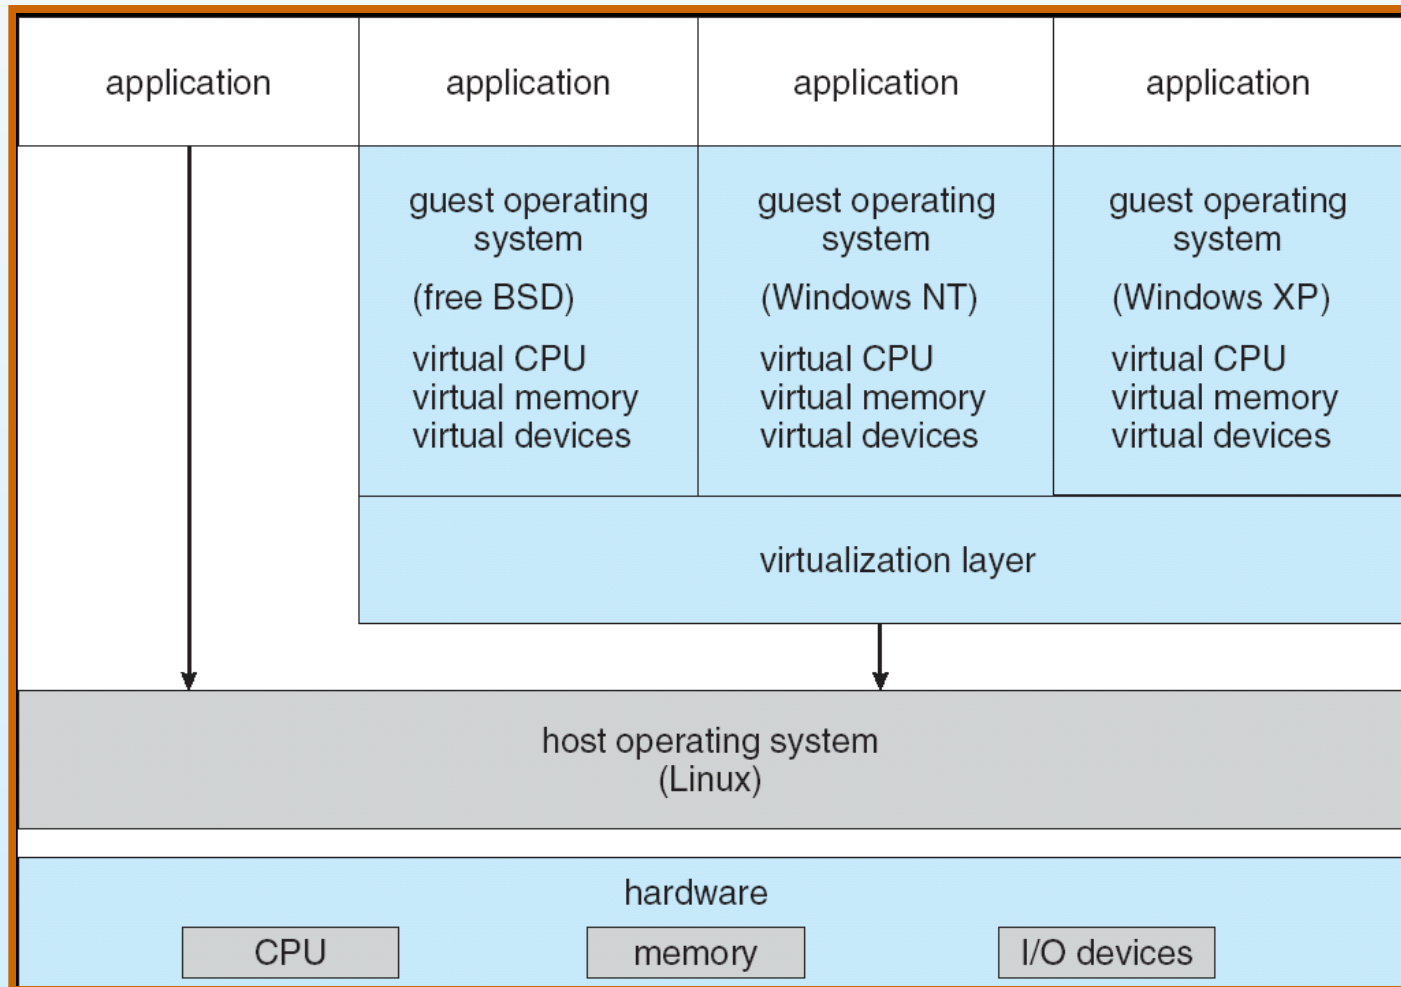


Mac OS X Structure



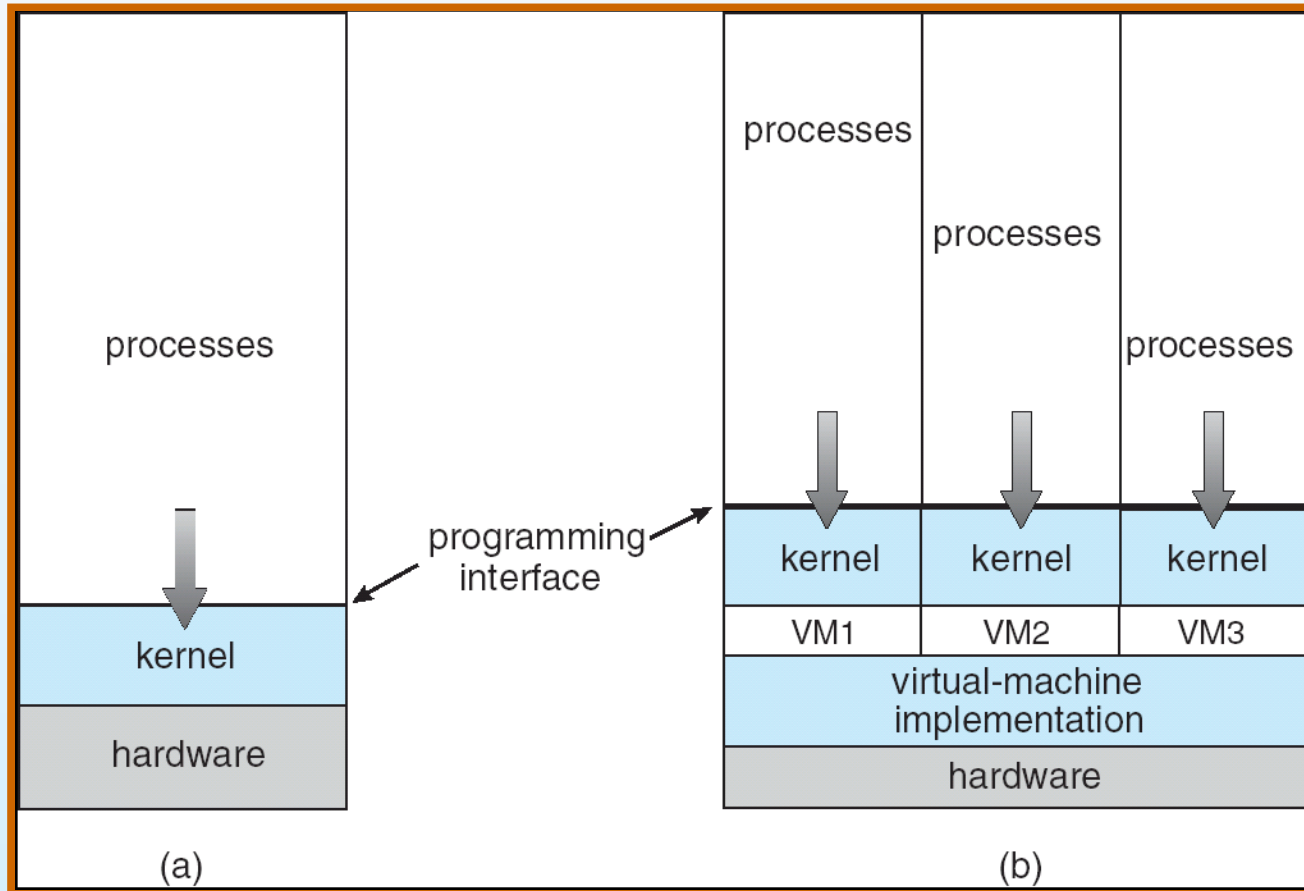


VMware Architecture





Virtual Machines



(a) Nonvirtual machine (b) virtual machine





Virtual Machines (Cont.)

- A ***virtual machine*** takes the layered approach to its logical conclusion.
- A VM provides an interface *identical* to the underlying bare hardware
- Each OS executes on its own processor with its own memory






Virtual Machines (Cont.)

- The VM provides complete protection of resources
 - Each virtual machine is isolated from all other virtual machines.
- A VM system is a perfect vehicle for OS research and development.
 - System development is done on the VM, instead of on a physical machine
 - Does not disrupt normal system operation.
- It is difficult to implement
 - Required to provide an *exact* duplicate to the underlying machine





The Java Virtual Machine





System Boot

■ *Booting*

- starting a computer by loading the kernel

■ *Bootstrap program*

- Code stored in ROM that locates the kernel, load it into memory, and start its execution

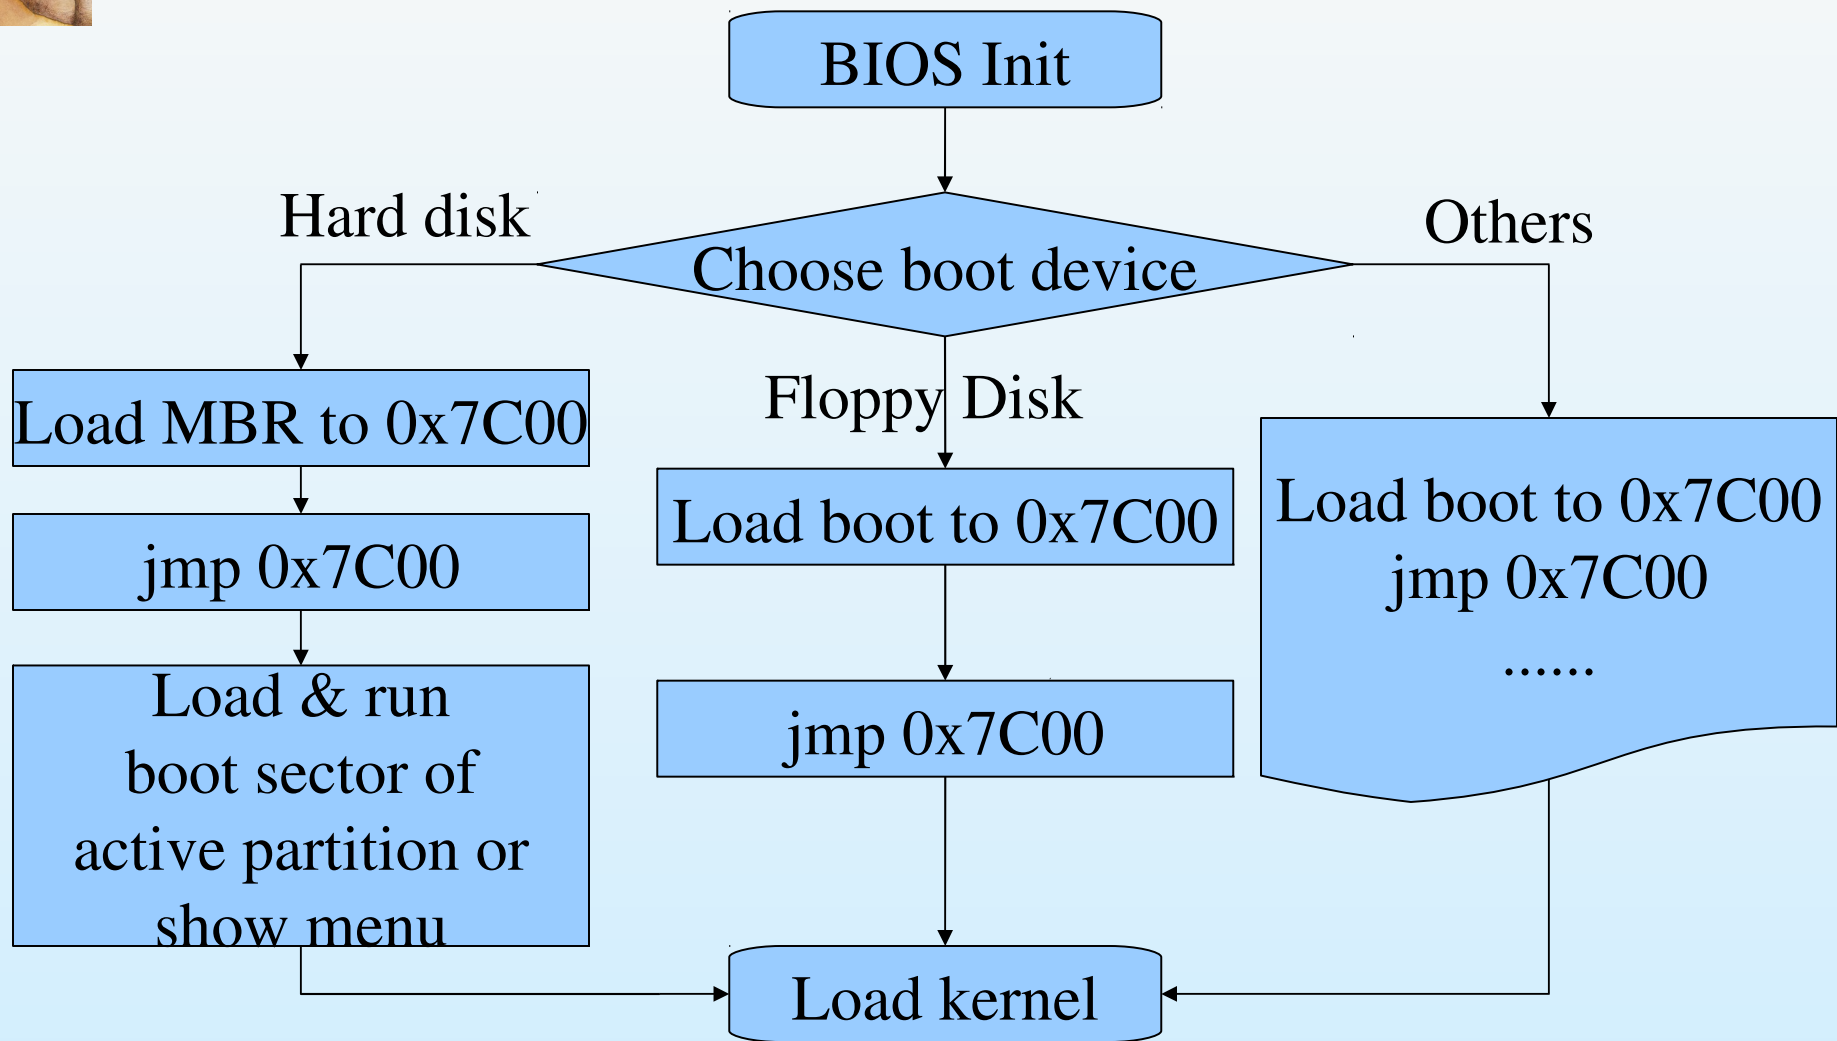
■ OS must be made available to hardware so hardware can start it

- Small piece of code – bootstrap loader, locates the kernel, loads it into memory, and starts it
- Sometimes two-step process where boot block at fixed location loads bootstrap loader
- Store entire OS in ROM and run it directly





X86 PC Boot





End of Chapter 2

