

Chapter 3: Processes





Chapter 3: Processes (进程)

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems






Process Concept

- **An OS executes a variety of programs:**
 - **Batch system – jobs (作业)**
 - **Time-shared systems – user programs or tasks**
- **Process**
 - **A program in execution**
- **A process includes:**
 - **program counter and other registers**
 - **text section**
 - **stack**
 - **data section**
 - **heap**





Process in Memory





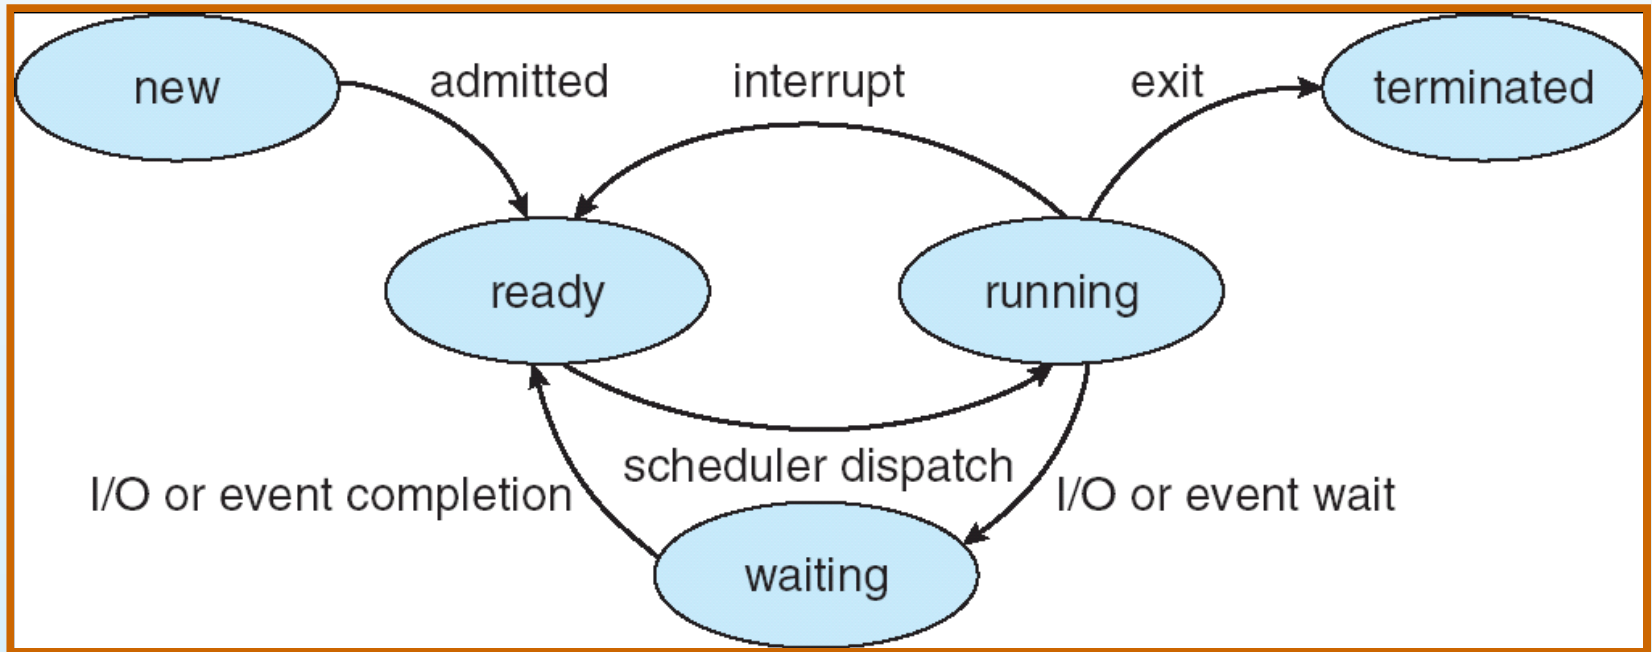
Process State

- As a process executes, it changes *state*
 - new: The process is being created
 - running: Instructions are being executed
 - waiting: The process is waiting for some event to occur
 - ready: The process is waiting to be assigned to a processor
 - terminated: The process has finished execution





Diagram of Process State





Process Control Block (PCB)

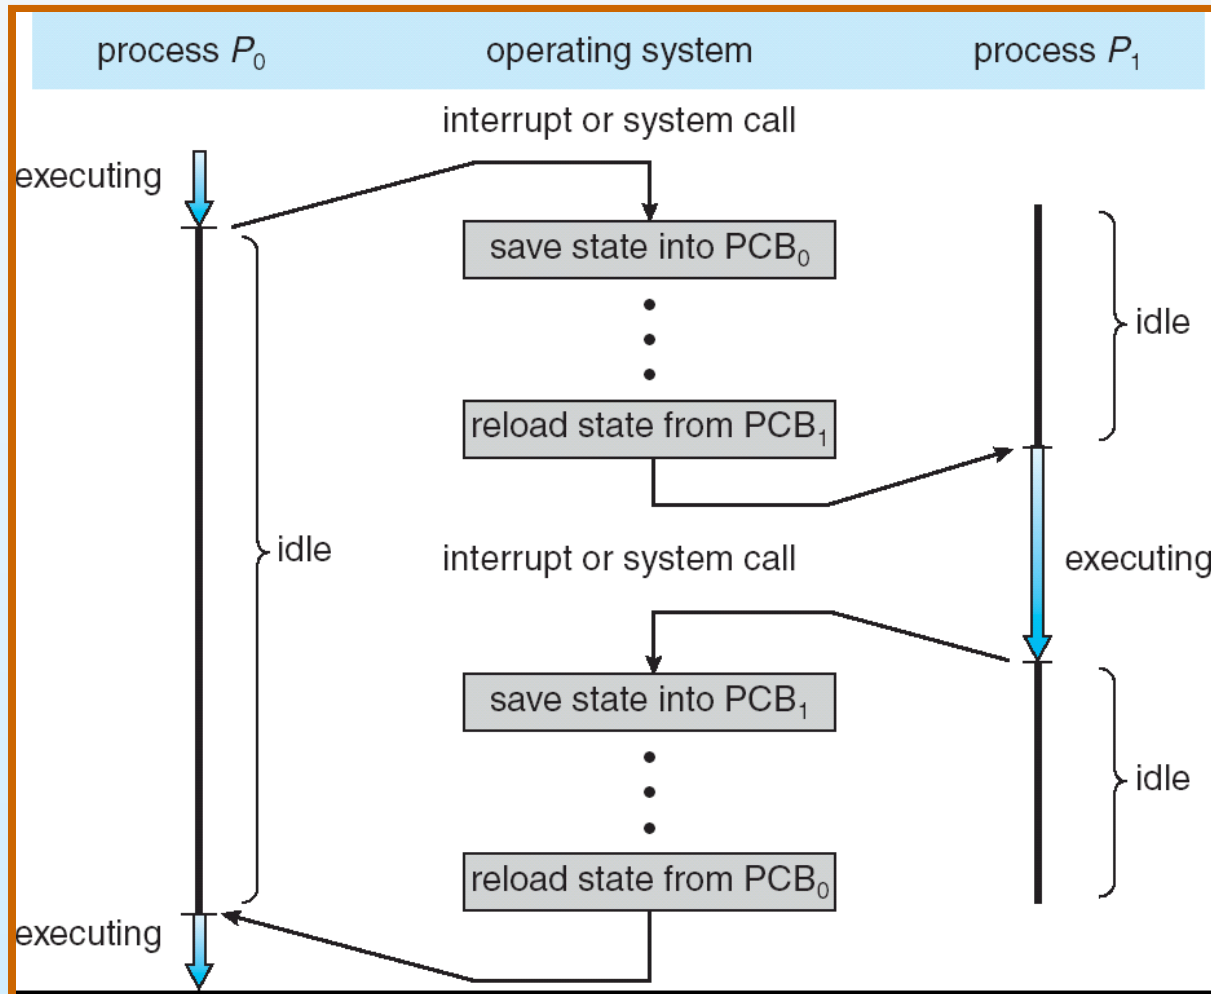
Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information





CPU Switch From Process to Process






Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues






Ready Queue And Various I/O Device Queues





Representation of Process Scheduling





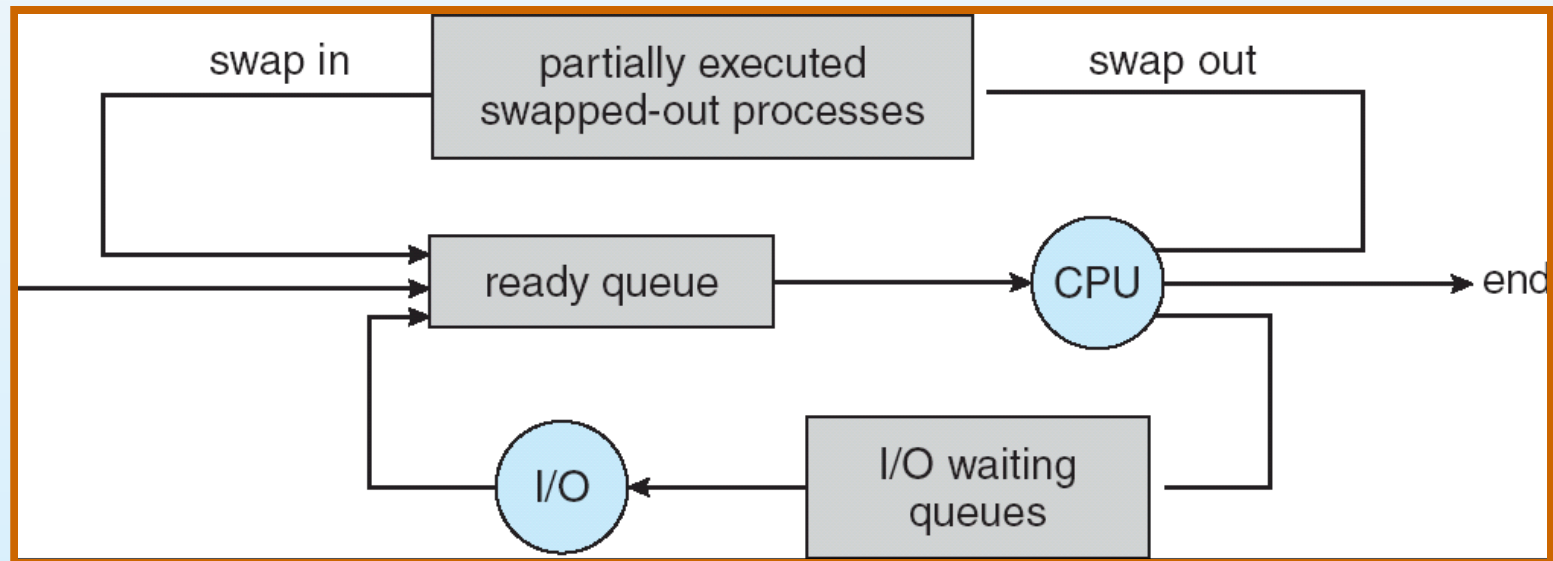
Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
- **Medium-term scheduler** (or swapping) – selects which process should be removed from memory





Addition of Medium Term Scheduling





Context Switch

- **When CPU switches to another process, the system must**
 - **Save the state of the old process**
 - **Load the saved state for the new process**
- **Context-switch time is overhead**
 - **The system does no useful work while switching**
- **Time dependent on hardware support**





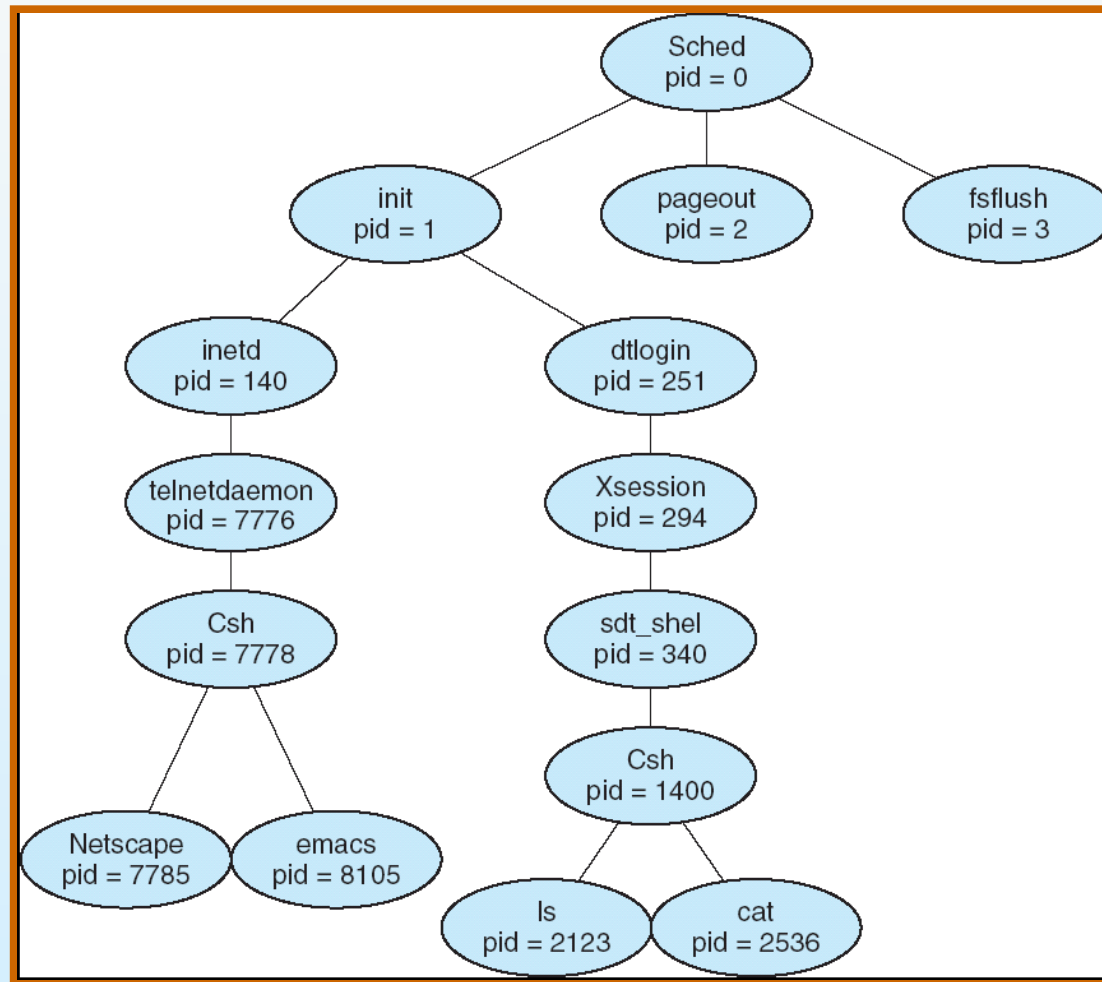
Process Creation

- **Parent process create children processes, which, in turn create other processes, forming a tree of processes**
- **Resource sharing**
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- **Execution**
 - Parent and children execute concurrently
 - Parent waits until children terminate





A tree of processes on a typical Solaris



Try “pstree -p” in Linux!





Process Creation (Cont.)

■ Address space

- Child duplicate of parent
- Child has a program loaded into it

■ UNIX examples

- fork system call creates new process
- exec system call used after a fork to replace the process' memory space with a new program





C Program Forking Separate Process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```





Process Creation (Cont.)

■ Windows examples

- **BOOL WINAPI CreateProcess(
LPCTSTR lpApplicationName,
LPTSTR lpCommandLine,
LPSECURITY_ATTRIBUTES lpProcessAttributes,
LPSECURITY_ATTRIBUTES lpThreadAttributes,
BOOL bInheritHandles,
DWORD dwCreationFlags,
LPVOID lpEnvironment,
LPCTSTR lpCurrentDirectory,
LPSTARTUPINFO lpStartupInfo,
LPPROCESS_INFORMATION lpProcessInformation
);**



C Program Creates Process in Win32

```
int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    // allocate memory
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    // create child process
    if (!CreateProcess(NULL, // use command line
        "C:\\WINDOWS\\system32\\mspaint.exe", // command line
        NULL, // don't inherit process handle
        NULL, // don't inherit thread handle
        FALSE, // disable handle inheritance
        0, // no creation flags
        NULL, // use parent's environment block
        NULL, // use parent's existing directory
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    // parent will wait for the child to complete
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    // close handles
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

- **Process executes last statement and asks the OS to delete it (exit)**
 - **Return value from child to parent (via wait)**
 - **Resources are deallocated by OS**
- **Parent may terminate execution of children (abort)**
 - **Kill any of the children as the parent likes**
 - **If parent is exiting**
 - ▶ **Some OS kill all children**
 - ▶ **Some OS give all children to *init***





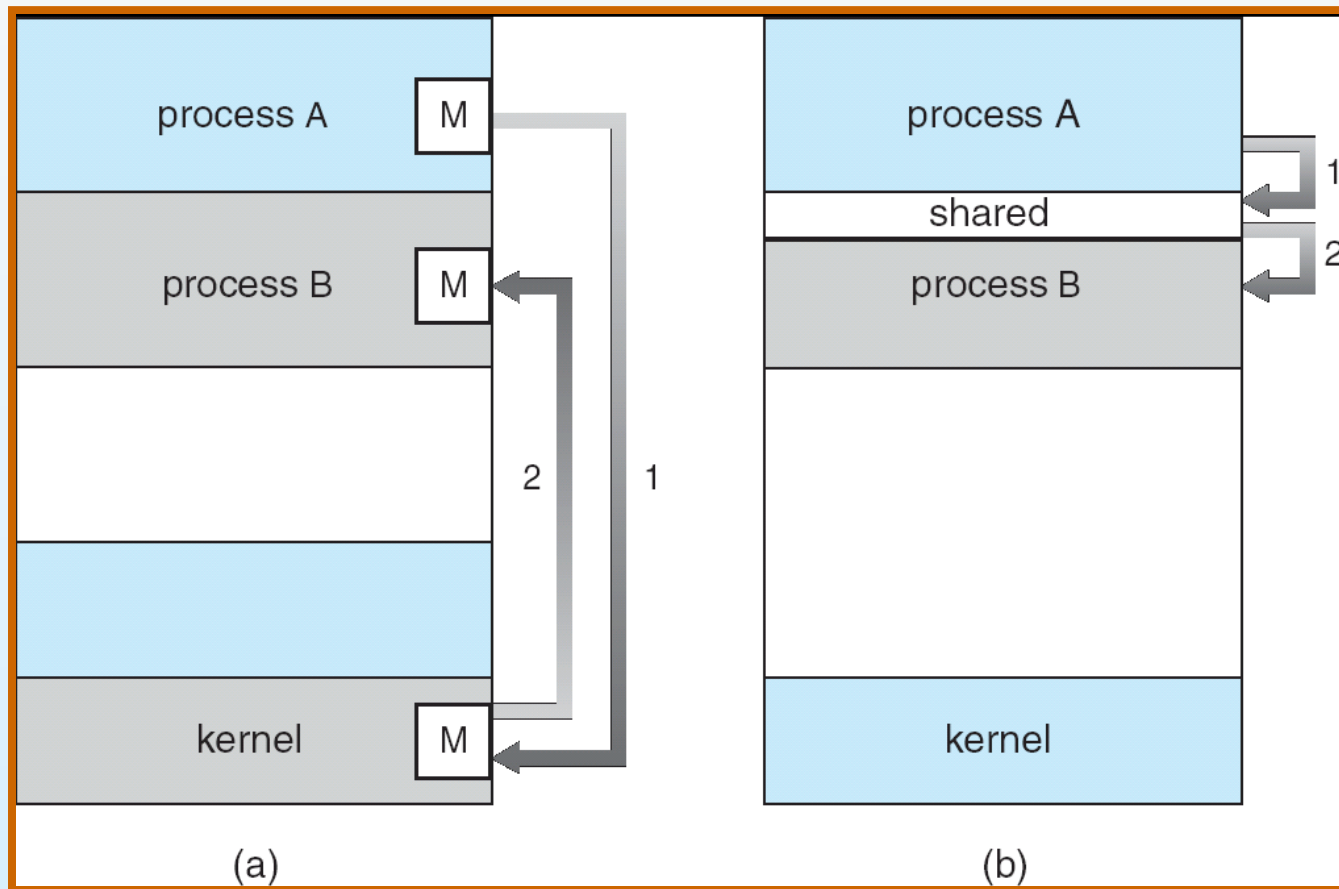
Cooperating Processes

- **Independent** process can NOT affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- **Advantages of process cooperation**
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience





Communications Models



(a) Message passing

(b) Shared memory





Shared-Memory Systems

- OS prevents one process from accessing another process's memory
- Shared memory requires that processes agree to remove this restriction
- Processes' responsibility
 - Synchronization
 - Don't write to the same location simultaneously
- Not all OSes support this feature
- Some OSes are born shared-memory
 - Most of embedded systems





Shared-Memory in POSIX

- **shmget()**
 - Allocate a shared-memory segment
- **shmat()**
 - Attach the shared-memory segment
- **shmdt()**
 - Detach the shared-memory segment
- **shmctl()**
 - Control the shared-memory segment





Message-Passing Systems

- Processes communicate with each other without resorting to shared variables
- Particularly useful in a distributed environment
- Provides two operations:
 - `send(message)`
 - ▶ Message size fixed or variable
 - `receive(message)`
- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive





Synchronization vs. Asynchronous

- Message passing may be either blocking or non-blocking
- Blocking is considered synchronous
 - The sender block until the message is received
 - The receiver block until a message is available
- Non-blocking is considered asynchronous
 - The sender send the message and continue
 - The receiver receive a valid message or null





Buffering

- Queue of messages attached to the link;
implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits





Client-Server Communication

- Sockets – leave to Network course
- Remote Procedure Calls
- Remote Method Invocation (Java)





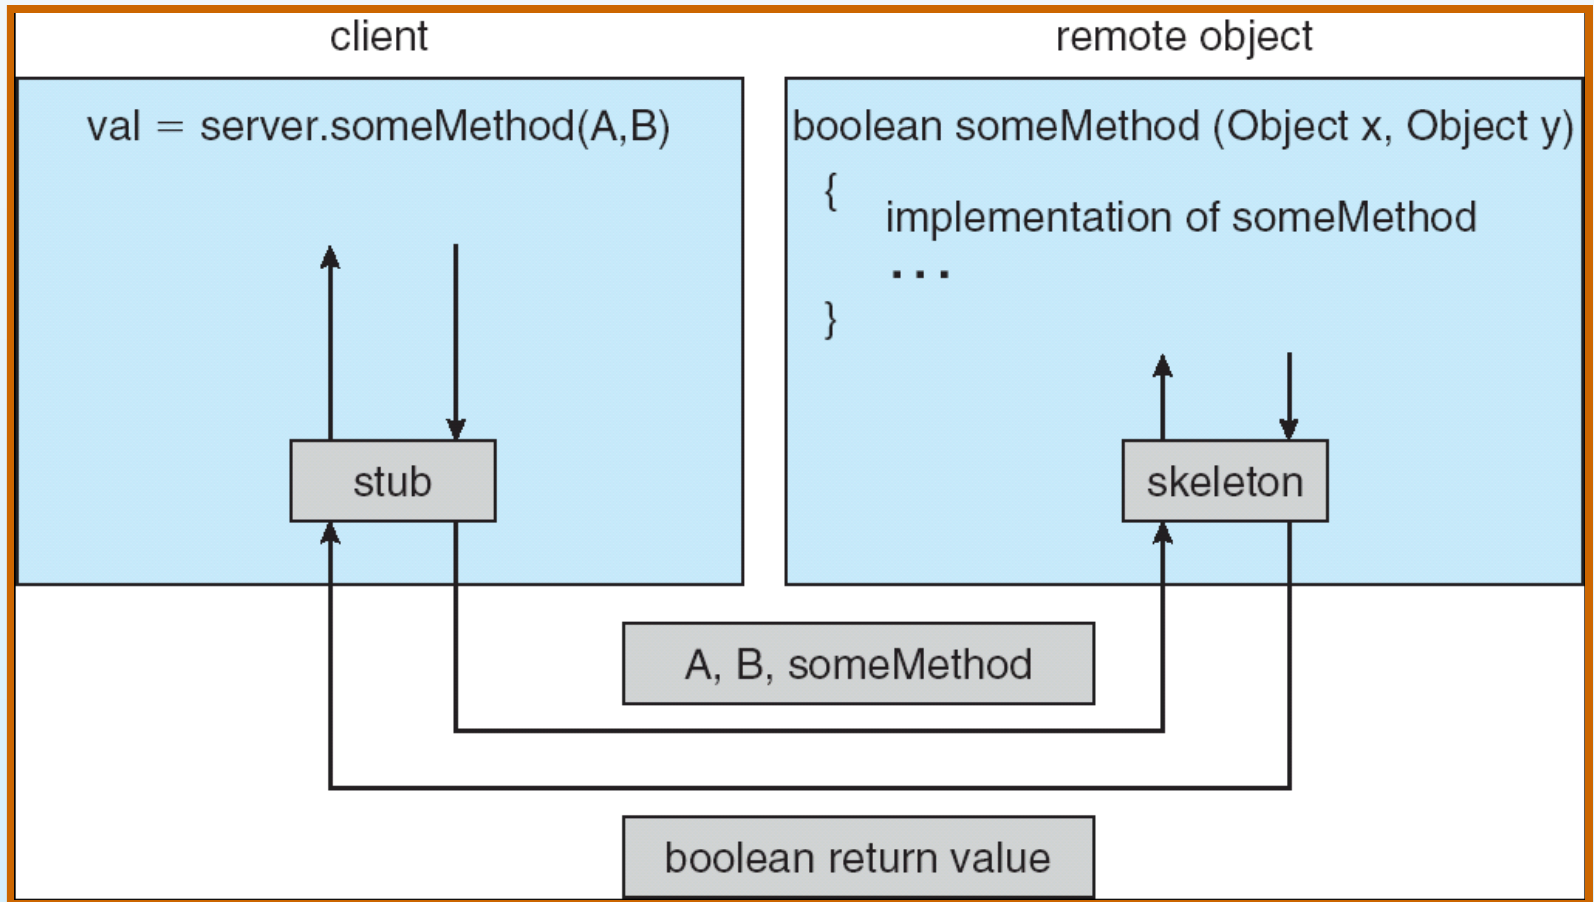
Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- The client-side stub
 - Locates the server
 - *Marshalls* the parameters.
- The server-side stub
 - Receives this message
 - Unpacks the marshalled parameters
 - Performs the procedure on the server.






Marshalling Parameters





Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.



End of Chapter 3

