

Chapter 5: CPU Scheduling





Chapter 5: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Operating Systems Examples





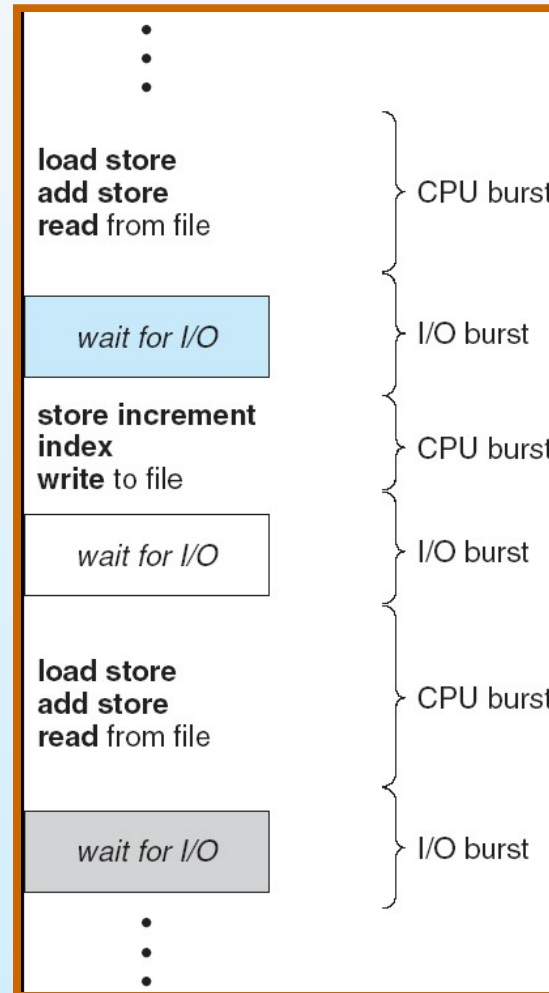
Basic Concepts

- **Maximum CPU utilization obtained with multiprogramming**
- **CPU–I/O Burst Cycle**
 - **Process execution consists of a *cycle* of CPU execution and I/O wait**
- **CPU burst distribution**





Alternating Sequence of CPU And I/O Bursts





CPU Scheduler (调度)

- Selects one of the processes in ready queue to run next
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is *nonpreemptive*
- All other scheduling is *preemptive*





Dispatcher

- Gives the CPU to the process selected by scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- *Dispatch latency*
 - time it takes for the dispatcher to stop one process and start another running





Scheduling Criteria

- **CPU utilization – to max**
 - keep the CPU as busy as possible
- **Throughput – to max**
 - # of processes that complete execution per time unit
- **Turnaround time – to min**
 - amount of time to execute a process
- **Waiting time – to min**
 - amount of time a process has been waiting in the ready queue
- **Response time – to min**
 - amount of time from a request was submitted to the first response is produced






How to Schedule?

Process **Arrival time**

Burst time


A

0



B

0



C

0




D

0



E

0






How to Schedule?

Process **Arrival time**

Burst time


A

0




B

1



C

2




D

3



E

4






How to Schedule?

Process **Arrival time**

Burst time


A

0




B

0




C

0




D

0



E

0





First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

P_1	24
-------	----

P_2	3
-------	---

P_3	3
-------	---

- Arrive in the order: P_1 , P_2 , P_3

The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

■ The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case





Shortest-Job-First (SJF) Scheduling

- Schedule the process who has the shortest next CPU burst
- Two schemes:
 - Nonpreemptive
 - ▶ Process cannot be preempted until completes its CPU burst
 - Preemptive
 - ▶ if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
 - ▶ Shortest-Remaining-Time-First (SRTF)
- SJF is optimal
 - minimum average waiting time






Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

■ SJF (non-preemptive)



■ Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$






Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

■ SJF (preemptive)



■ Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$





How to Know Length of Next CPU Burst?

- Can only estimate the length
 - using the length of previous CPU bursts





Priority Scheduling


- A priority number is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv Starvation
 - low priority processes may never execute
- Solution \equiv Aging
 - as time progresses increase the priority of the process





Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*)
 - usually 10-100 milliseconds.
- After this time has elapsed, the process is preempted and added to the end of the ready queue.






Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

■ The Gantt chart is:



■ Typically, higher average turnaround than SJF, but better *response*






Time Quantum and Context Switch Time

■ Performance

- *quantum* large \Rightarrow FCFS
- *quantum* small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high





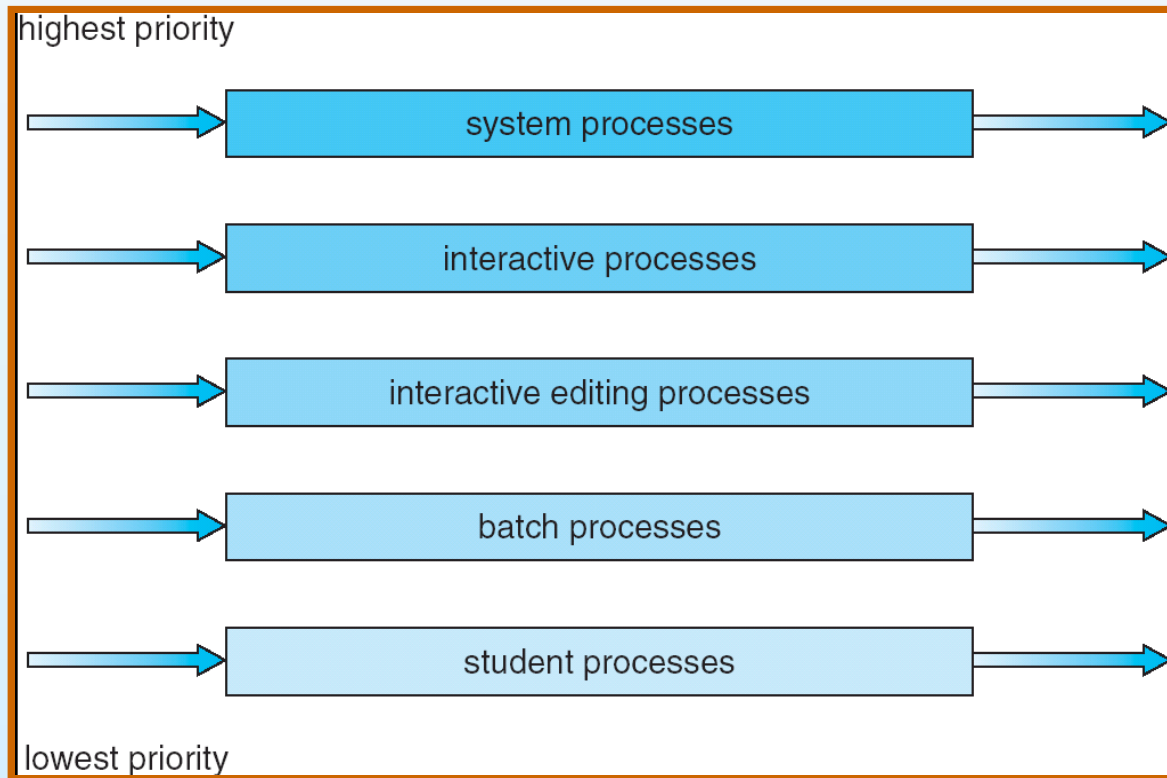
Multilevel Queue

- Ready queue is partitioned into separate queues
 - i.e., foreground (interactive), background (batch)
- Each queue has its own scheduling algorithm
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues
 - Fixed priority scheduling
 - ▶ (i.e., serve all from foreground then from background).
Possibility of starvation.
 - Time slice
 - ▶ each queue gets a certain amount of CPU time





Multilevel Queue Scheduling





Multilevel Feedback Queue

- A process can move between the queues
 - aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine
 - ▶ when to upgrade a process
 - ▶ when to demote a process
 - ▶ which queue a process will enter when that process needs I/O





Example of Multilevel Feedback Queue

■ Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS


■ Scheduling

- A new job enters queue Q_0 , receives 8 ms.
- If it does not finish in 8 ms, job is moved to queue Q_1 .
- At Q_1 job receives 16 additional milliseconds.
- If it still does not complete, it is preempted and moved to queue Q_2 .





Multilevel Feedback Queues





Multiple-Processor Scheduling

- Multiple CPUs give us
 - *Load sharing*
 - More complex CPU scheduling
- We discuss *homogeneous processors* only





Approaches to Multiple-Processor Scheduling

■ *Asymmetric multiprocessing*

- Only one processor runs OS
- Other processors execute only user code
- Reduce the need for data sharing

■ *Symmetric multiprocessing (SMP)*

- All processes are equal to run processes
 - ▶ Every processor has a private ready queue (common way)
 - ▶ All processors share one ready queue
- The scheduler must be programmed carefully
- Supported by Windows XP/2000, Solaris, Linux and Mac OS X





Processor Affinity (亲和度)

- Avoid migrating a process from CPU 0 to CPU 1
 - Invalidate cache of CPU 0
 - Re-populate cache of CPU 1
- *Processor affinity*
 - A process has an affinity for the processor on which it is running
- *Soft affinity*
 - Attempting to keep affinity but not guaranteeing
- *Hard affinity*
 - Specifying a process which will never migrate
 - ▶ `sched_setaffinity()` in Linux,
 - ▶ `SetProcessAffinityMask()` in Windows 2000/XP/Vista





Load Balancing

- Attempting to keep the workload distributed across all processors
- *Push migration*
 - A specific task periodically pushes process from overloaded to idle or less-busy processors
- *Pull migration*
 - An idle processor pulls a waiting task from a busy processor
- Some OSes implement both
 - i.e., Linux and FreeBSD
- Load balancing vs. processor affinity





Operating System Examples

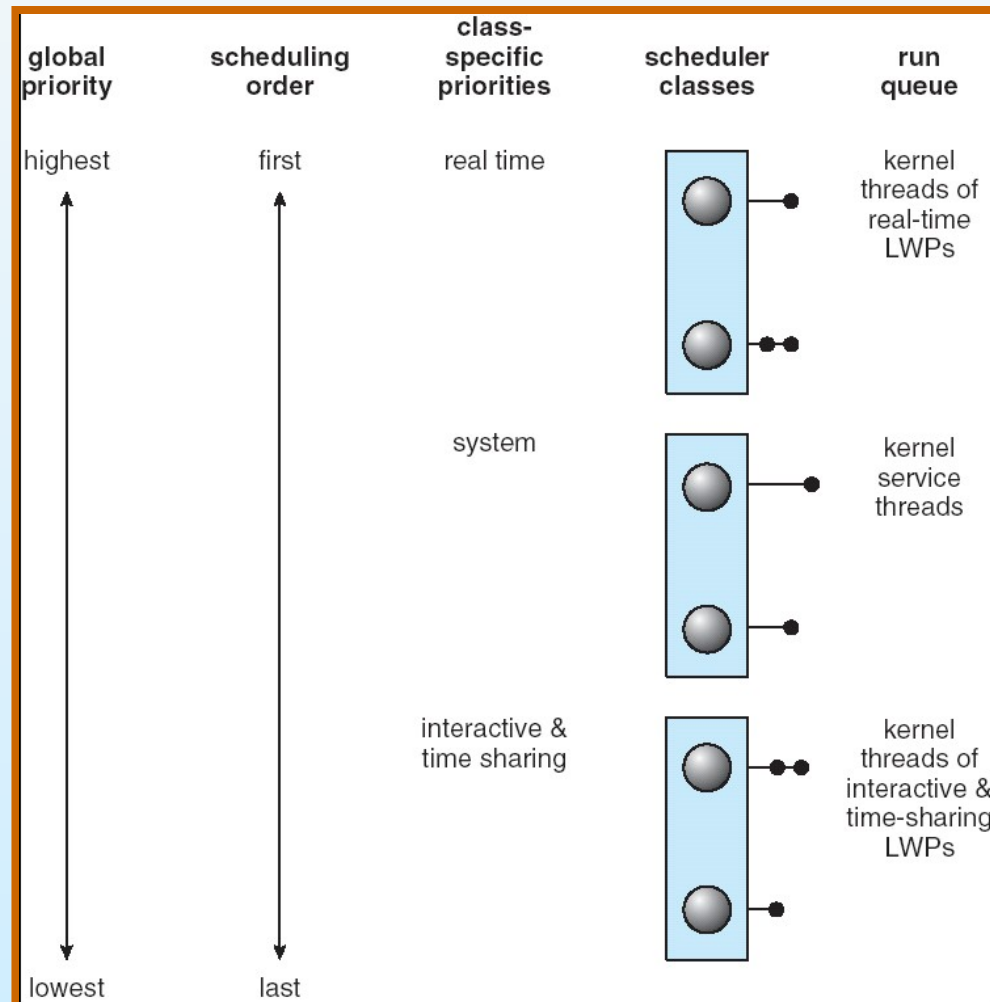
- **Solaris scheduling**
- **Windows XP scheduling**
- **Linux scheduling**





Solaris Scheduling

■ Priority-based thread scheduling





Solaris Dispatch Table

priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59





Windows XP Priorities

- **Schedule threads using a priority-based, preemptive scheduling algorithm**
- **A thread will run until**
 - Preempted by a higher-priority thread
 - Terminates
 - Time quantum ends
 - Calls a blocking system call
- **Priorities**
 - Variable class: 1-15
 - Real-time class: 16-31
 - Memory management: 0
 - Idle thread: no priority





Windows XP Priorities

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

- For threads in variable class, the priorities
 - Lowered after time quantum's running out
 - Never lowered below base priority (normal) of the class
 - Boosted when “waiting -> ready”
 - ▶ Interactive threads (active window) get more boost
- Gives foreground processes 3x quantum





Linux Scheduling

- A preemptive, priority-based algorithm
- Priorities
 - Real-time: 0-99
 - Nice: 100-140

numeric priority	relative priority		time quantum
0	highest	real-time tasks	200 ms
•			
•			
•			
99		other tasks	
100			
•			
•			
•			
140			
	lowest		10 ms





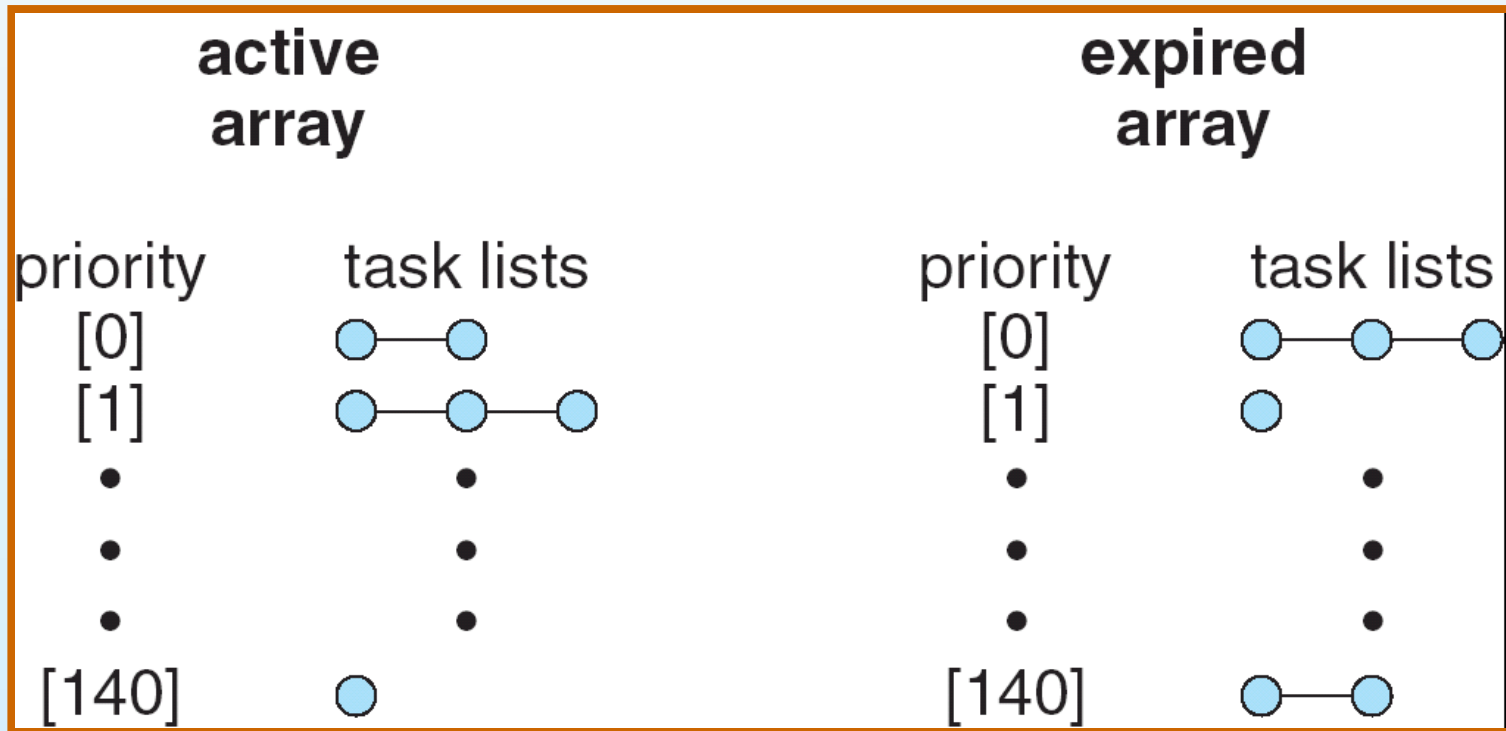
Linux Scheduling (cont.)

- Each processor has one *runqueue* which contains all runnable tasks distributed to it
- Each runqueue contains two arrays
 - *Active* array contains all tasks with time remaining in their time slices
 - *Expired* array contains all expired tasks
- Task with the highest priority in the active array is chosen to run
- When active array is empty, the expired array becomes active array
 - And adjust priorities





List of Tasks Indexed According to Priorities





Adjust Priorities

- A task's priority can be nice, nice – 5 or nice + 5
- Determined by how long it has been waiting
 - Interactive tasks, nice – 5
 - CPU-bound tasks, nice + 5
- Priorities are updated when moving tasks from active array to expired array



End of Chapter 5

