

Chapter 7: Deadlocks





Chapter 7: Deadlocks

- The Deadlock Problem
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock





The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Resource types
 - CPU cycles, memory space, I/O devices*
- Each process utilizes a resource as follows:
 - Request
 - ▶ `open()`, `malloc()`, `wait()`
 - Use
 - Release
 - ▶ `close()`, `free()`, `signal()`





Deadlock Characterization


- Deadlock can arise if four conditions hold simultaneously.
 - *Mutual exclusion*: only one process at a time can use a resource.
 - *Hold and wait*: a process holding at least one resource is waiting to acquire additional resources held by other processes.
 - *No preemption*: a resource can be released only voluntarily by the process holding it.
 - *Circular wait*: there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .





Resource-Allocation Graph


- Process




- Resource Type with 4 instances



- P_i requests instance of R_j




- P_i is holding an instance of R_j






Example of a Resource Allocation Graph






Resource Allocation Graph With A Deadlock





Graph With A Cycle But No Deadlock





Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.





Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
- Allow the system to enter a deadlock state and then recover.
- Ignore the problem and pretend that deadlocks never occur in the system
 - Used by most operating systems, including UNIX and Windows.
 - It is up to the application developer to handle deadlock





Deadlock Prevention

■ Attack Mutual Exclusion

- not required for sharable resources
- must hold for nonsharable resources.

■ Attack Hold and Wait

- Get all or none
 - ▶ Require process to request and get all its resources before it begins execution
 - ▶ Allow process to request resources only when the process has none.
- Low resource utilization; starvation possible.





Deadlock Prevention (Cont.)

■ Attack No Preemption

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

■ Attack Circular Wait

- Impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.





Deadlock Avoidance

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.





Safe State

- A process requests a resource, system must decide if the allocation leaves the system in a safe state.
- **Safe state**
 - A sequence $\langle P_1, P_2, \dots, P_n \rangle$ of processes
 - The resources that P_i can still request equal available resources + resources held by all the P_j with $j < i$.
- That is:
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, release resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on.





Safe or Unsafe

Has Max			Has Max			Has Max			Has Max			Has Max		
A	3	9	A	3	9	A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	—	B	0	—	B	0	—
C	2	7	C	2	7	C	2	7	C	7	7	C	0	—
Free: 3			Free: 1			Free: 5			Free: 0			Free: 7		
(a)			(b)			(c)			(d)			(e)		





Safe or Unsafe?

Has Max		
A	3	9
B	2	4
C	2	7

Free: 3

(a)

Has Max		
A	4	9
B	2	4
C	2	7

Free: 2

(b)

Has Max		
A	4	9
B	4	4
C	2	7

Free: 0

(c)

Has Max		
A	4	9
B	—	—
C	2	7

Free: 4


(d)





Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.






Avoidance algorithms

- **Single instance of a resource type. Use a resource-allocation graph**
- **Multiple instances of a resource type. Use the banker's algorithm**






Resource-Allocation Graph





Unsafe State In Resource-Allocation Graph





Resource-Allocation Graph Algorithm

- Suppose that process P_i requests a resource R_j
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph





Banker's Algorithm

- Multiple instances.
- Each process must a priori claim maximum use.
 - Is this possible?
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.





Example of Banker's Algorithm

■ Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	7 5 3	3 3 2	7 4 3
P_1	2 0 0	3 2 2		1 2 2
P_2	3 0 2	9 0 2		6 0 0
P_3	2 1 1	2 2 2		0 1 1
P_4	0 0 2	4 3 3		4 3 1

The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.





Example: P_1 Request (1,0,2)

- Check that Request \leq Available.

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	4	3	2	3	0
P_1	3	0	2	0	2	0			
P_2	3	0	1	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

- Sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ is safety.
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?





Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme





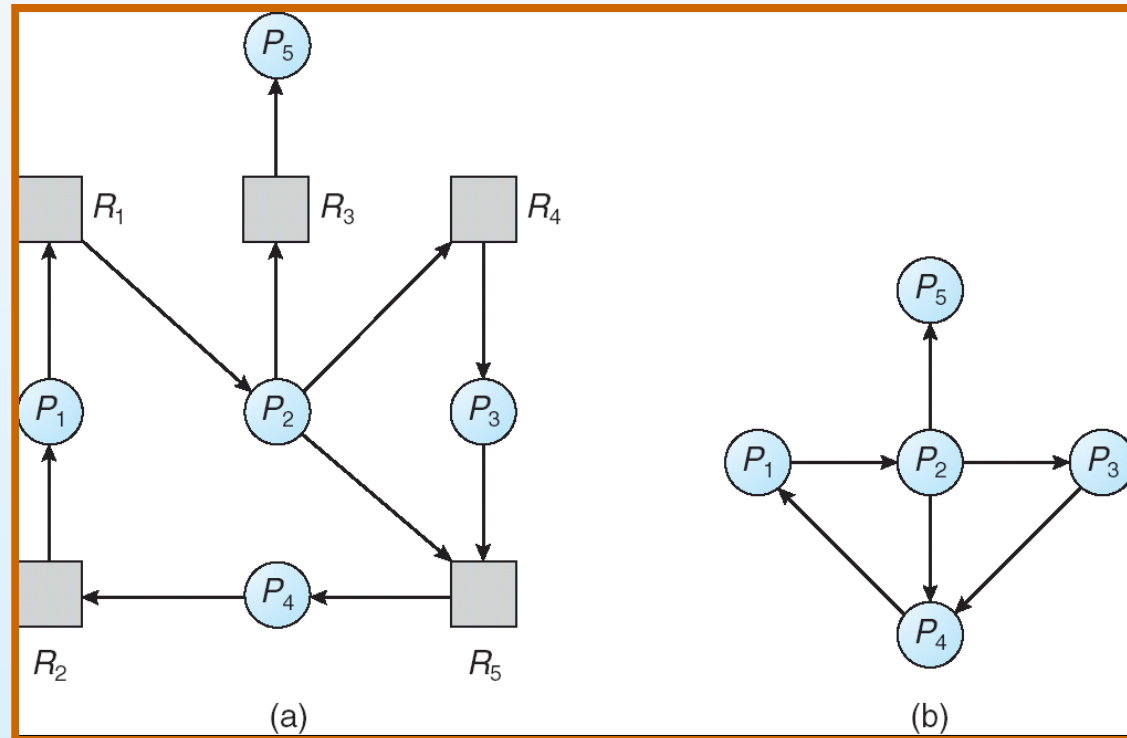
Single Instance of Each Resource Type

- Maintain *wait-for* graph
 - Nodes are processes.
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j .
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of n^2 operations.





Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph





Example of Detection Algorithm

■ Snapshot at time T_0 :

Allocation Request Available

	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

■ Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$.





Example (Cont.)

- P_2 requests an additional instance of type C.

Request

A B C

P_2 0 0 1

- State of system?
 - Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes; requests.
 - Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4 .





Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - How many processes will need to be terminated.
 - Is process interactive or batch?





Recovery from Deadlock: Resource Preemption

■ Selecting a victim

- Minimize cost.

■ Rollback

- Return to some safe state
- Restart process for that state.

■ Starvation

- Same process may always be picked as victim
- So include number of rollback in cost factor.



End of Chapter 7

