

Chapter 8: Main Memory





Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Segmentation
- Example: The Intel Pentium





Objectives

- Know how to organize memory
- Know memory-management techniques
 - Paging
 - Segmentation
- Know the Intel Pentium, which supports
 - Pure segmentation
 - Pure paging
 - Segmentation with paging





Background

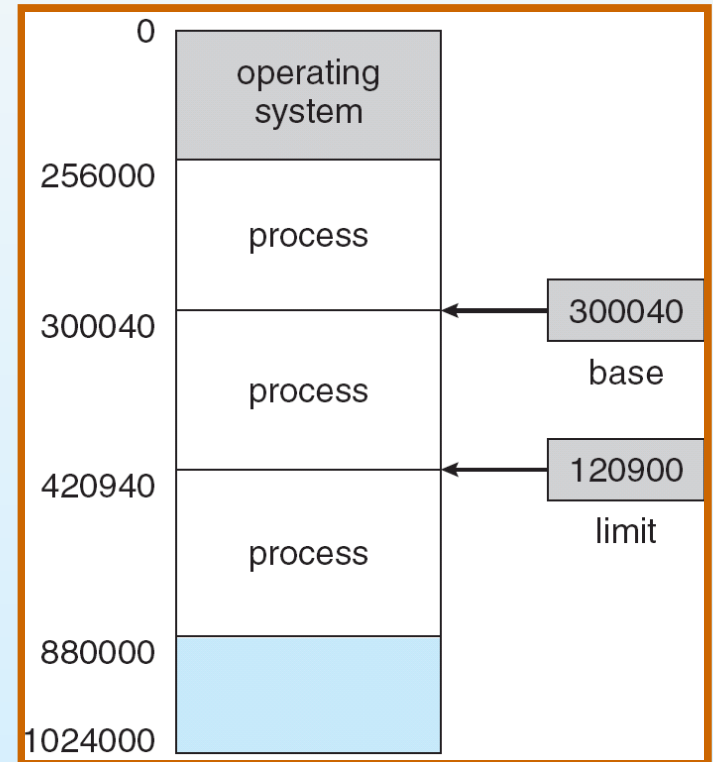
- Before running, a program must be
 - Brought (from disk) into memory
 - Placed within a process
- The only storage that CPU can access directly
 - Main memory
 - Registers
- Access time
 - Register: one CPU clock (or less)
 - Main memory: many cycles
 - So, cache comes!
- The way to ensure correct operation
 - Protection of memory (by hardware)





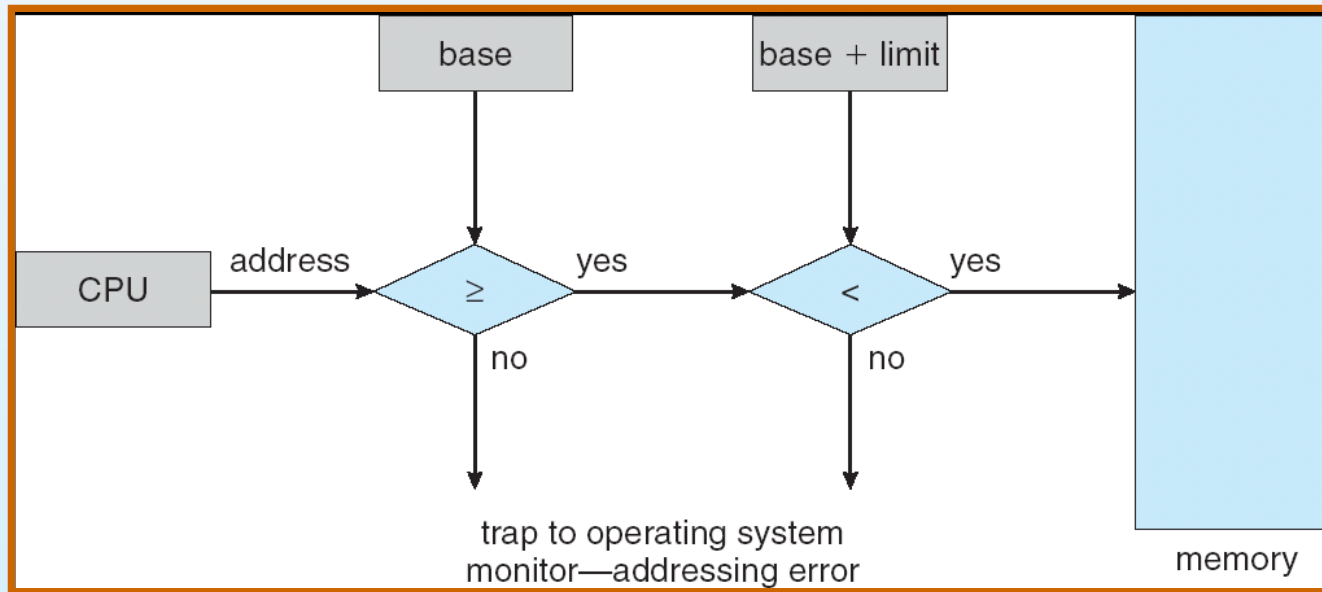
Base and Limit Registers

- Logical address space defined
 - By **base** and **limit** registers
- Only OS can set base and limit
 - Privileged Instructions
- Verify *every* address generated in user mode
- If bad address found , trap into kernel
- Can kernel access user code and data?





HW address protection with base and limit registers





Binding of Instructions and Data to Memory

- **Address binding happens at**
 - **Compile time**
 - ▶ If memory location known, generate absolute code
 - ▶ Must recompile code if starting location changes
 - **Load time**
 - ▶ At compile time, memory location is unknown
 - ▶ Generate relocatable code
 - **Execution time**
 - ▶ Process can be moved from one memory segment to another.
 - ▶ Need hardware support (e.g., base and limit registers)





Logical vs. Physical Address Space

■ Logical address

- Generated by the CPU
- Also referred to as virtual address

■ Physical address

- Address seen by the memory unit

■ They are the same in

- Compile-time address-binding
- Load-time address-binding

■ They are different in

- Execution-time address-binding
- Translated by MMU






Memory-Management Unit (MMU)

- Hardware device that maps logical to physical address
- In MMU scheme
 - Physical address
= relocation register + logical address
 - MS-DOS works in this way
- The user program
 - Deals with *logical* addresses
 - Never sees the *real* physical addresses





Dynamic relocation using a relocation register





Dynamic Loading

- Routine is not loaded until it is called
- Better memory-space utilization
- Unused routine is never loaded
 - Useful when large amounts of code are needed to handle infrequently occurring cases
- Implemented through program design
- No special support from the OS





Dynamic Linking and Shared Libraries

- **Linking postponed until execution time**
 - Useful for system libraries, language libraries
 - All processes execute only one copy of the routine
 - Easy to update
 - ▶ Version must be controlled
- **Stub**
 - Used to locate the appropriate library routine
 - Replaces itself with the address of the routine
 - Executes the routine
- **OS maps the routine into processes' memory address**





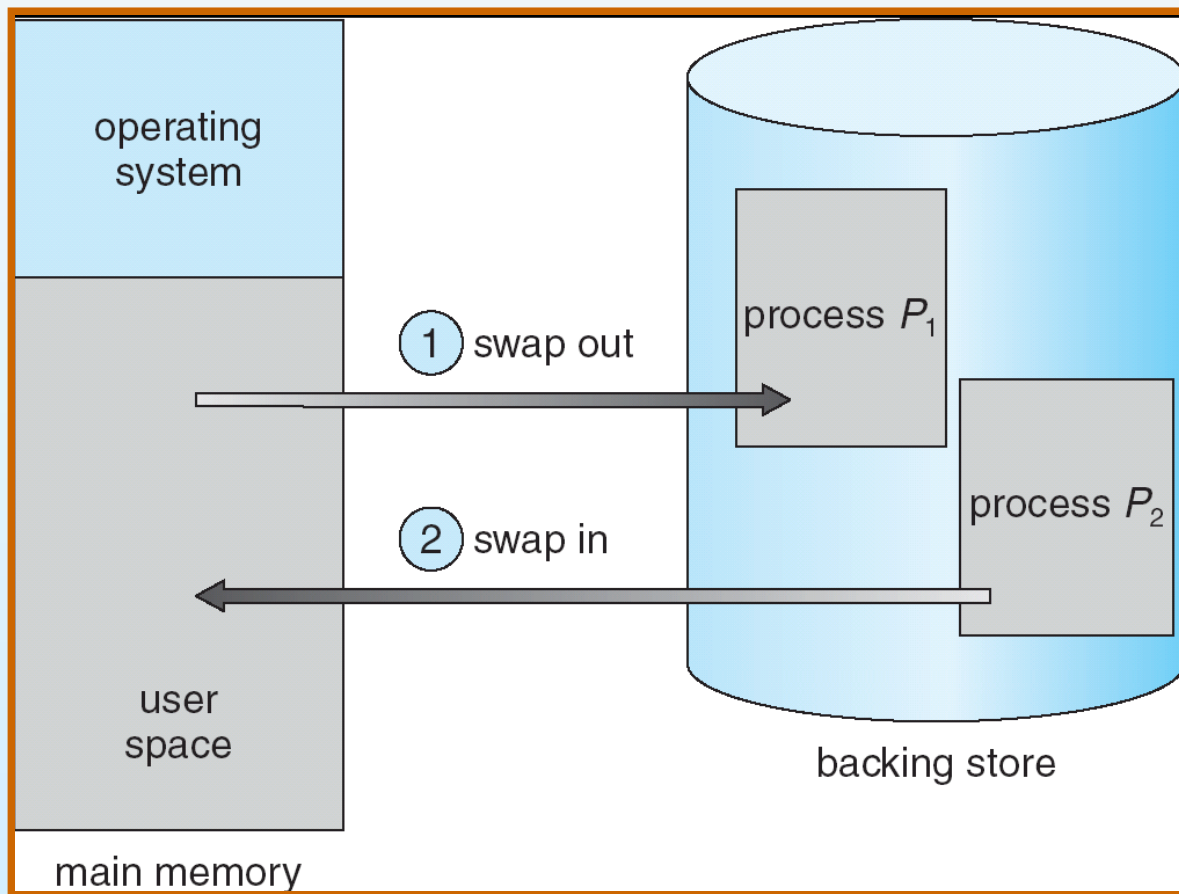
Swapping

- A process can be
 - Swapped temporarily out to a backing store
 - Brought back for continued execution
- Logical space relocation is needed
- Swap out, swap in
 - Use priority-based scheduling algorithms
- Overhead
 - Transfer time is fairly high
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)





Schematic View of Swapping





Contiguous Allocation

- Main memory includes two partitions:
 - OS kernel
 - ▶ Usually in low memory
 - User processes
 - ▶ Usually in high memory
- Relocation registers used to do protection
 - Base register contains the smallest physical address
 - Limit register contains range of logical addresses
 - Logical address must be less than the limit
 - MMU maps logical address *dynamically*

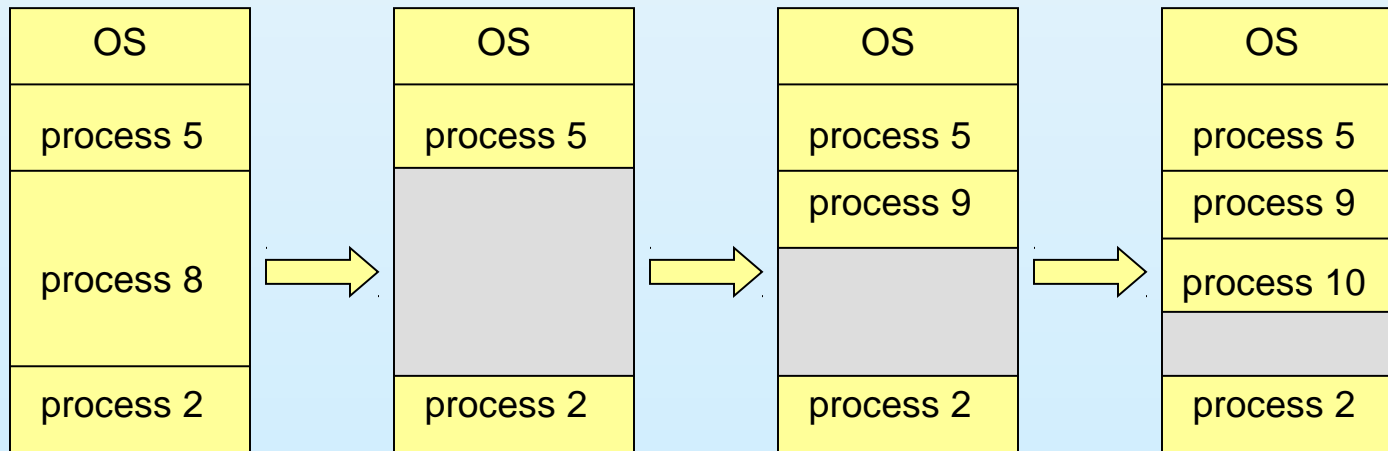




Contiguous Allocation (Cont.)

■ Multiple-partition allocation

- Hole – block of available memory
- A process arrives
 - ▶ Allocated memory from a hole large enough
- OS maintains information about:
 - a) allocated partitions b) free partitions (hole)





Dynamic Storage-Allocation Problem

■ First-fit

- Allocate the *first* hole that is big enough

■ Next-fit

- First-fit begins from previous allocation point

■ Best-fit

- Allocate the *smallest* hole that is big enough
- Produces the smallest leftover hole
- Must search entire list, unless ordered by size

■ Worst-fit

- Allocate the *largest* hole
- Produces the largest leftover hole
- Must also search entire list





Fragmentation

■ External Fragmentation

- Total free space satisfies a request
- But it is not contiguous

■ Internal Fragmentation

- Allocated memory based on block size
- Some waste space in the last block

■ Reduce external fragmentation by compaction

- Place all free memory together
- *Only* if relocation is dynamic
- I/O problem (DMA)
 - ▶ Don't move I/Oing processes
 - ▶ Do I/O only into OS buffers





Paging


- Logical address space can be noncontiguous
- Divide physical memory into fixed-sized frames
 - Size is power of 2, between 512 bytes and 16 MB
- Divide logical memory into fixed-sized pages
 - The same as frames
- Keep track of all free frames
 - Run a program of n pages, need to find n free frames
- Page table
 - Used to translate logical to physical addresses
- Internal fragmentation





Address Translation Scheme

- Logical address is divided into:
 - Page number (p)
 - ▶ An index into a *page table*
 - ▶ *Page table* contains physical address of each page
 - Page offset (d)
 - ▶ Combined with base address to define the physical memory address

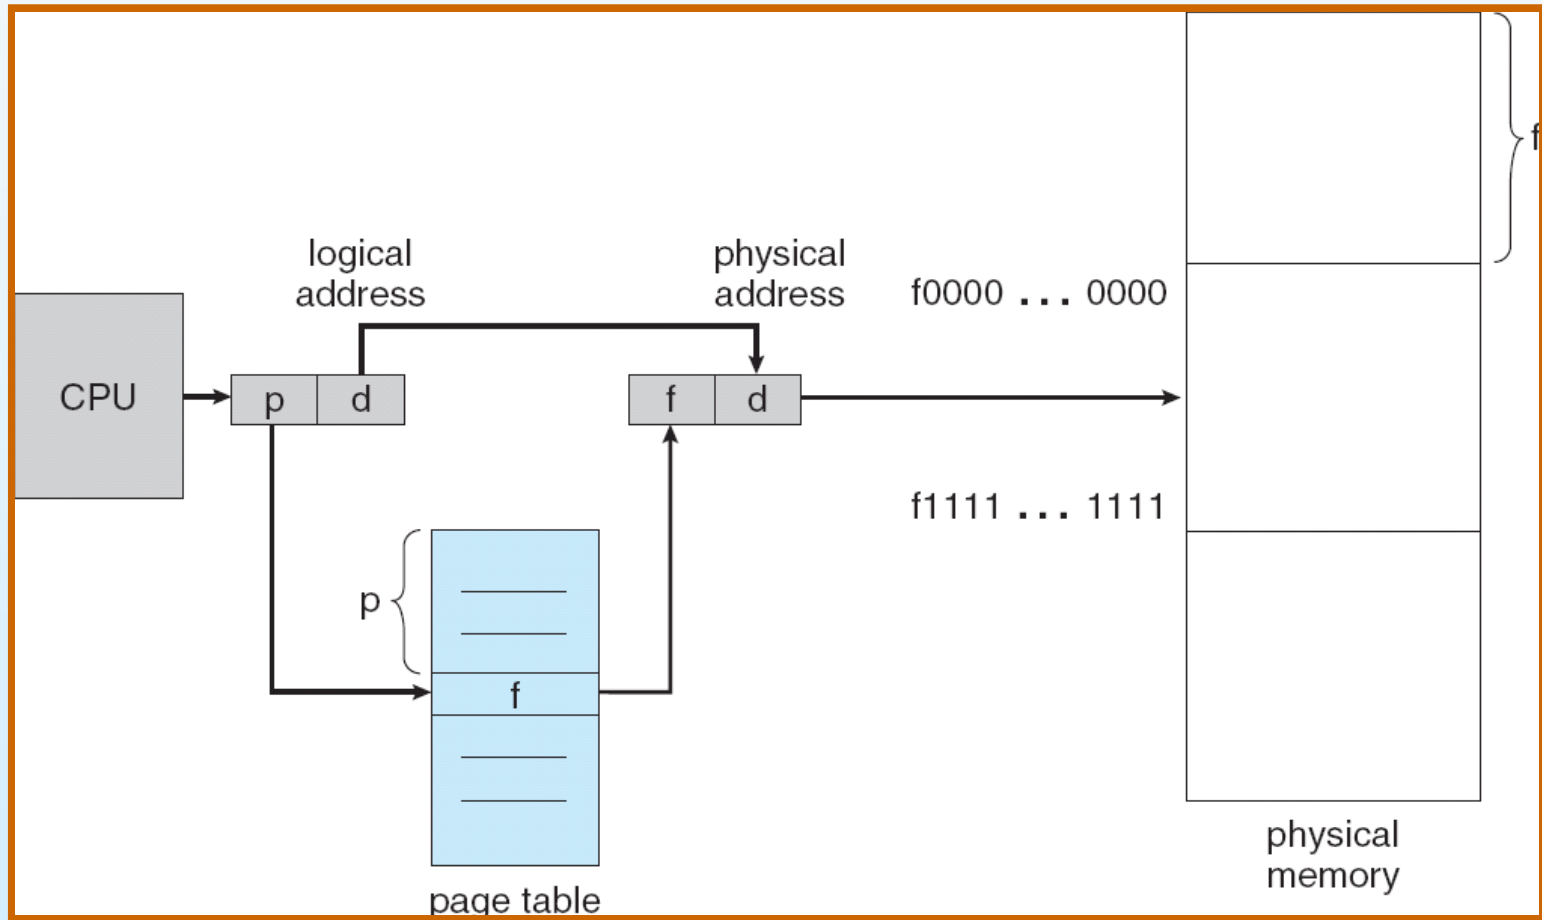


For given logical address space 2^m and page size 2^n



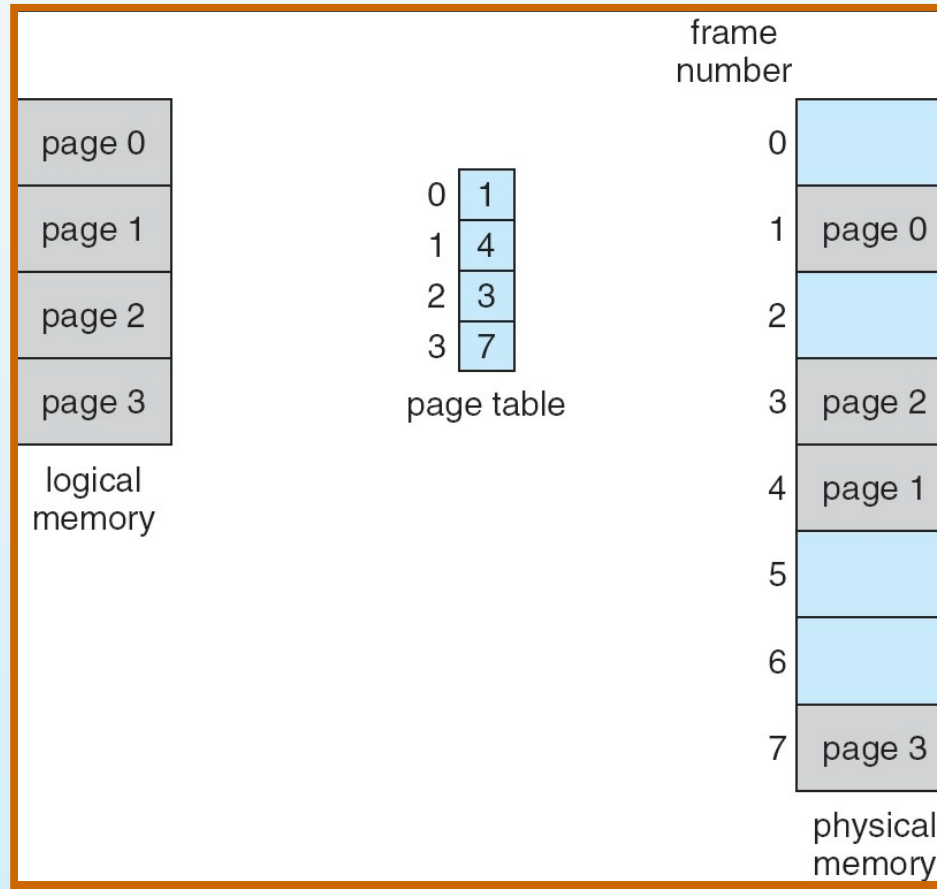


Paging Hardware



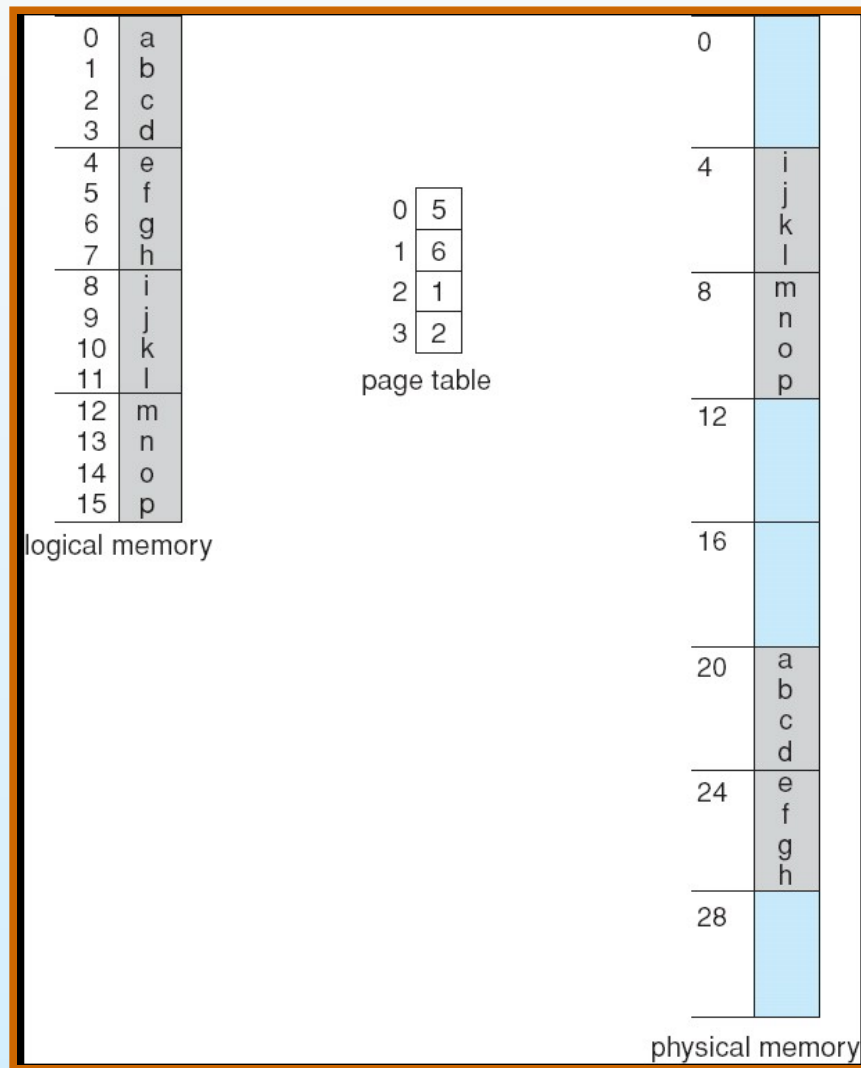


Paging Model of Logical and Physical Memory





Paging Example



32-byte memory and 4-byte pages





Page Size

■ Small page size

- Decrease internal fragmentation
- Increase page table size

■ Big page size

- Increase internal fragmentation
- Decrease page table size

■ Usually 4 KB or 8 KB

■ Solaris in SPARC cpu supports

- 8 KB and 4 MB

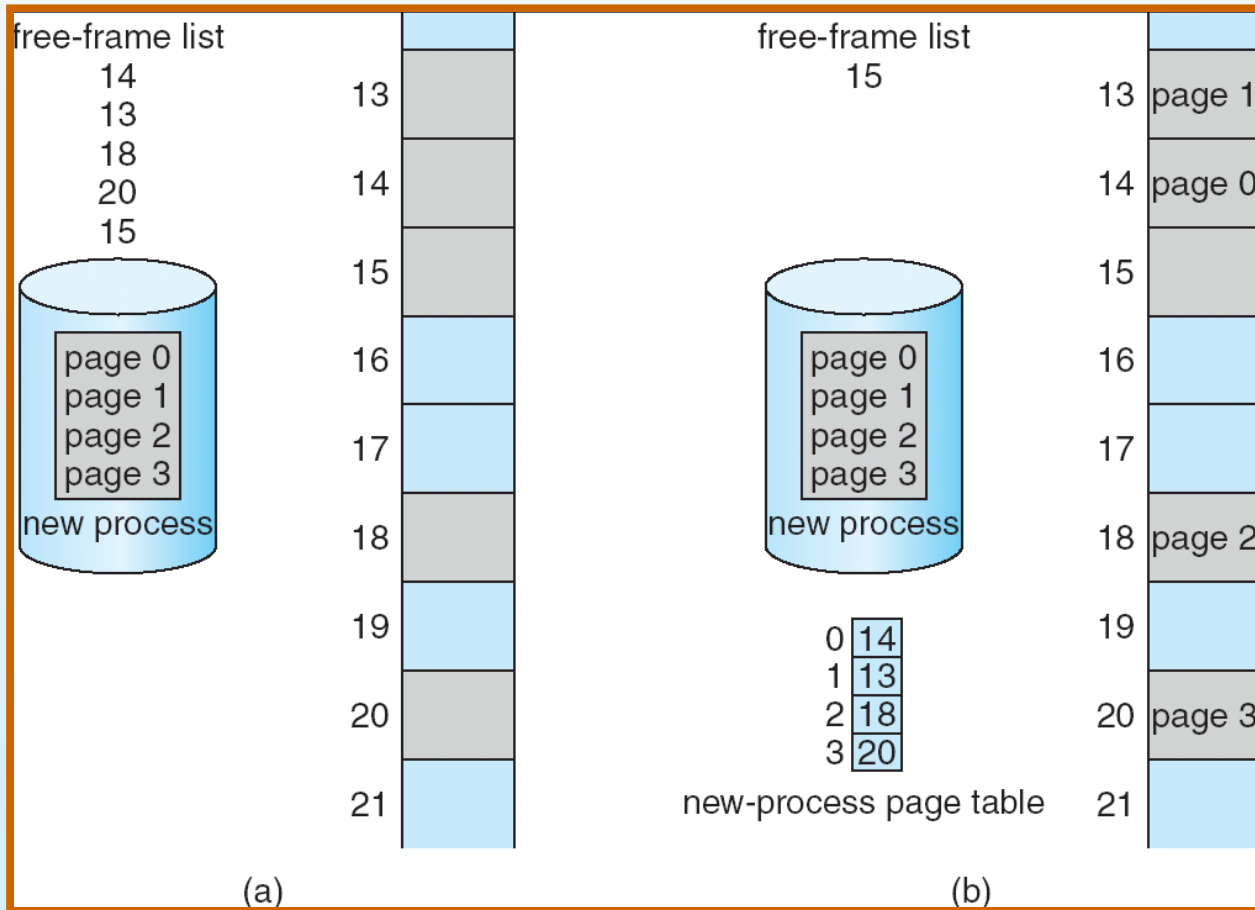
■ Linux 2.6.x support 4KB or 4MB

- Run “getconf PAGESIZE”





Free Frames



Before allocation

After allocation





Implementation of Page Table

- Page table is kept in main memory
 - Page-table base register (PTBR) points to it
 - Page-table length register (PRLR) indicates size
- Every data/instruction access requires two memory accesses.
 - For page table, and for data/instruction.
- Solve it by using TLBs
 - Translation look-aside buffers
 - A special fast-lookup hardware cache





TLB

■ Parallel search

Page #	Frame #


Address translation (p, d)

- If p is in, get frame # out
- Otherwise
 - ▶ Load entry of page p from page table into TLB
 - ▶ Get frame #





Paging Hardware With TLB





Effective Access Time

- Assume memory cycle time is 1 time unit
- Associative Lookup = ϵ time unit
- Hit ratio
 - Percentage of times that a page number is found in TLB
 - Ratio related to size of TLB
- Hit ratio = α
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha) \\ &= 2 + \epsilon - \alpha \end{aligned}$$





Memory Protection

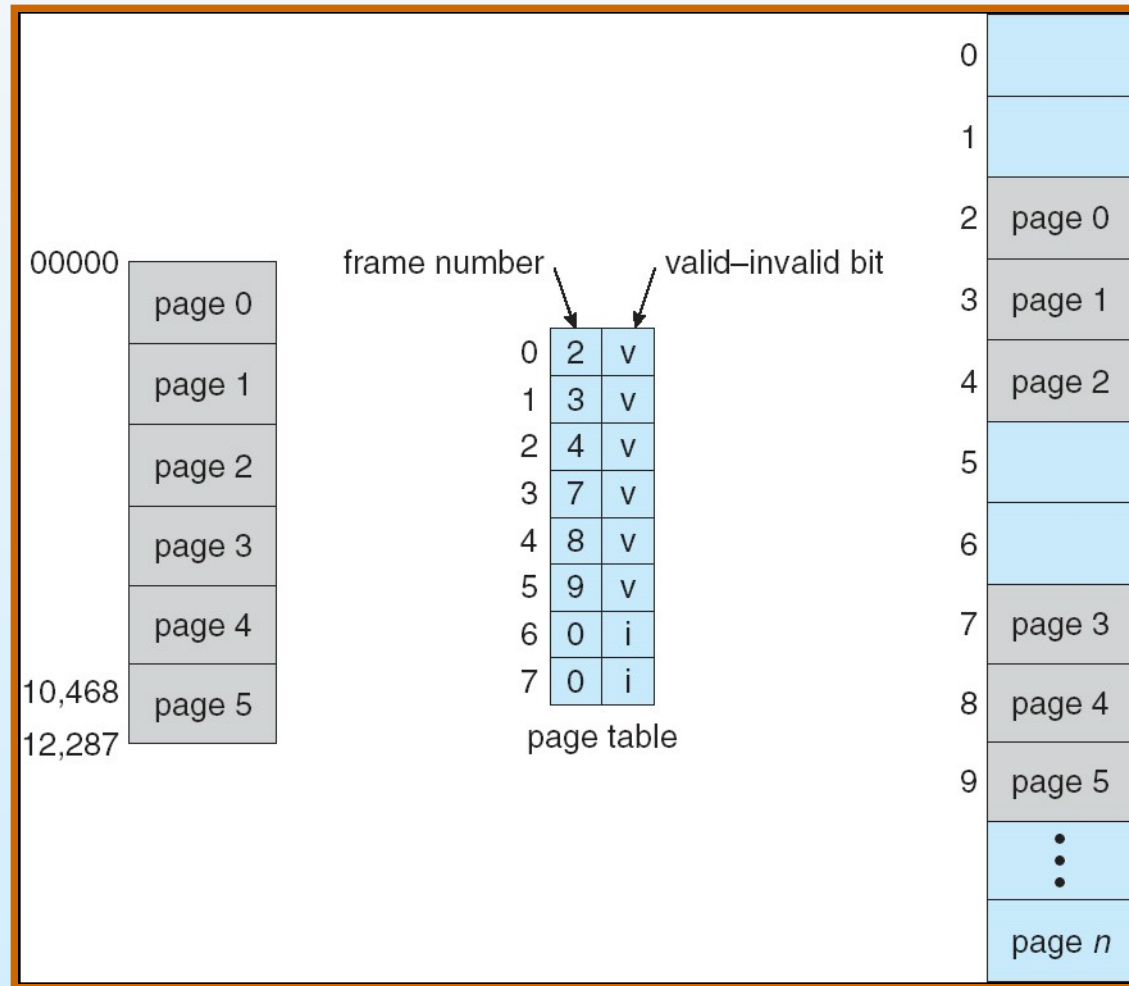
- **Memory protection implemented by**
 - **Associating protection bit with each frame**
 - **Read-write, read-only, execution-only**

- **Valid-invalid bit attached to each entry in the page table:**
 - **Valid**
 - ▶ **The page is in the process' logical address space**
 - **Invalid**
 - ▶ **The page is NOT in the process' logical address space**





Valid (v) or Invalid (i) Bit In A Page Table





Shared Pages

■ Shared code

- One copy of read-only (reentrant) code shared among processes
 - ▶ i.e. DLL, Multi-process program

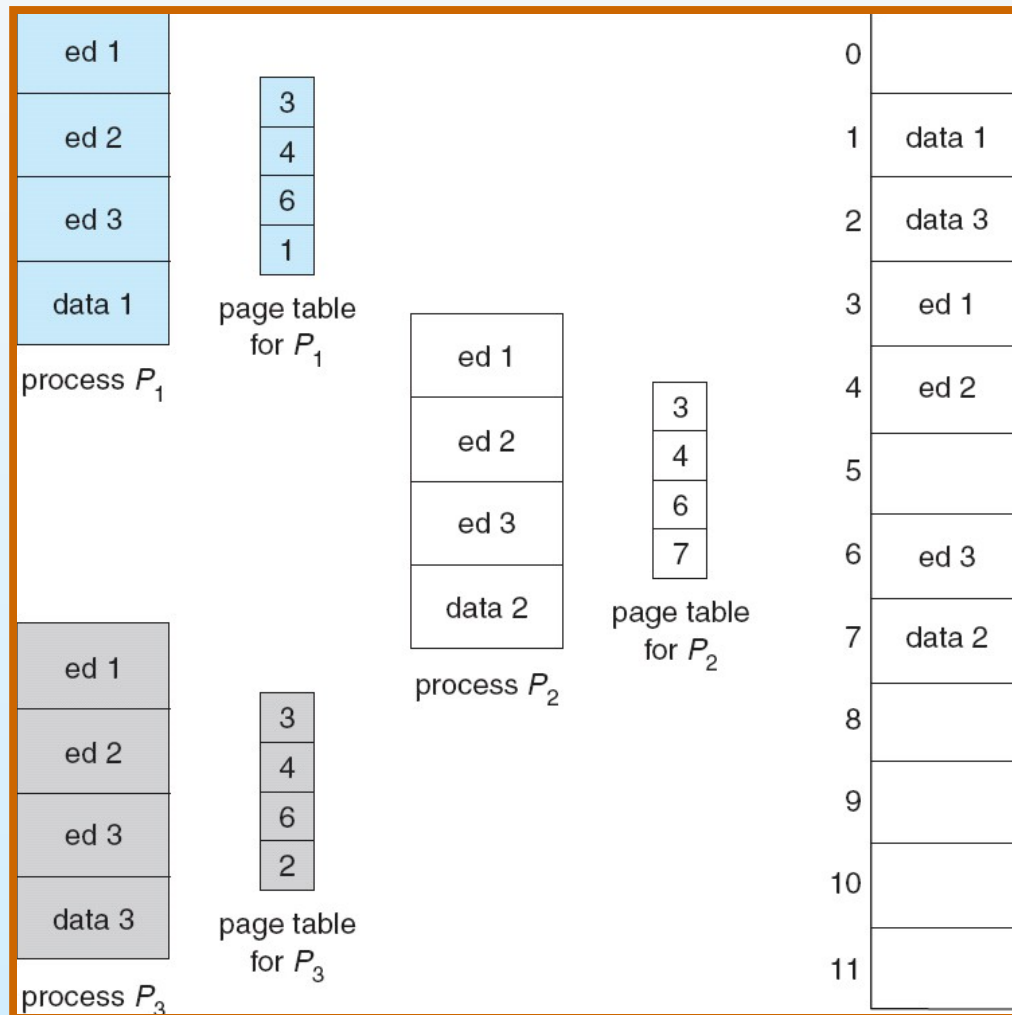
■ Private code and data

- Each process keeps a separate copy of the code and data





Shared Pages Example





Structure of the Page Table

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables





Hierarchical Page Tables

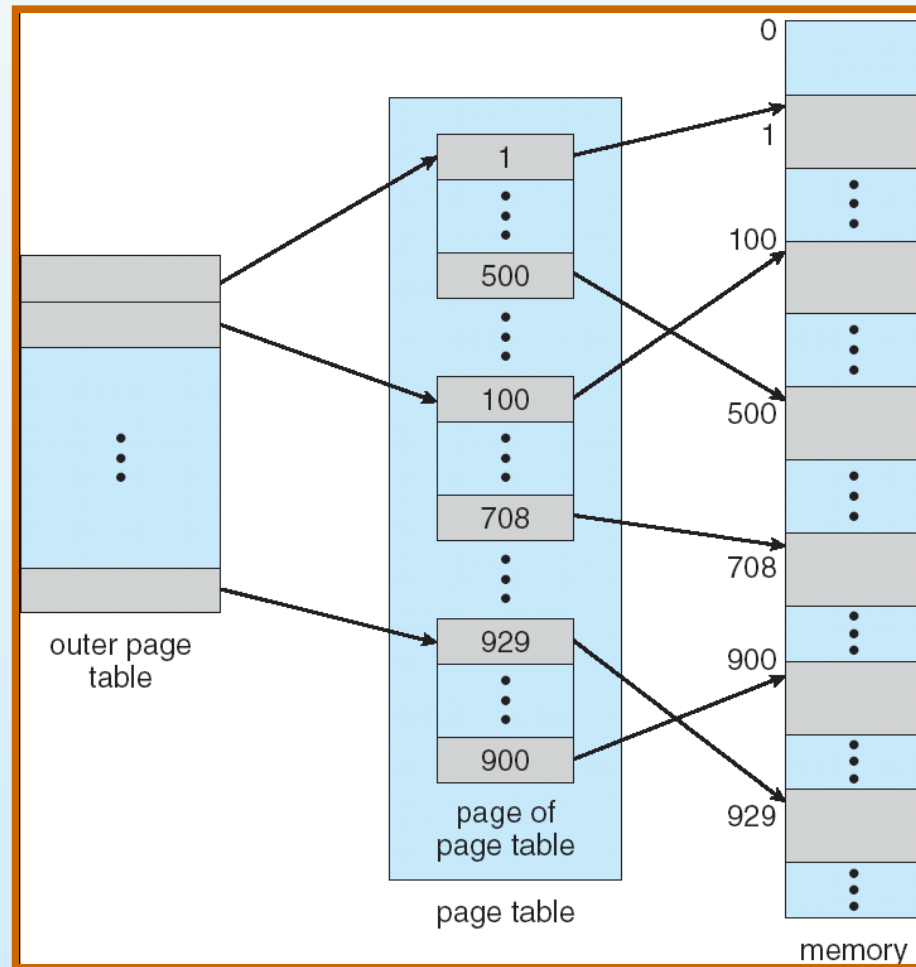
- Break up the logical address space into multiple page tables
 - Adding new tables as requesting new space

- A simple technique is a two-level page table
 - Widely used in IA32






Two-Level Page-Table Scheme





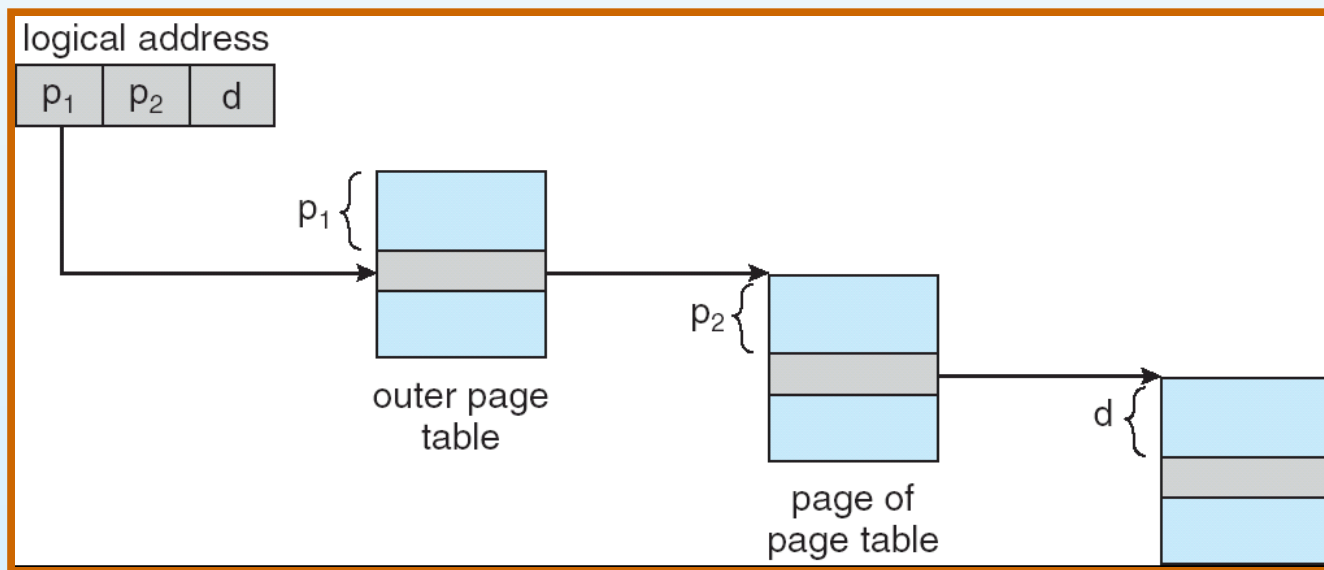
Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a 20 bits page number
 - a 12 bits page offset
- The page number is further divided into
 - a 10-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:





Address-Translation Scheme





Three-level Paging Scheme

outer page	inner page	offset
p_1	p_2	d
42	10	12

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12





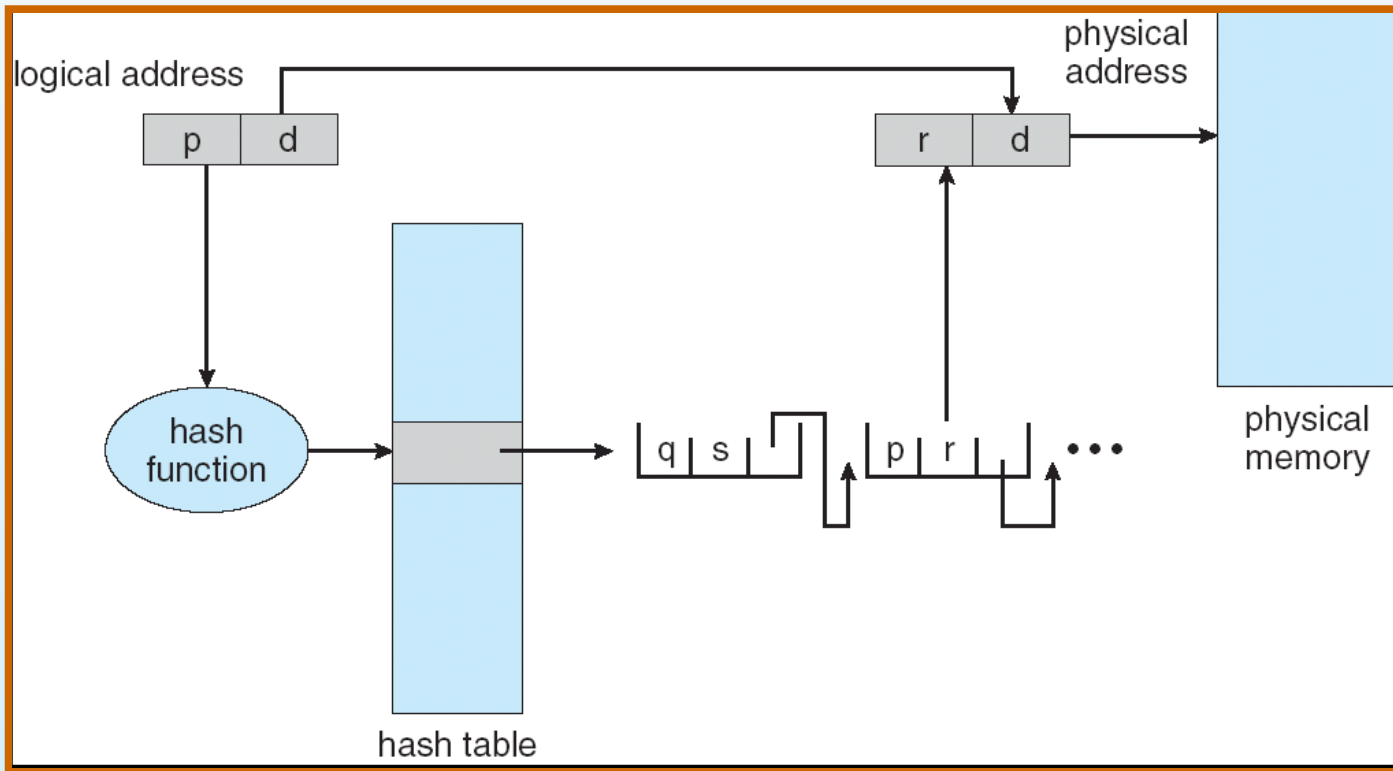
Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table.
- This page table contains a chain of elements hashing to the same location.
- Virtual page numbers are compared in this chain searching for a match.
- If a match is found, the corresponding physical frame is extracted.





Hashed Page Table





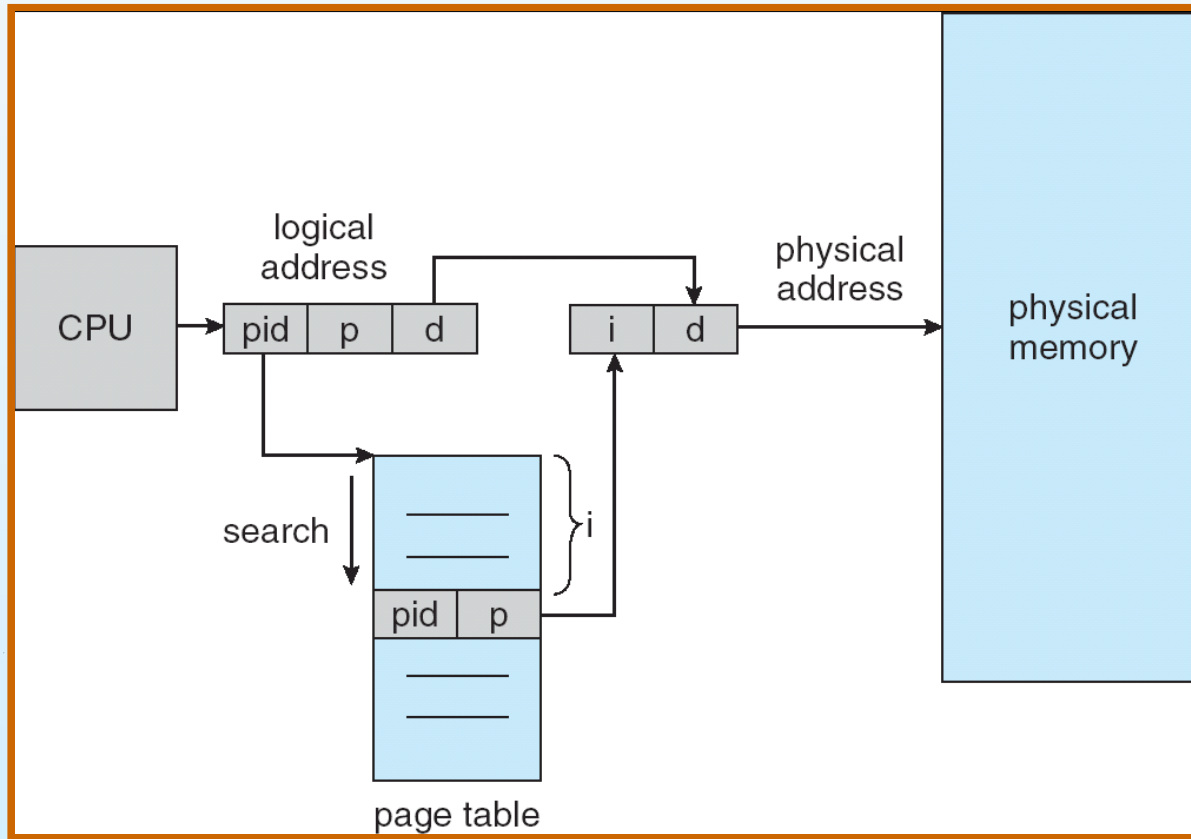
Inverted Page Table

- One entry for each real page of memory
- Entry consists of
 - PID
 - Virtual page number which stored in the real memory
- Decreases memory using but increases time to search the table
- Use hash table to limit the search to one — or at most a few — page-table entries





Inverted Page Table Architecture






Segmentation

- **Memory-management scheme that supports user view of memory**
- **A program is a collection of segments.**
- **A segment is a logical unit such as:**
 - **Code**
 - **Global variable**
 - **Heap**
 - **Stacks**
 - **Libraries**






User's View of a Program






Logical View of Segmentation



user space



physical memory space





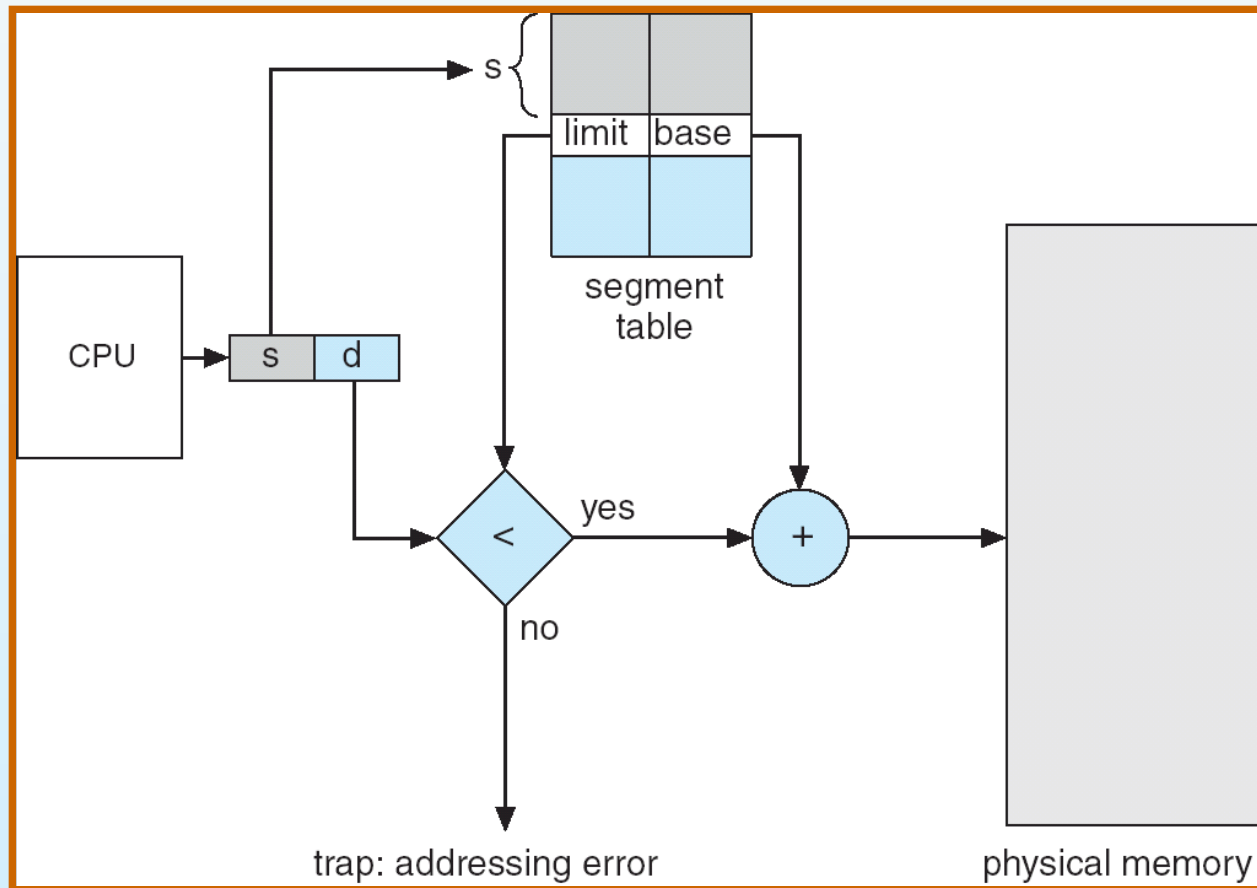
Segmentation Architecture

- Logical address consists of a two tuple:
 <segment-number, offset>,
- Segment table do the mapping
- Each table entry has:
 - Base – the starting physical address of segment
 - Limit – the length of the segment
- Segment-table base register (STBR) points to the segment table
- Segment-table length register (STLR) indicates number of segments used by a program;
 segment number **s** is legal if **s < STLR**





Segmentation Hardware





Example: The Intel Pentium

■ Supports

- Segmentation
- Paging
- Segmentation with paging


■ CPU generates logical address

- Given to segmentation unit
 - ▶ Which produces linear addresses
- Linear address given to paging unit
 - ▶ Which generates physical address





Logical to Physical Address Translation in Pentium

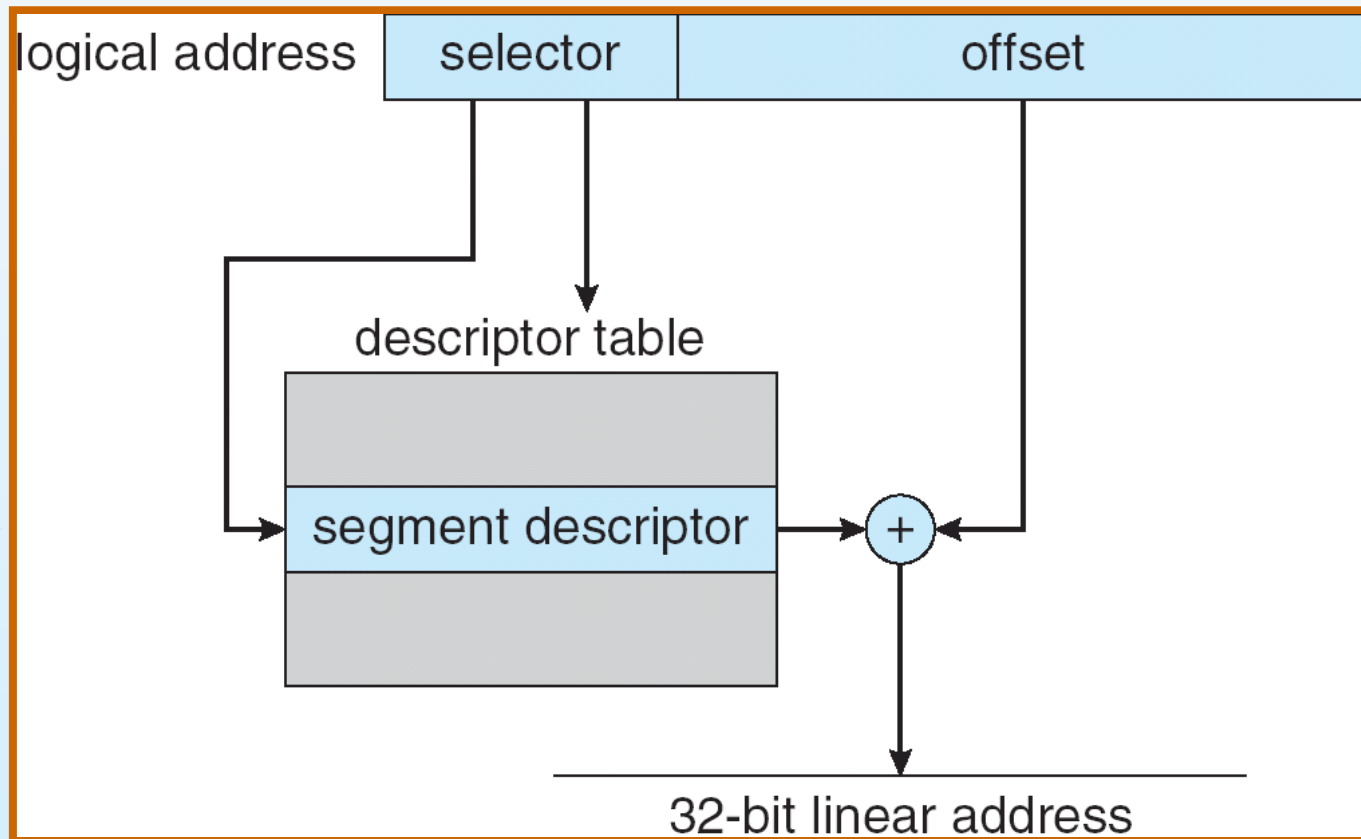


page number		page offset
p_1	p_2	d
10	10	12



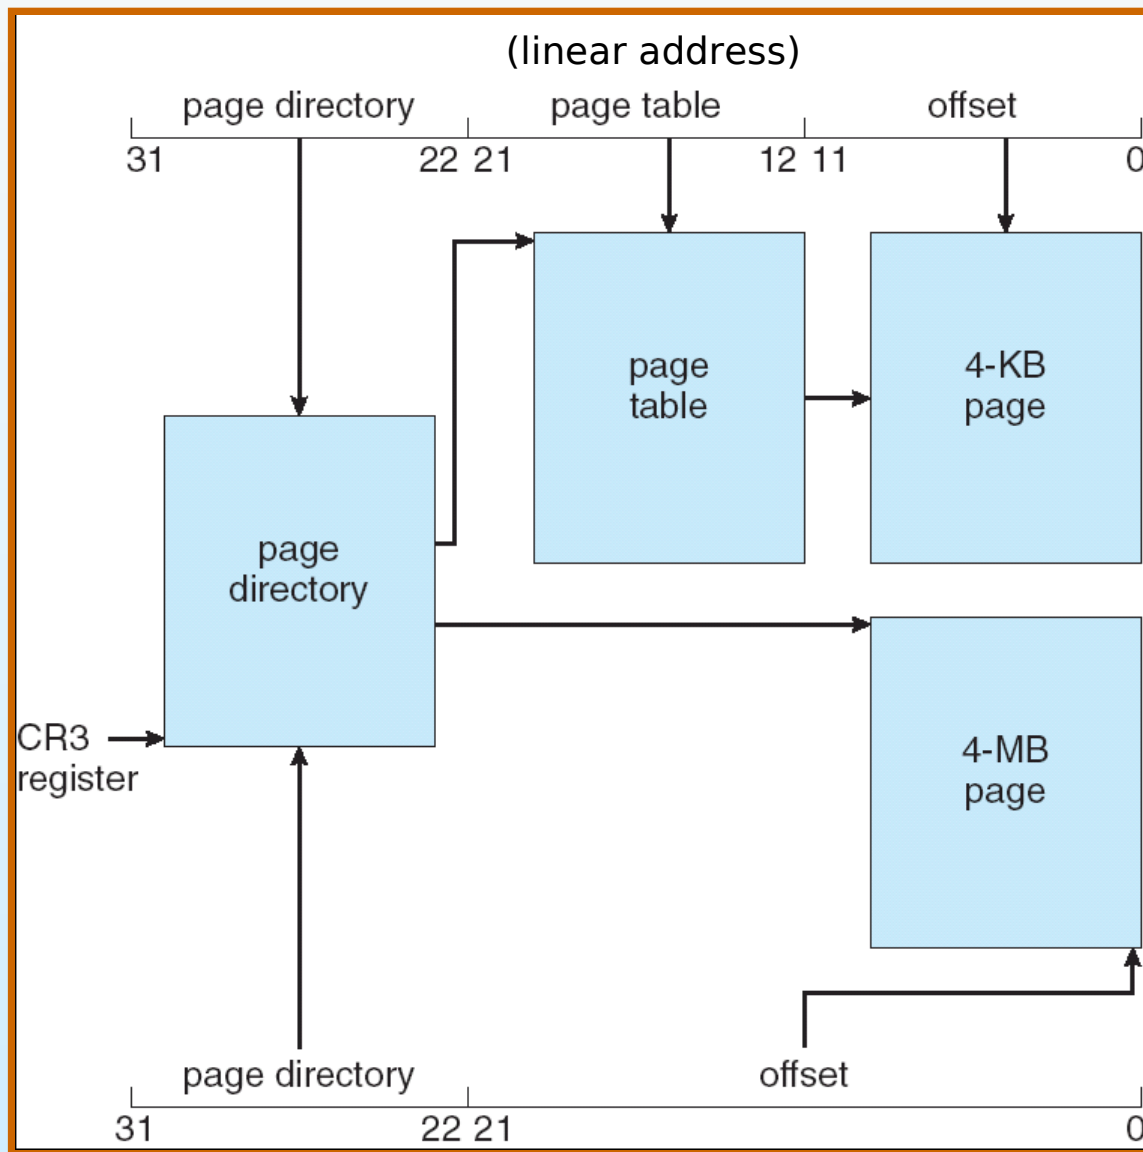


Intel Pentium Segmentation





Pentium Paging Architecture





Linear Address in Linux

- To support more platform, linear address is broken into four parts:

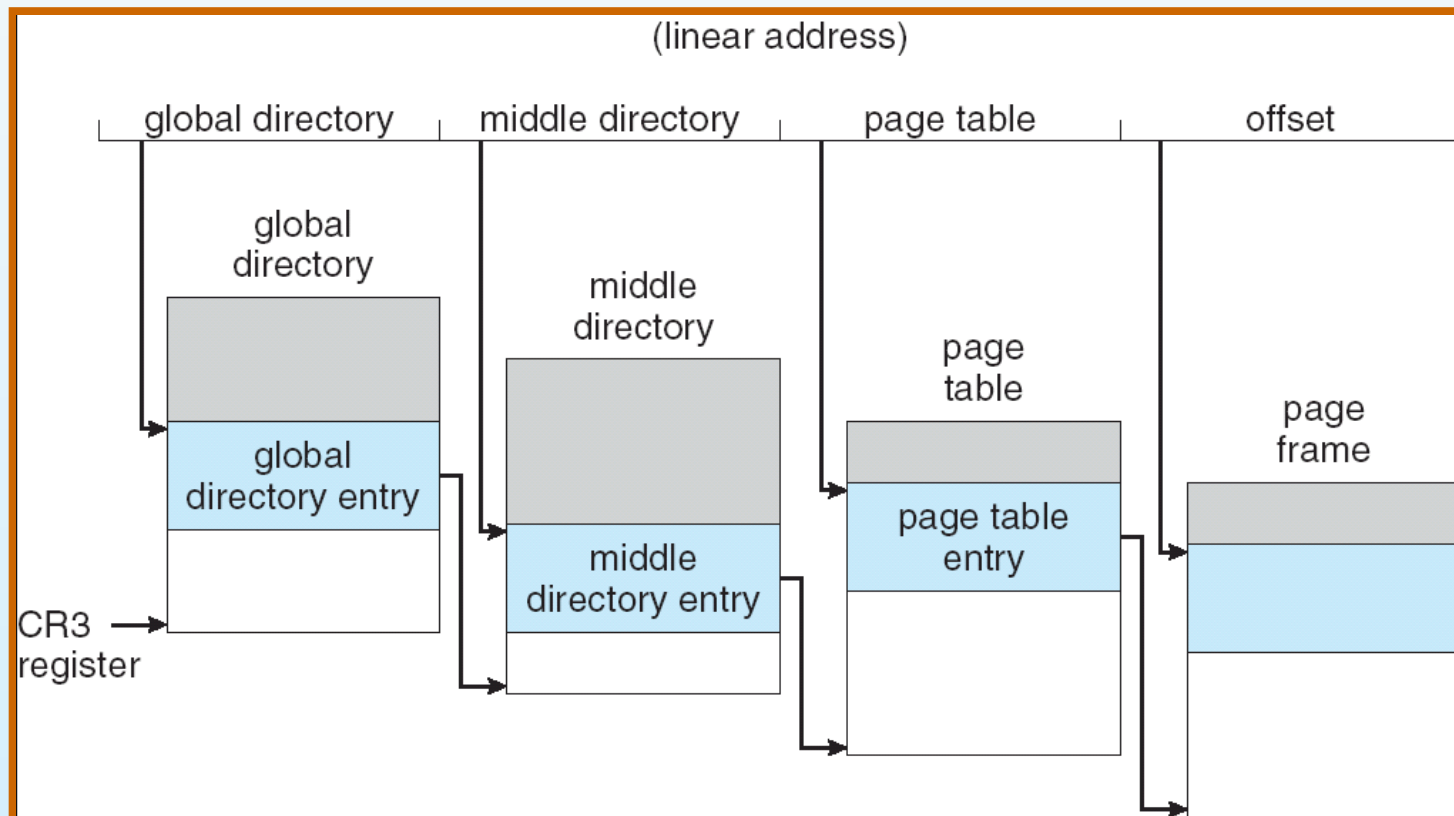
global directory	middle directory	page table	offset
---------------------	---------------------	---------------	--------

- In IA32, the number of bits of middle-directory is 0





Three-level Paging in Linux



End of Chapter 8

