

Chapter 9: Virtual Memory





Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Other Considerations
- Operating-System Examples





Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of
 - demand paging
 - page-replacement algorithms
 - allocation of page frames
- To discuss the principle of the working-set model





Background

■ Virtual Memory

- Only part of the program needs to be in memory
- Logical address space can be much larger than physical address space
 - ▶ 0 to 0xFFFFFFFF...
- Increase CPU utilization and throughput
 - ▶ No affect on response time and turnaround time
- Free the programmer!


■ Virtual memory can be implemented via:

- Demand paging
- Demand segmentation





Virtual Memory That is Larger Than Physical Memory





Demand Paging

- **Bring a page into memory only when it is needed**
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More processes
- **Page is needed \Rightarrow reference to it**
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- **Lazy swapper**
 - Never swaps page into memory unless it'll be needed
 - Swapper that deals with pages is a pager





Valid-Invalid Bit

- A valid–invalid bit is included in each page table entry
 - **v** – in-memory
 - **i** – not-in-memory
- Initially the bit is set to **i** on all entries
- During address translation, if the bit is **i**, **PAGE FAULT!**

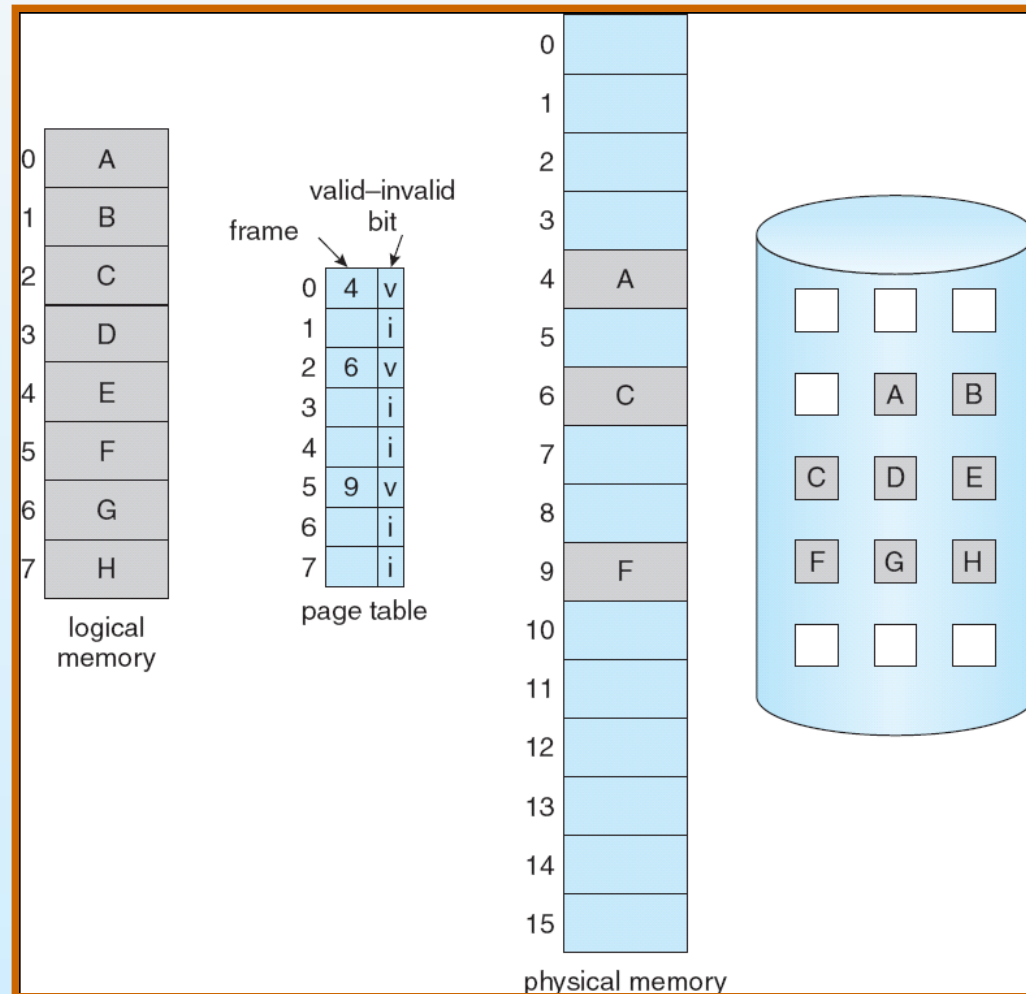
Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table





Page Table When Some Pages Are Not in Main Memory





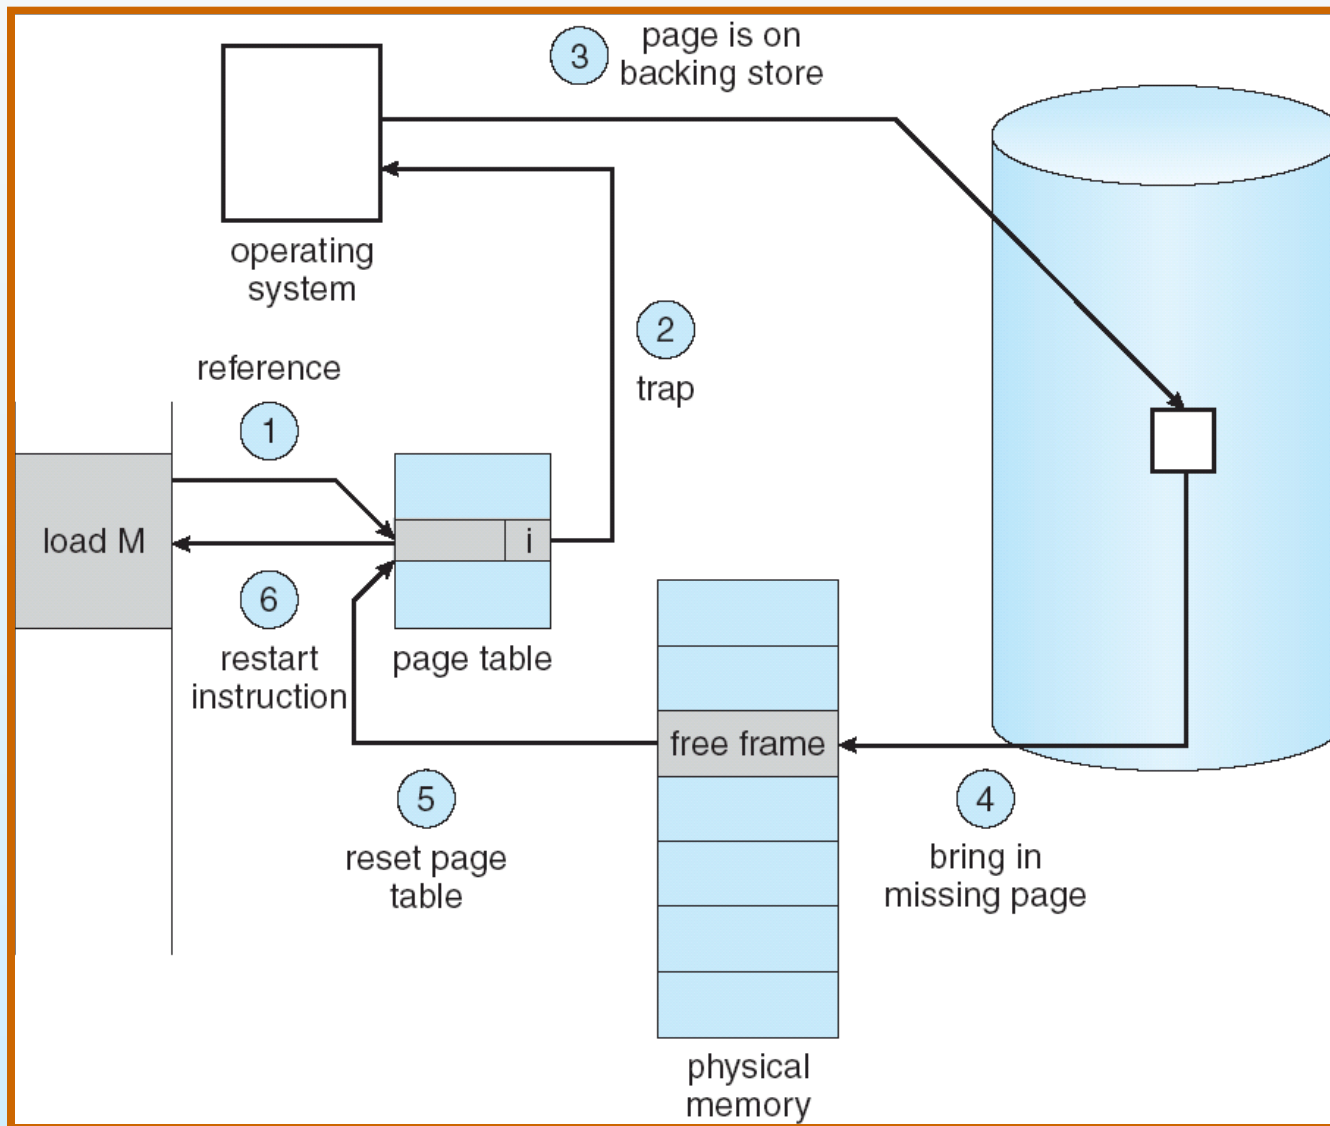
Page Fault

- First reference to an invalid page will trap to OS
 - page fault
- OS looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
- Get empty frame
- Swap page into frame
- Reset tables
- Set validation bit = **v**
- Restart the instruction that caused the page fault





Steps in Handling a Page Fault





Performance of Demand Paging

■ Page Fault Rate $0 \leq p \leq 1.0$

- if $p = 0$ no page faults
- if $p = 1$, every reference is a fault

■ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead} \\ &) \end{aligned}$$





Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
 $= (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + p \times 7,999,800$
- If page fault rate = $1/1000$, then $EAT = 8.2 \text{ ms}$
This is a slowdown by a factor of 40!!





Copy-on-Write

- **Copy-on-Write (COW)**
 - Both parent and child processes initially *share* the same pages in memory
 - If either process modifies a page, the page is copied

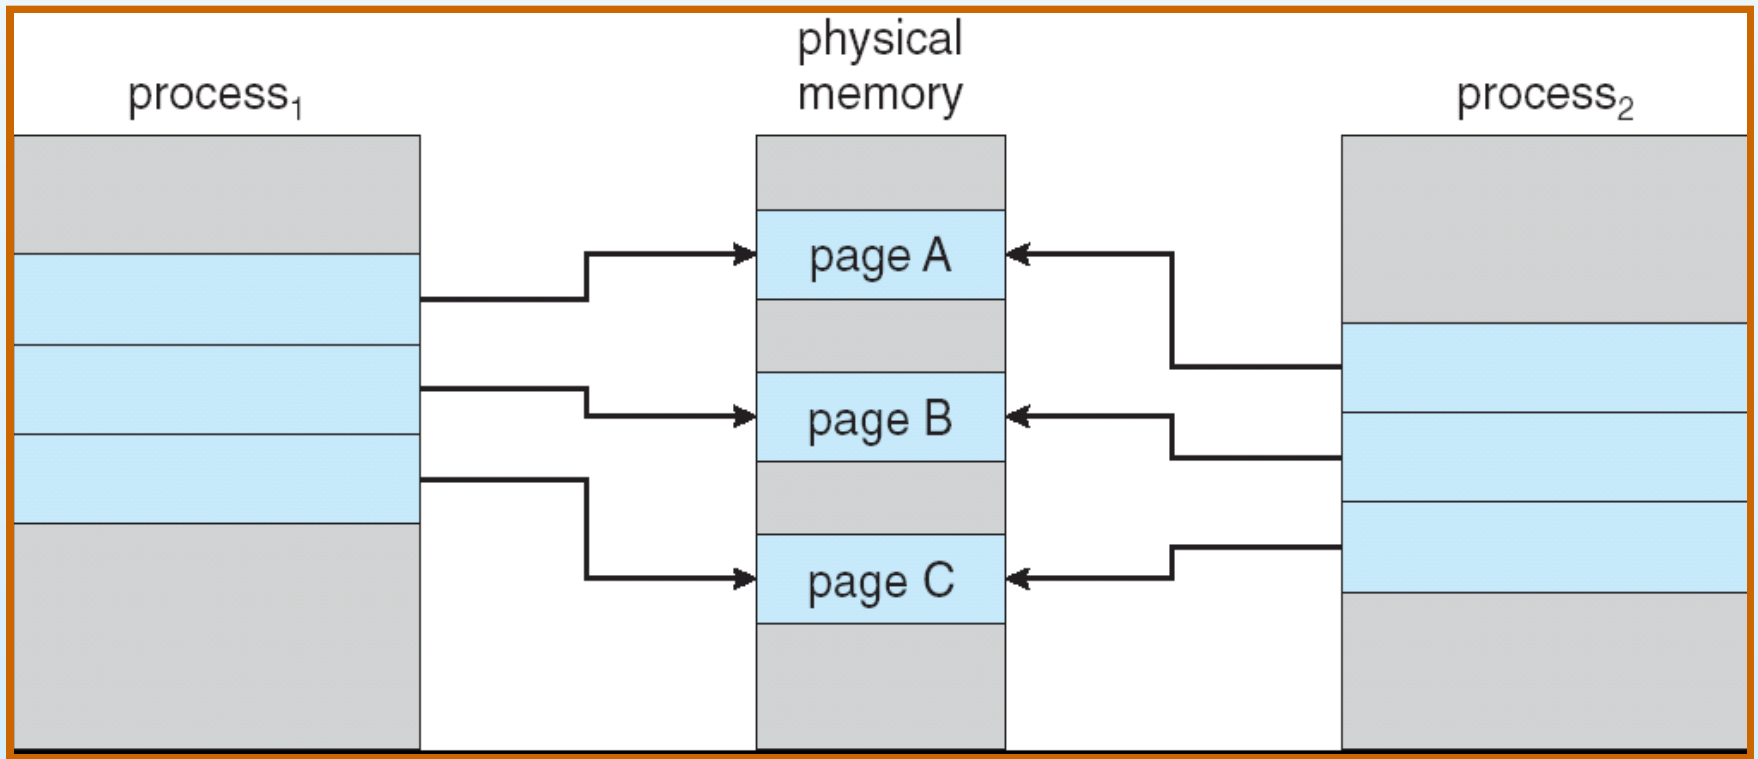
- **COW allows more efficient process creation**

- **Free pages are allocated from a pool of zeroed-out pages**



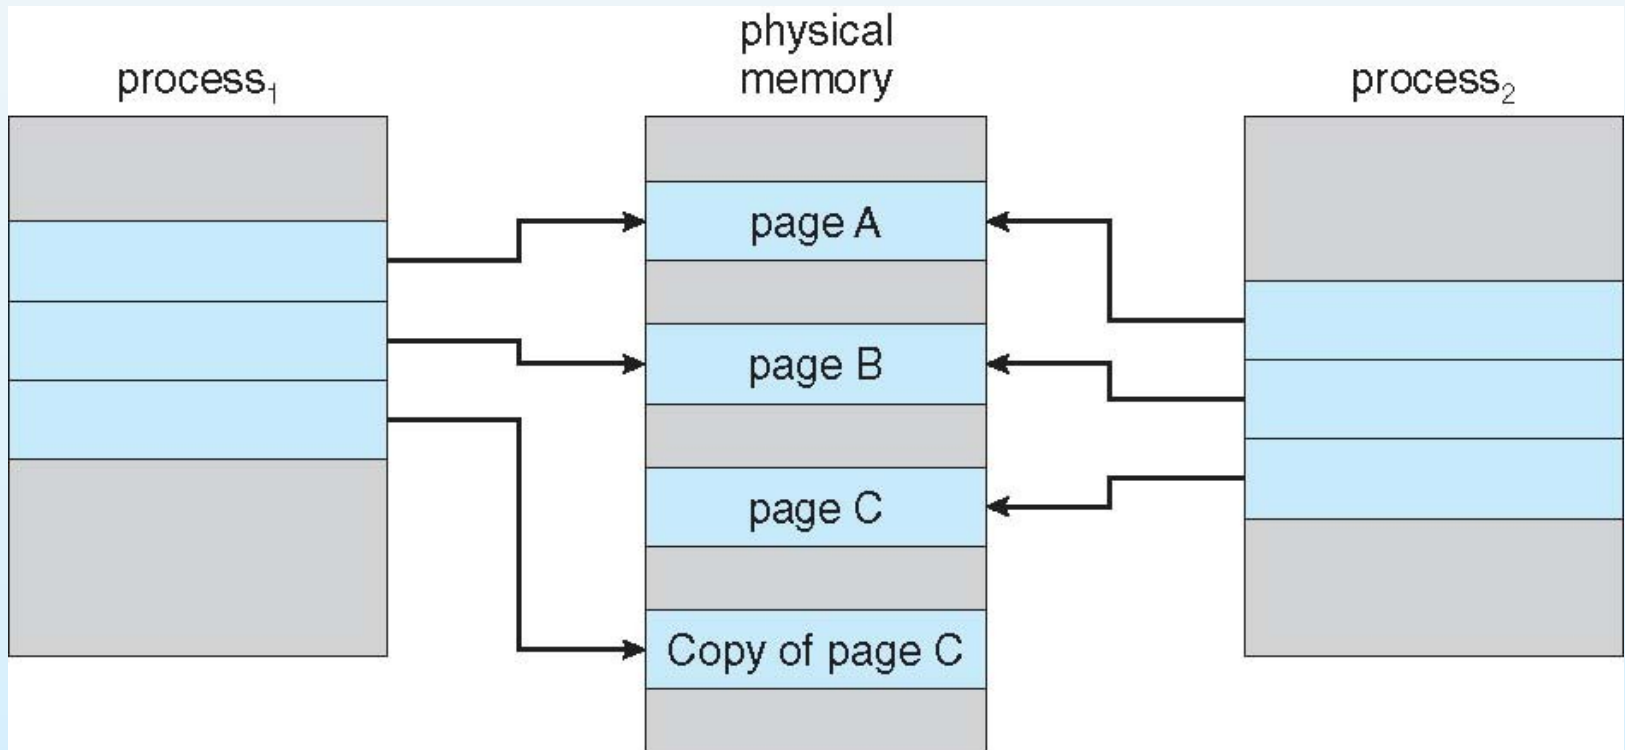


Before Process 1 Modifies Page C





After Process 1 Modifies Page C





What happens if there is no free frame?

■ Page replacement

- Find some page in memory, but not really in use, swap it out
- Algorithm
- Performance – minimize number of page faults

■ Same page may be brought into memory several times






Basic Page Replacement

- Find the location of the desired page on disk
- Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a victim frame
 - Swap out the victim if its modified (dirty) bit is set
- Bring the desired page into the (newly) free frame
- Update the page and frame tables
- Continue to run the process





Page Replacement






Page Replacement Algorithms

- Want lowest page-fault rate
- Evaluate algorithm
 - Running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- In all our examples, the reference string is
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5





Graph of Page Faults Versus The Number of Frames





First-In-First-Out (FIFO) Algorithm

- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	9 page faults
2	2	1	3	
3	3	2	4	

- 4 frames


1	1	5	4	10 page faults
2	2	1	5	
3	3	2		
4	4	3		

- Belady's Anomaly: more frames -> more page faults






FIFO Illustrating Belady's Anomaly





Optimal Algorithm

- Replace page that will not be used for longest period of time
- 4 frames example



- How do you know this?
- Used for measuring how well your algorithm performs





Least Recently Used (LRU) Algorithm

■ Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3





LRU Algorithm Implementation

■ Counter implementation

- Every page entry has a counter
- Every time page is referenced, copy the clock into the counter
- Search for the page with smallest counter
- Replace it

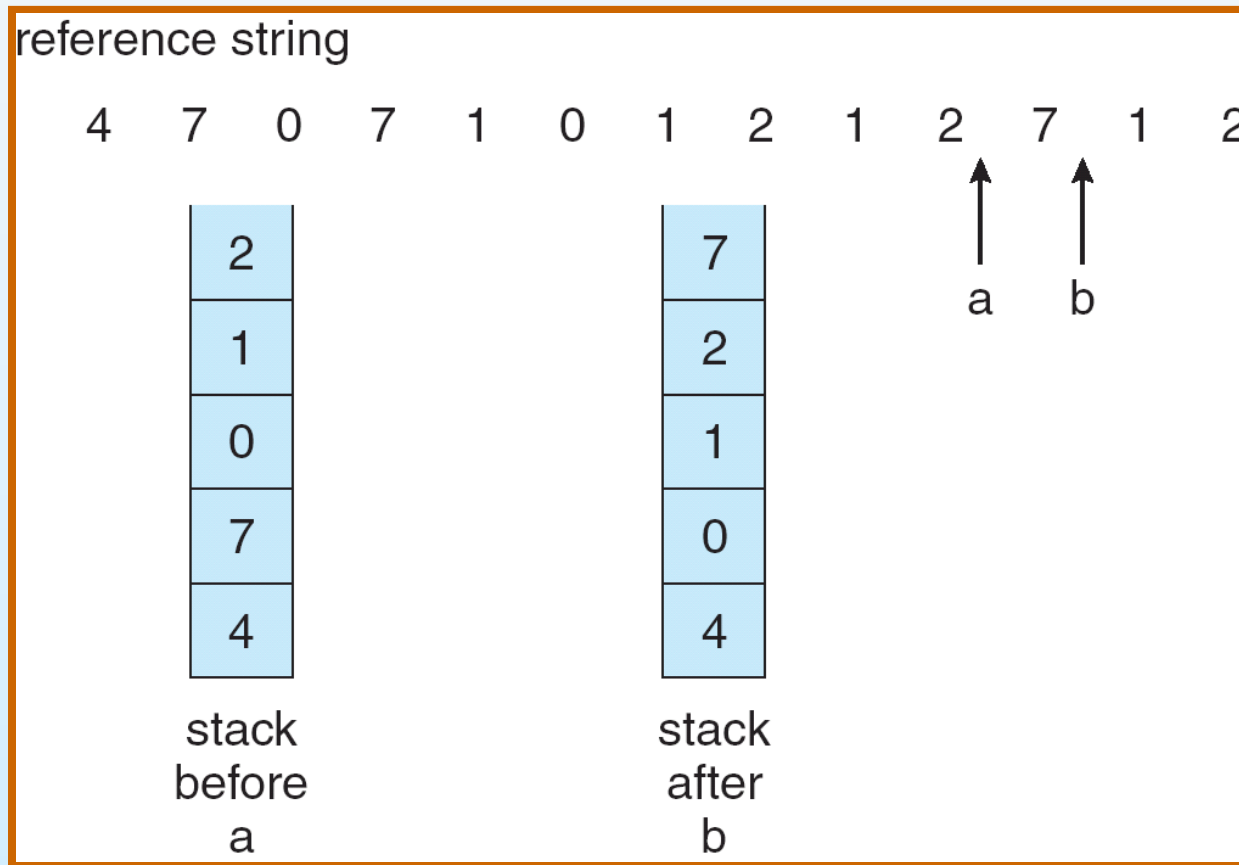
■ Stack implementation – keep a stack of page numbers in a double link form:

- Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
- No search for replacement





Use Of A Stack To Record The Most Recent Page References





LRU Approximation Algorithms

■ Reference bit

- With each page associate a bit, initially = 0
- When page is referenced, CPU set bit to 1
- Replace the one which is 0 (if one exists)
- We do not know the order, however






Additional-Reference-Bits Algorithm

	R bits for pages 0-5, clock tick 0	R bits for pages 0-5, clock tick 1	R bits for pages 0-5, clock tick 2	R bits for pages 0-5, clock tick 3	R bits for pages 0-5, clock tick 4
	1 0 1 0 1 1	1 1 0 0 1 0	1 1 0 1 0 1	1 0 0 0 1 0	0 1 1 0 0 0
Page					
0	10000000	11000000	11100000	11110000	01111000
1	00000000	10000000	11000000	01100000	10110000
2	10000000	01000000	00100000	00100000	10001000
3	00000000	00000000	10000000	01000000	00100000
4	10000000	11000000	01100000	10110000	01011000
5	10000000	01000000	10100000	01010000	00101000
	(a)	(b)	(c)	(d)	(e)





Second-Chance (clock) Page-Replacement Algorithm



When a page fault occurs, the page the hand is pointing to is inspected. The action taken depends on the R bit:

R = 0: Evict the page

R = 1: Clear R and advance hand





Enhanced Second-Chance Algorithm

■ (reference bit, modify bit)

- (0, 0) – neither recently used, nor modified
- (0, 1) – not recently used, but modified
- (1, 0) – recently used, but clean
- (1, 1) – recently used and modified

■ In which order to evict?





Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- Least frequently used (LFU) Algorithm
 - Replaces page with smallest count
- Most frequently used (MFU) Algorithm
 - The page with the smallest count was probably just brought in and has yet to be used





Page Buffering Algorithms

- OS keeps a pool of free frames
 - Allocate from the pool
 - Swap out a frame lately and add it to the pool
- When system is idle, write back the modified pages
- If a requested page is in the pool, get it immediately





Global vs. Local Allocation

■ Global replacement

- Get a replacement frame from the set of all frames
- One process can take a frame from another

■ Local replacement

- Each process selects from only its own set of allocated frames

■ Thrashing

- A process is busy swapping pages in and out





Locality Model

- To prevent thrashing, provide a process with as many frames as it needs.
- A locality is a set of pages that are actively used together
- Processes migrates from one locality to another locality always





Working-Set Model

- $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instruction
- WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ
(varies in time)
 - if Δ too small will not encompass entire locality
 - if Δ too large will encompass several localities
 - if $\Delta = \infty \Rightarrow$ will encompass entire program
- $D = \sum WSS_i \equiv$ total demand frames
- if $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend one of the processes







Working-set model

page reference table

. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$





Keeping Track of the Working Set


- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = 1 \Rightarrow page in working set
- Improvement = 10 bits and interrupt every 1000 time units





Page-Fault Frequency Scheme

- Establish “acceptable” page-fault rate
 - If actual rate too low, process loses frame
 - If actual rate too high, process gains frame





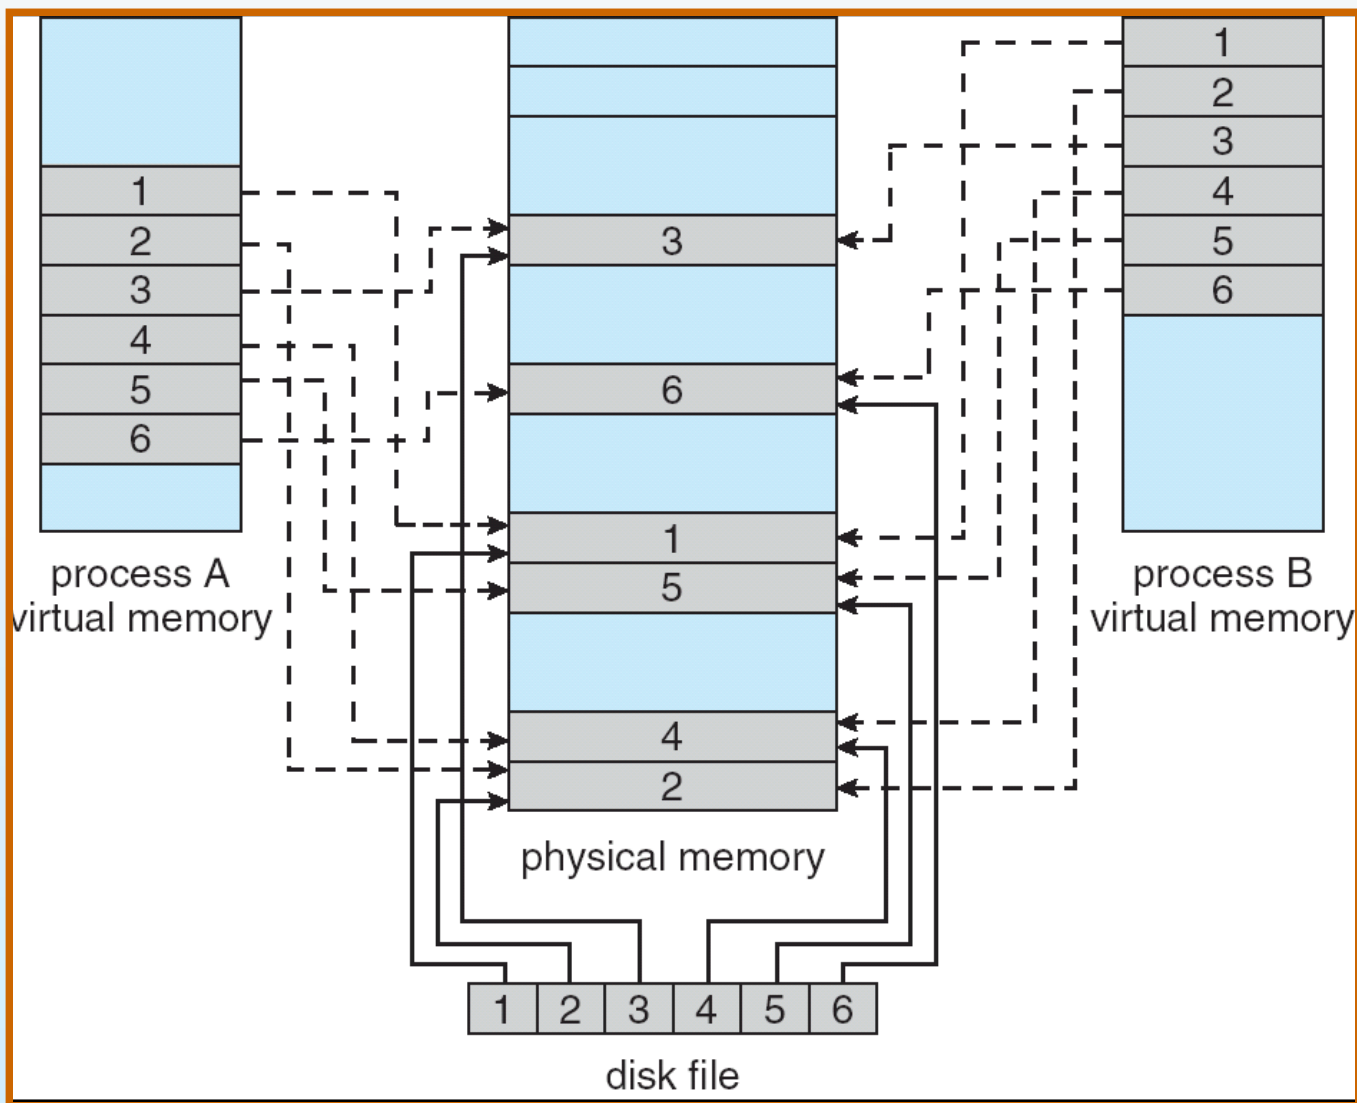
Memory-Mapped Files

- **Memory-mapped file I/O allows file I/O to be treated as memory access by mapping a file to pages in memory**
 - A file is initially read using demand paging.
 - A page-sized portion of the file is read into a frame.
 - Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- **Simplifies file access rather than `read()` `write()` system calls**
- **Also allows several processes to map the same file allowing the pages in memory to be shared**






Memory Mapped Files





Memory-Mapped Shared Memory in Windows





Other Issues – Program Structure

■ Program structure

- `int data[1024][1024];`
- Each row is stored in one page
- Program 1

```
for (j = 0; j < 1024; j++)  
    for (i = 0; i < 1024; i++)  
        data[i][j] = 0;
```

- Program 2

```
for (i = 0; i < 1024; i++)  
    for (j = 0; j < 1024; j++)  
        data[i][j] = 0;
```





Other Issues – Program Structure(Cont.)

- **Array vs. List/Hash table**
- **Local variable, global variable vs. heap variable**
- **Align to page edge**
- **Compiler and loader can affect locality**





Other Issues – I/O interlock

- **DMA from a device**
 - **DMA to kernel space, then copy to user space**
 - **lock the frame. It will not be replaced until unlock**





Operating System Examples

- Windows XP

- Solaris





Windows XP

- Uses demand paging with clustering.
- Clustering brings in pages following the faulting page.
- Processes are assigned working set minimum and working set maximum
- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory falls below a threshold, automatic working set trimming is performed
- Working set trimming removes pages from processes that have pages in excess of their working set minimum
 - Using clock algorithm






Solaris

- Maintains a list of free pages
- *Lotsfree* – threshold parameter (amount of free memory) to begin paging
- *Desfree* – threshold parameter to increasing paging
- *Minfree* – threshold parameter to being swapping
- Paging is performed by *pageout* process
- Pageout scans pages using modified clock algorithm
 - Two hands
- *Scanrate* is the rate at which pages are scanned. Ranges from *slowscan* to *fastscan*
- Pageout is called more frequently depending upon the amount of free memory available





Solaris 2 Page Scanner



End of Chapter 9

