

# Chapter 13: I/O Systems





# Chapter 13: I/O Systems

- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Transforming I/O Requests to Hardware Operations
- Streams
- Performance





# Objectives

- Explore the structure of an operating system's I/O subsystem
- Discuss the principles of I/O hardware and its complexity
- Provide details of the performance aspects of I/O hardware and software





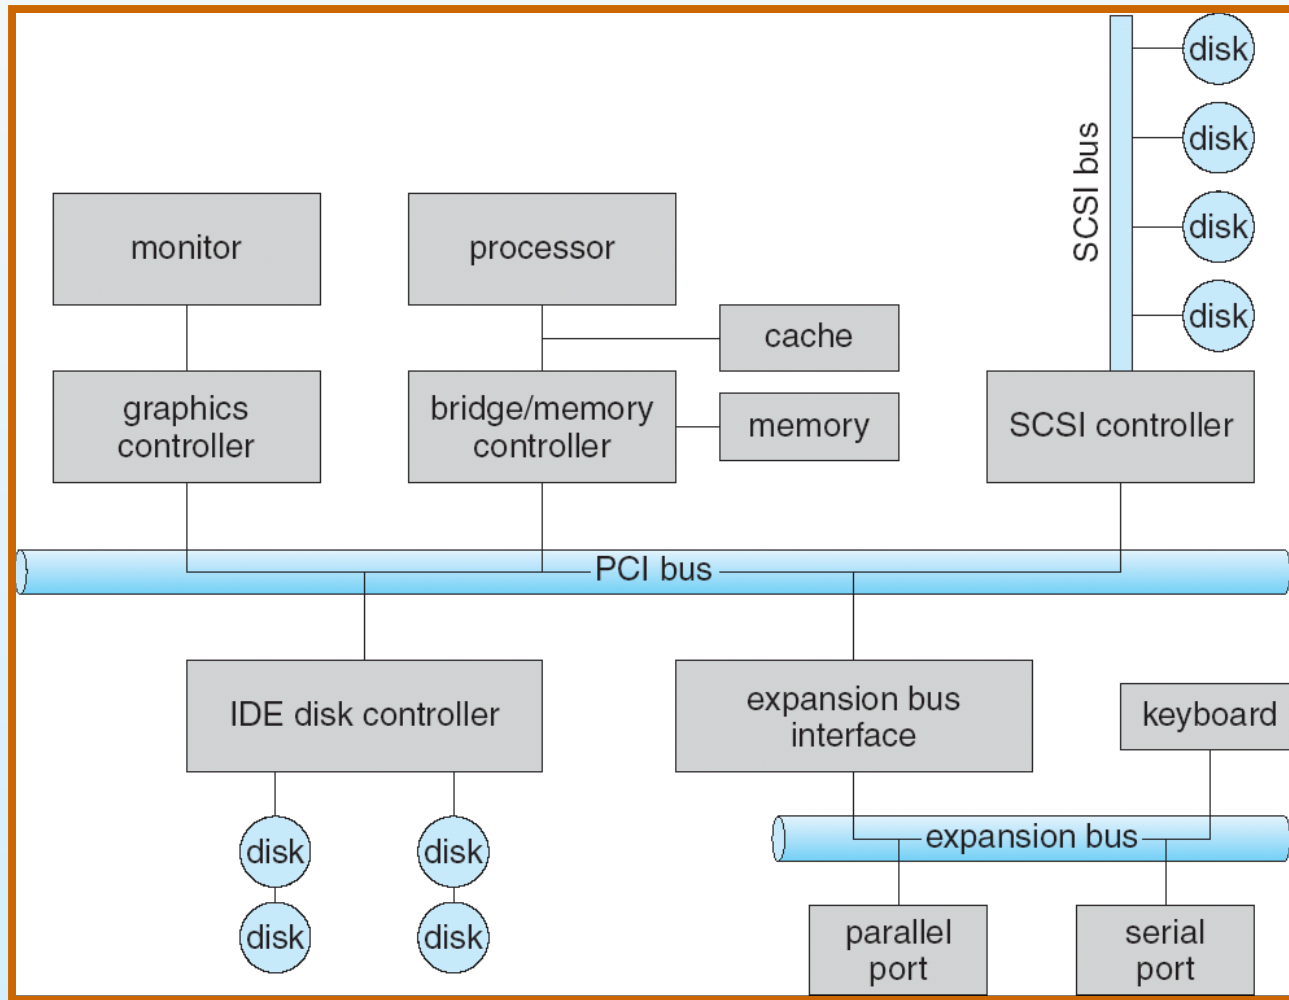
# I/O Hardware

- Incredible variety of I/O devices
- Common concepts
  - Port
  - Bus
  - Controller (host adapter)
- I/O instructions control devices
- Devices have addresses, used by
  - Direct I/O instructions
  - Memory-mapped I/O





# A Typical PC Bus Structure





# Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)





# Polling

- **Determines state of device**
  - **command-ready**
  - **busy**
  - **Error**
- **Busy-wait cycle to wait for I/O from device**





# Interrupts

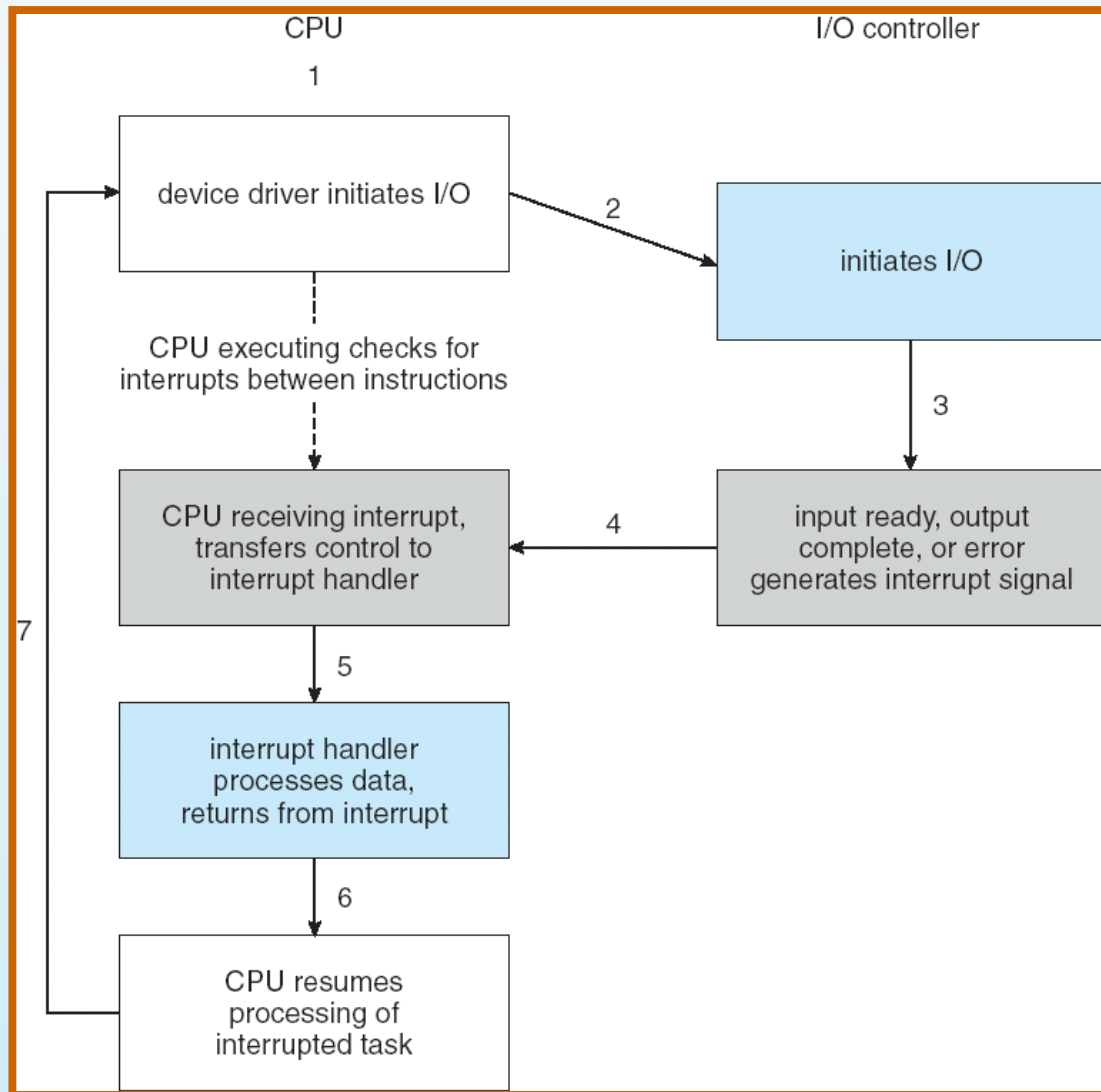
- **CPU Interrupt-request line triggered by I/O device**
- **Interrupt handler receives interrupts**
- **Maskable to ignore or delay some interrupts**
- **Interrupt vector to dispatch interrupt to correct handler**
  - **Based on priority**
  - **Some nonmaskable**
- **Interrupt mechanism also used for exceptions**







# Interrupt-Driven I/O Cycle





# Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts





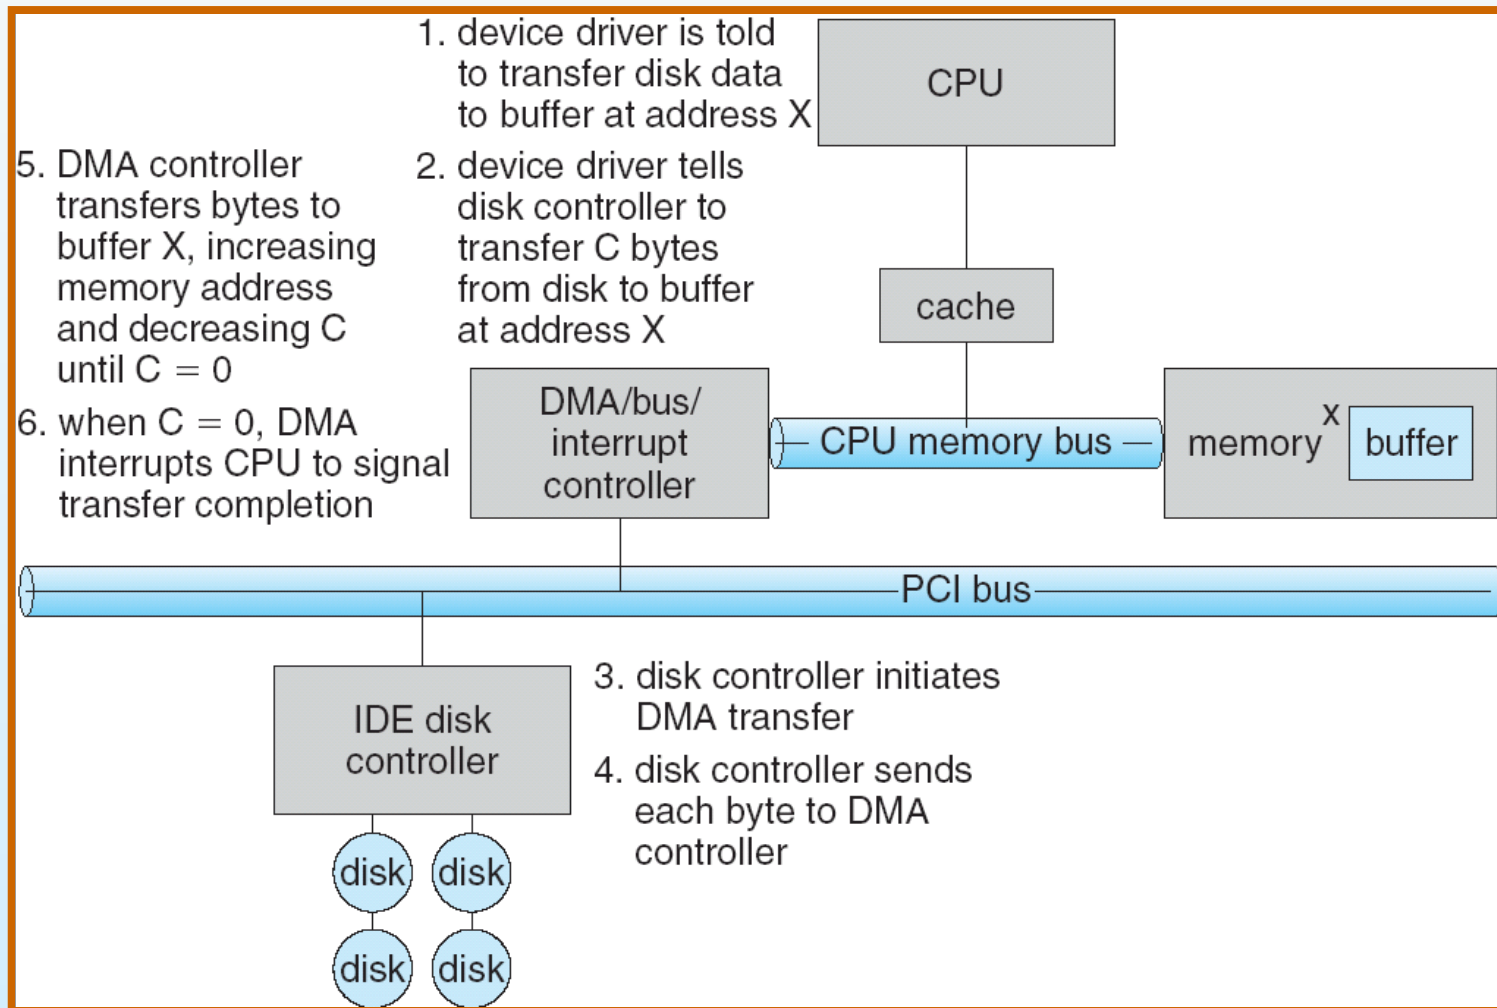
# Direct Memory Access

- Used for large data movement
- Requires DMA controller
- Transfer data directly between I/O device and memory
  - bypass CPU





# Six Step Process to Perform DMA Transfer





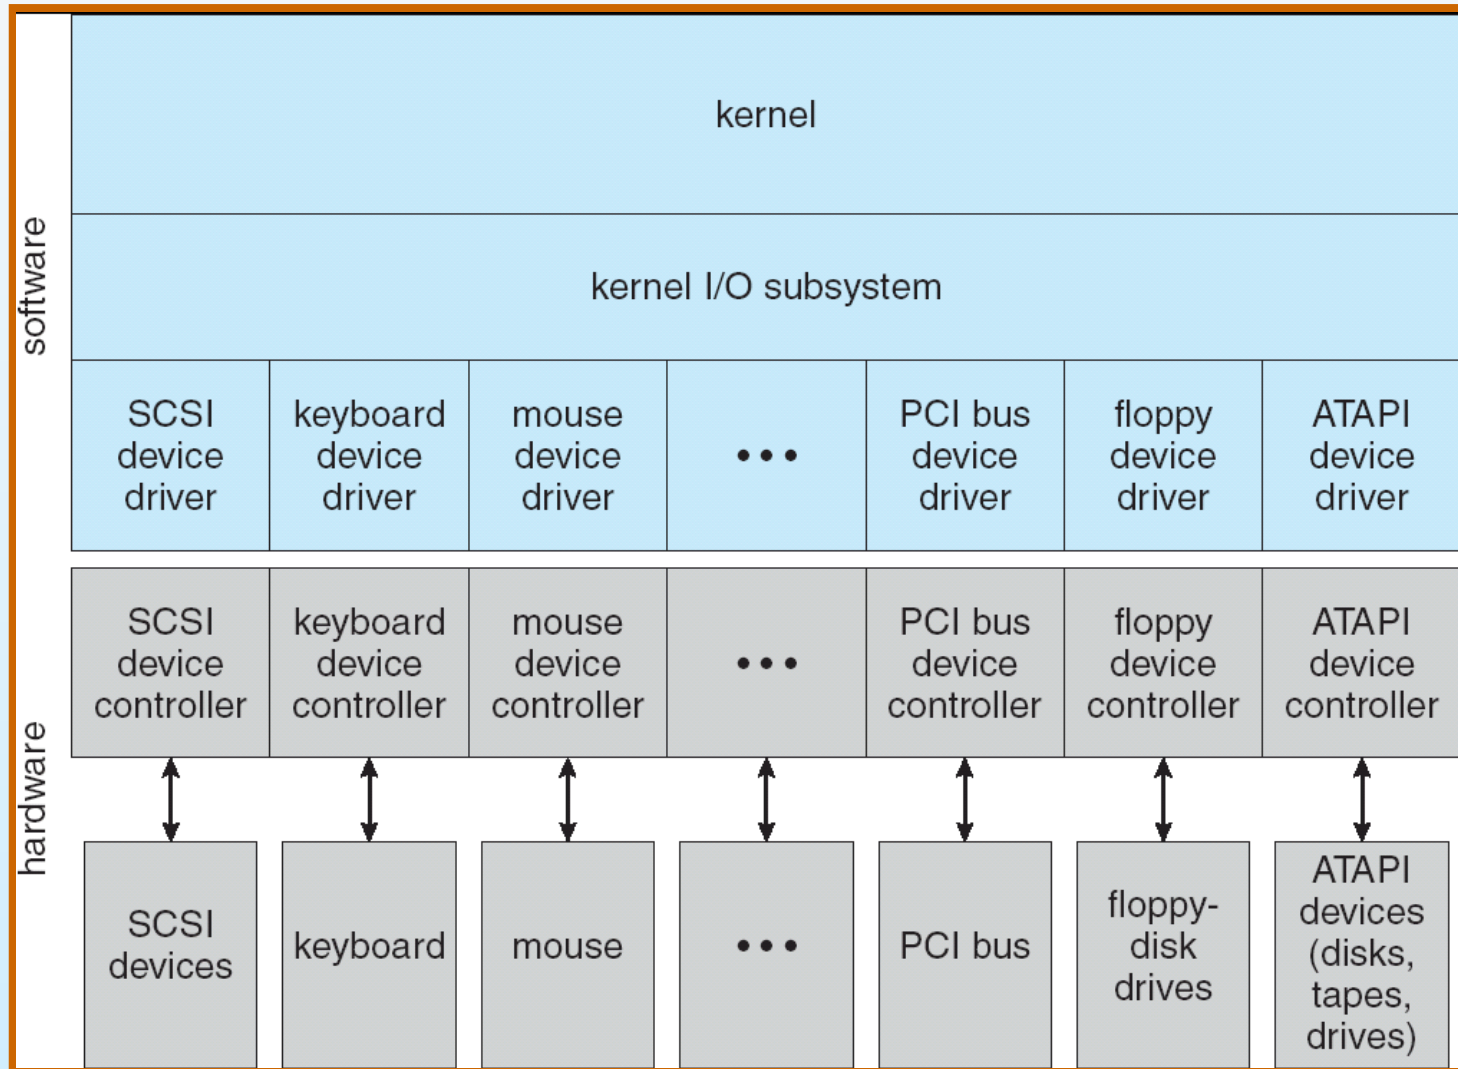
# Application I/O Interface

- **Devices vary in many dimensions**
  - Character-stream or block
  - Sequential or random-access
  - Sharable or dedicated
  - Speed of operation
  - read-write, read only, or write only
- **I/O system calls encapsulate device behaviors in generic classes**
  - `open()`, `read()`, `write()`, `close()`.....
  - `ioctl()` & `DeviceIoControl()`
- **Device-driver layer hides differences among I/O controllers from kernel**





# A Kernel I/O Structure





# Block and Character Devices

- **Block devices include disk drives**
  - **Commands include read, write, seek**
  - **Raw I/O or file-system access**
  - **Memory-mapped file access possible**
  
- **Character devices include keyboards, mice, serial ports**
  - **Commands include `get`, `put`**
  - **Libraries layered on top allow line editing**





# Network Devices

- Have own interface
- Unix and Windows NT/9x/2000/XP include socket interface
  - Separates network protocol from network operation
  - Includes `select` functionality







# Clocks and Timers

- Provide current time, elapsed time, timer
- Programmable interval timer used for timings, periodic interrupts
- `ioctl` (on UNIX) covers odd aspects of I/O such as clocks and timers





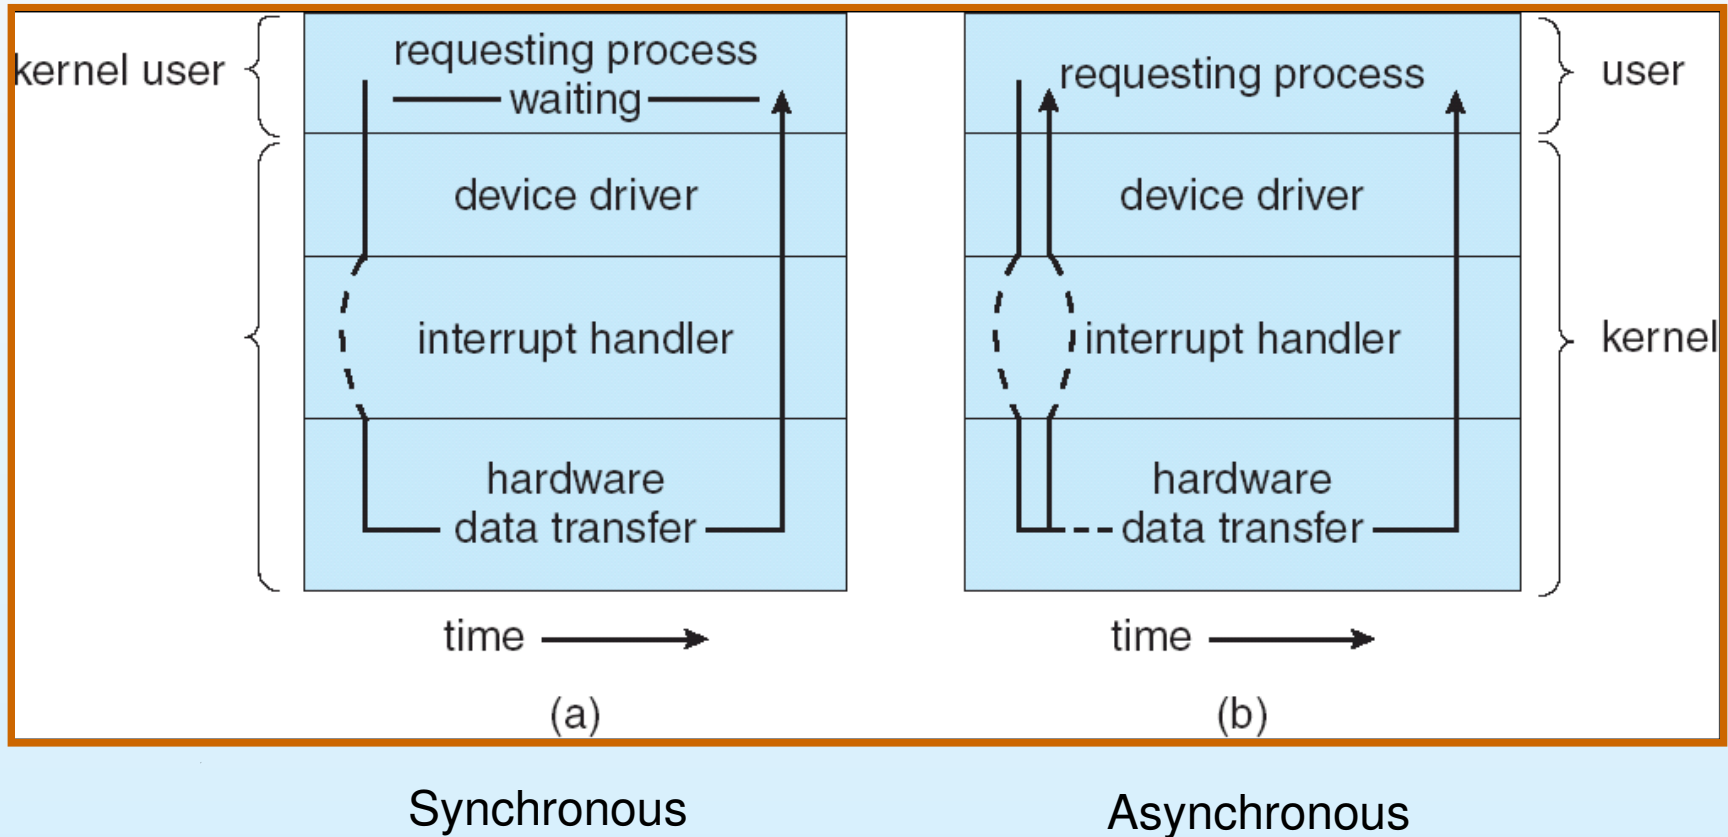
# Blocking and Nonblocking I/O

- **Blocking** - process suspended until I/O completed
  - Easy to use and understand
  - Insufficient for some needs
- **Nonblocking** - I/O call returns as much as available
  - Returns quickly with count of bytes read or written
  - Implemented via multi-threading
- **Asynchronous** - process runs while I/O executes
  - Difficult to use
  - I/O subsystem signals process when I/O completed





# Two I/O Methods





# Kernel I/O Subsystem

## ■ Scheduling

- Some I/O request ordering via per-device queue
- Some OSs try fairness

## ■ Buffering - store data in memory while transferring between devices

- To cope with device speed mismatch
  - ▶ Double buffering
- To cope with device transfer size mismatch
- To maintain “copy semantics”





# Kernel I/O Subsystem

- **Caching - fast memory holding copy of data**
  - Always just a copy
  - Key to performance
  
- **Spooling - hold output for a device**
  - If device can serve only one request at a time
  - i.e., Printing
  
- **Device reservation - exclusive access to a device**
  - System calls for allocation and deallocation
  - Watch out for deadlock





# Error Handling

- OS can recover from disk read, device unavailable, transient write failures
- Most return an error number or code when I/O request fails
- System error logs hold problem reports





# I/O Protection

- **User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions**
  - **All I/O instructions defined to be privileged**
  - **I/O must be performed via system calls**
    - ▶ **Memory-mapped and I/O port memory locations must be protected too**
  - **However, DirectX**





# Kernel Data Structures

- **Kernel keeps state info for I/O components**
  - **Open file tables**
  - **Network connections**
  - **Character device state**
  
- **Many, many complex data structures to track buffers, memory allocation, “dirty” blocks**
  
- **Some use object-oriented methods and message passing to implement I/O**







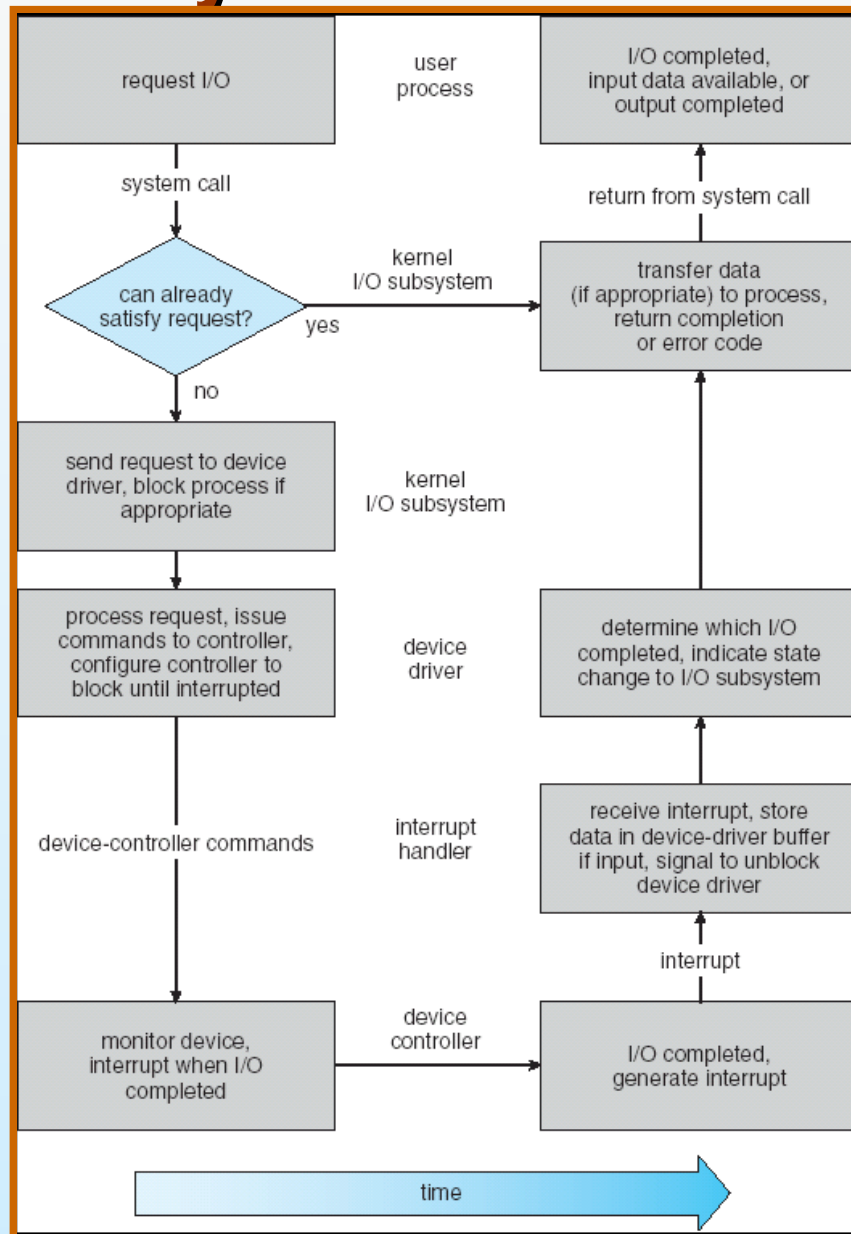
# I/O Requests to Hardware Operations

- Consider reading a file from disk for a process:
  - Determine device holding file
  - Translate name to device representation
  - Physically read data from disk into buffer
  - Make data available to requesting process
  - Return control to process





# Life Cycle of An I/O Request





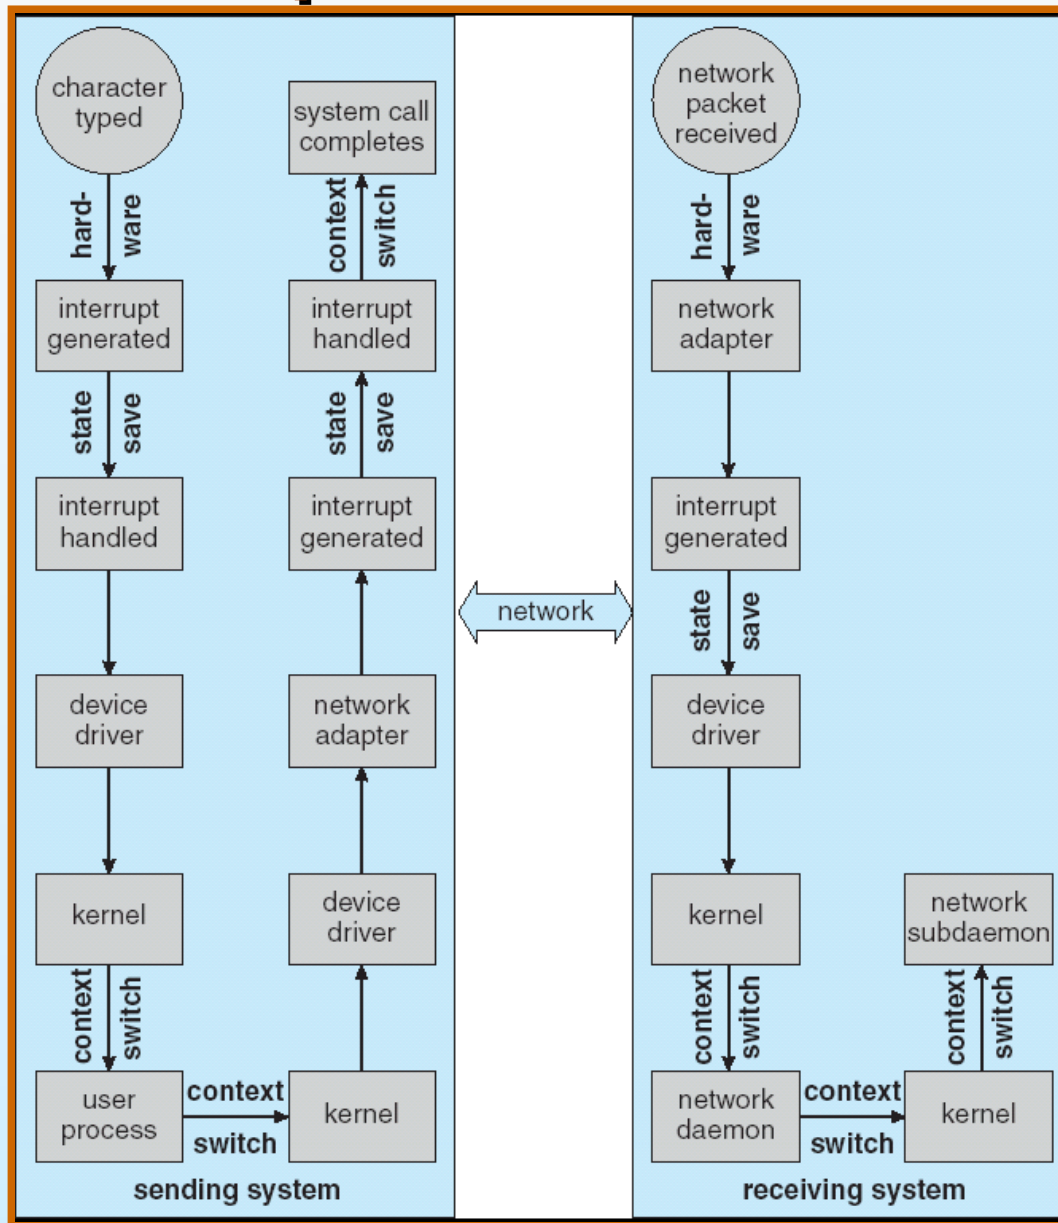
# Performance

- **I/O, a major factor in system performance:**
  - **Heavily demands CPU to execute device driver, kernel I/O code**
  - **Context switches due to interrupts**
  - **Data copying**
  - **Network traffic especially stressful**





# Intercomputer Communications





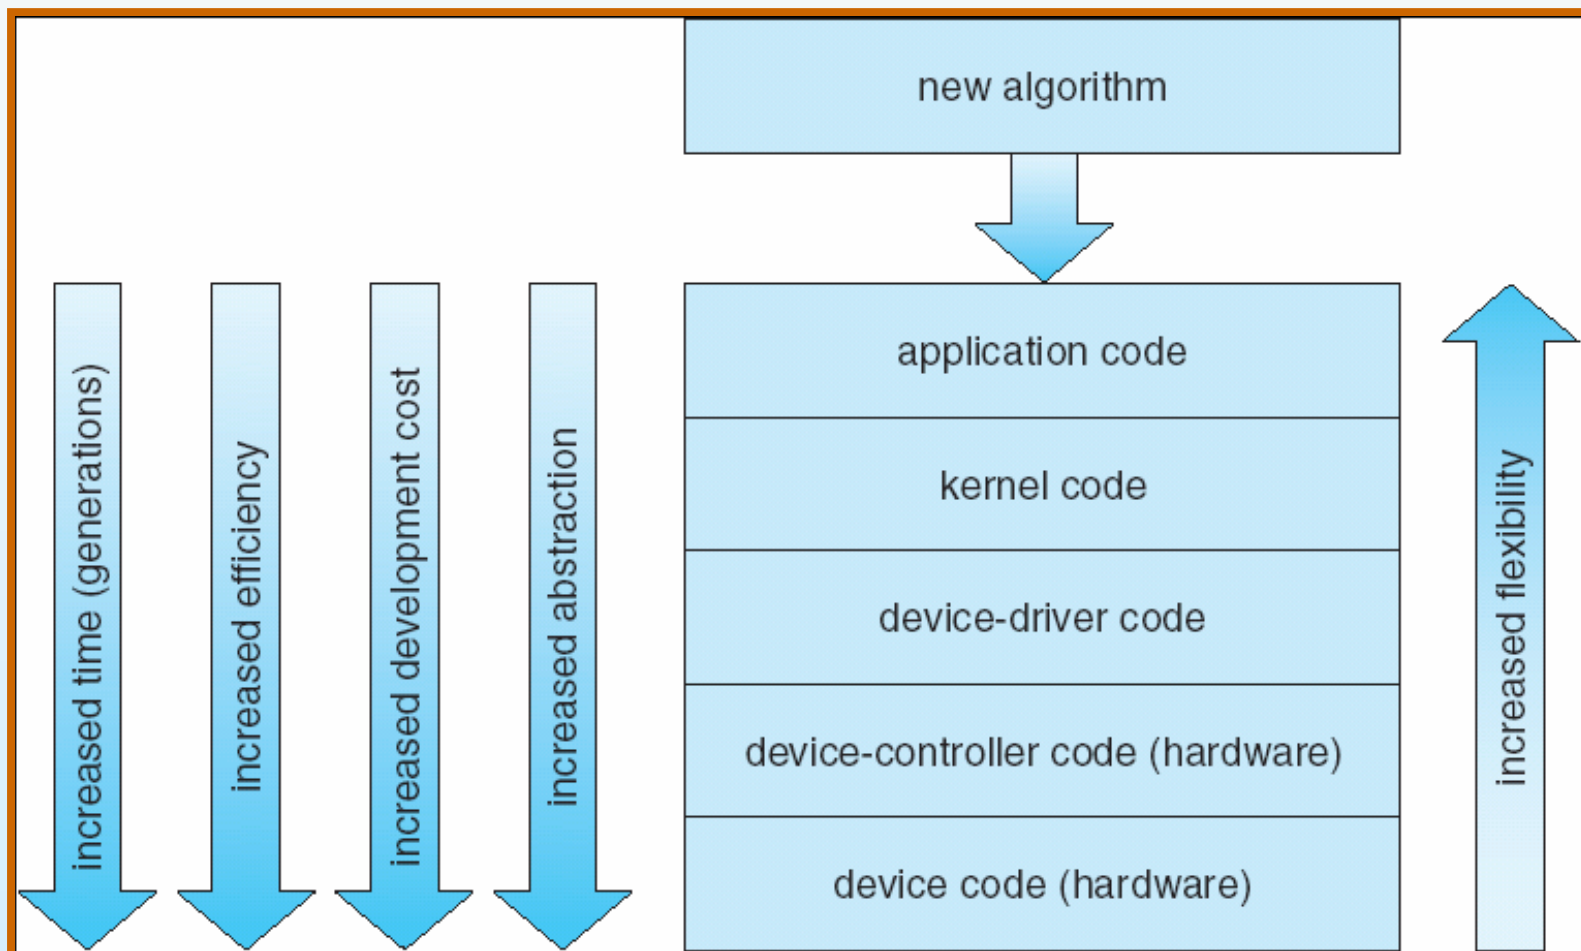
# Improving Performance

- Reduce number of context switches
- Reduce data copying
- Reduce interrupts by using large transfers, smart controllers, polling
- Use DMA
- Balance CPU, memory, bus, and I/O performance for highest throughput





# Device-Functionality Progression



# End of Chapter 13

